



Machine learning in space physics

Savvas Raptis

Space Physics I
KTH Royal Institute of Technology
10 / 10 / 2022 | 13:00 – 15:00



Introduction to Neural Networks & Applications in Space Physics

Savvas Raptis

Space Physics I
KTH Royal Institute of Technology
10 / 10 / 2022 | 13:00 – 15:00

What is machine learning & A.I ? (1)

*Making the computer “**learn**” without being explicitly programmed*

What is machine learning & A.I ? (2)

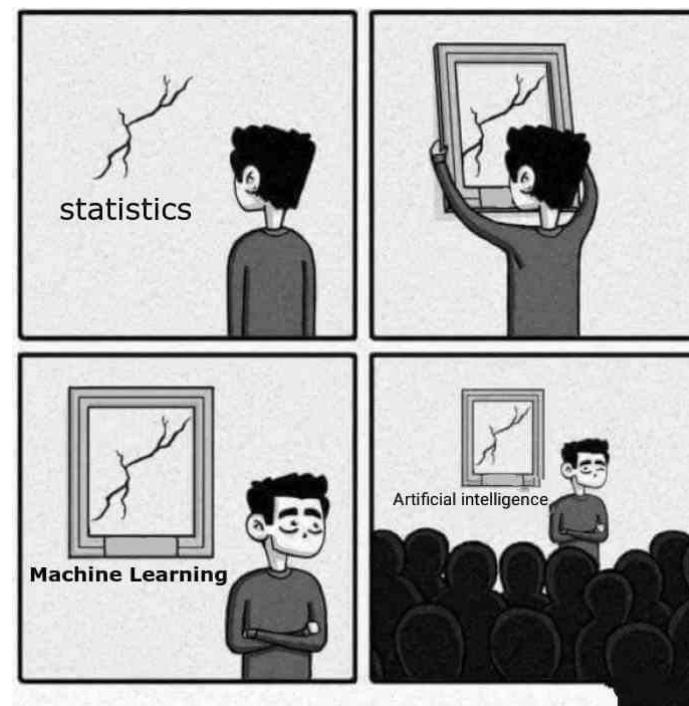
AI (Artificial Intelligence) : A term most of the times used to draw attention (at least for the moment).

ML (Machine Learning) : A program designed to “learn” something based on a set of rules and/or data.

Difference between machine learning and AI:

If it is written in Python, it's probably machine learning

If it is written in PowerPoint, it's probably AI

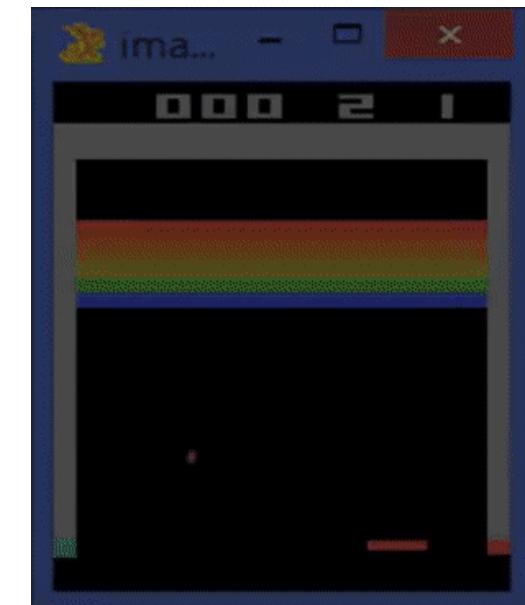
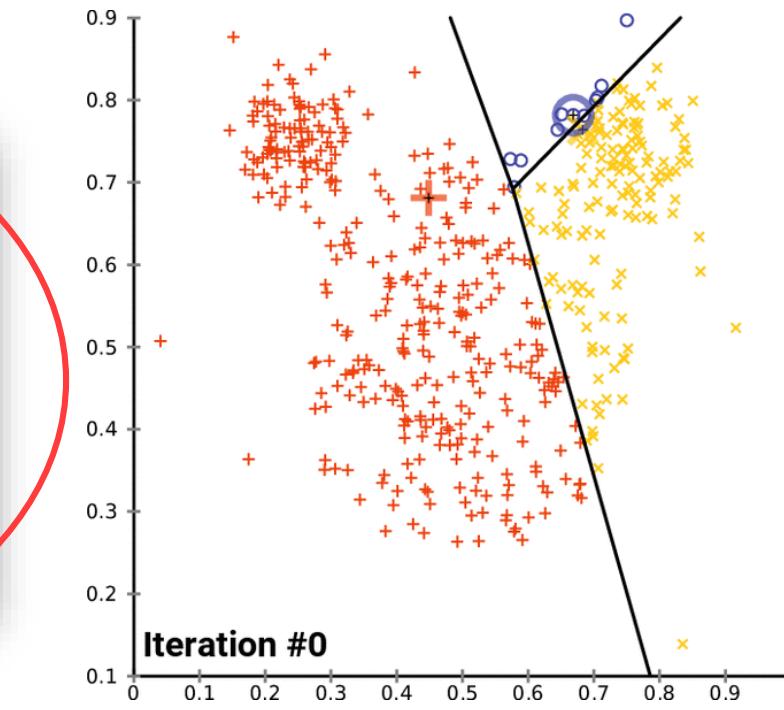
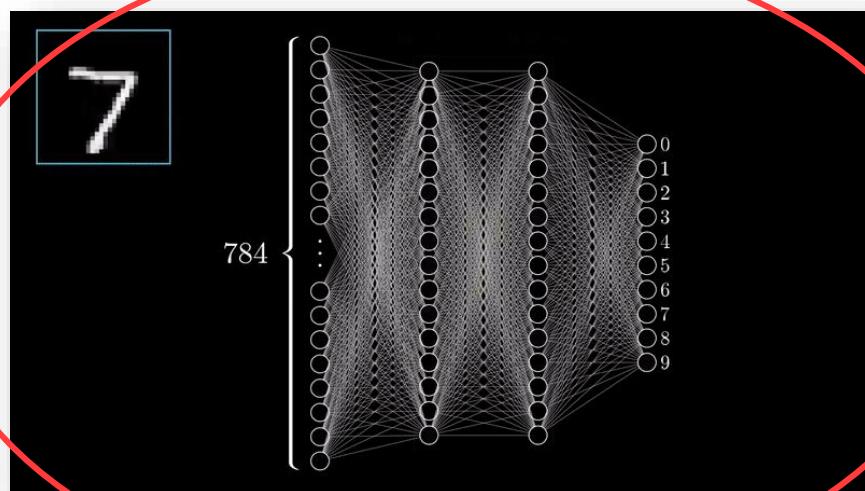


Different categories of Machine Learning

Supervised learning Given **Input** and **Output** - learn the mapping by minimizing an error between the initially random output with the given output. Obtaining a relation.

Unsupervised Learning: Given **Input** but not output - Either find features or discover patterns.

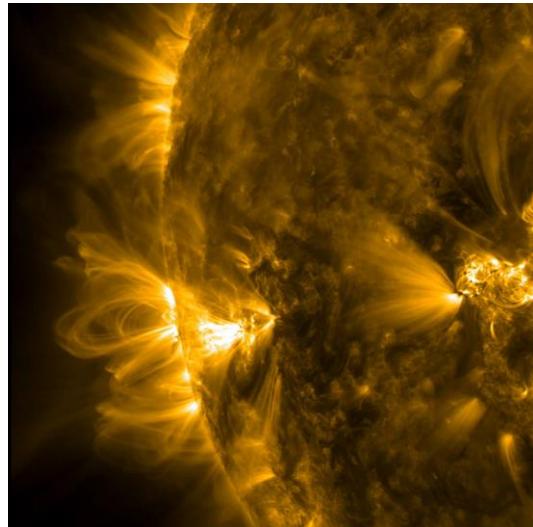
Reinforcement learning: **No data** needed, just a value that we want to maximize (e.g. score in game) and a set of actions/environment. The goal is to find a sequence of actions that maximizes the reward.



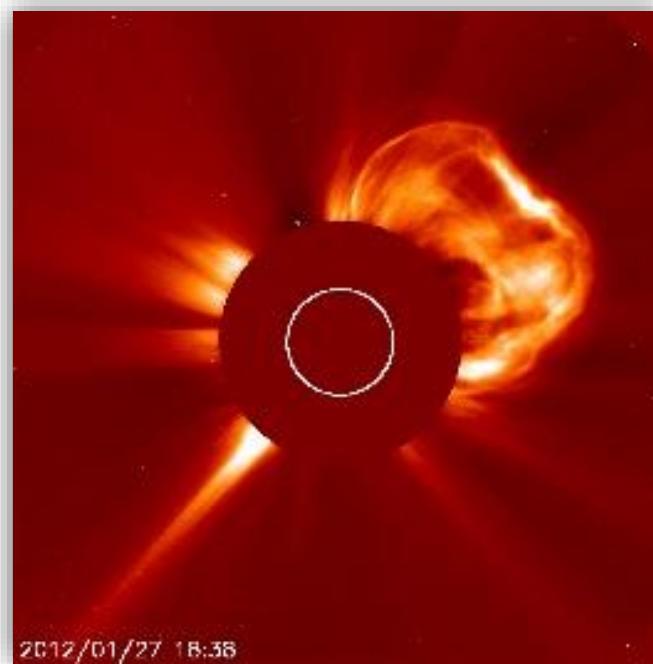
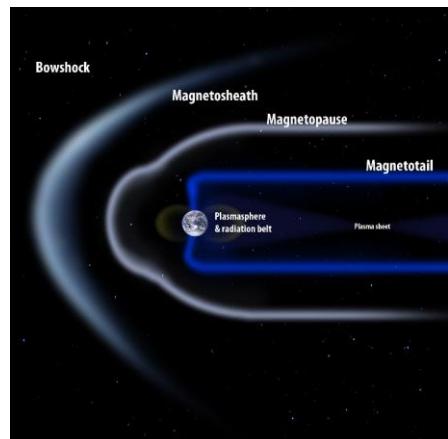
Relevance to Space Physics

Can we predict SEPs
(Solar Energetic Particles) ?

Can we reconstruct coronal loops?

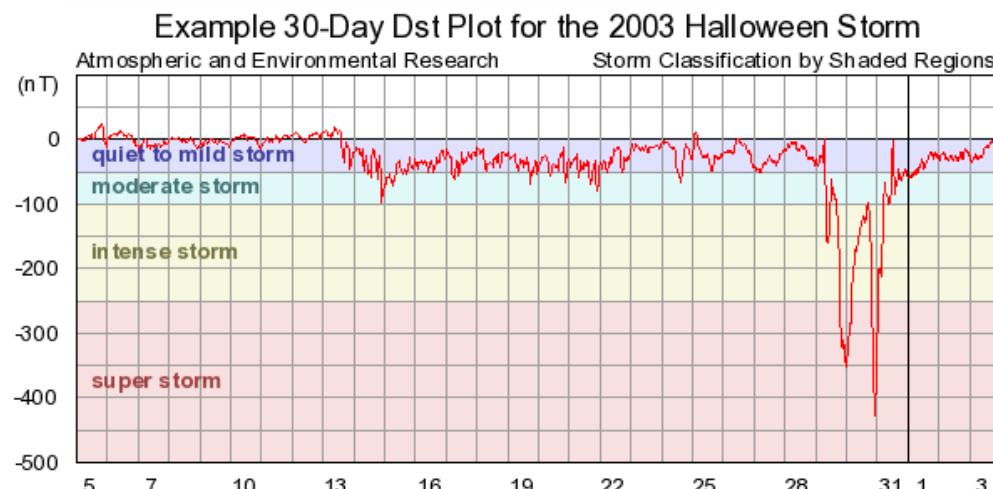


Can we classify plasma environments?



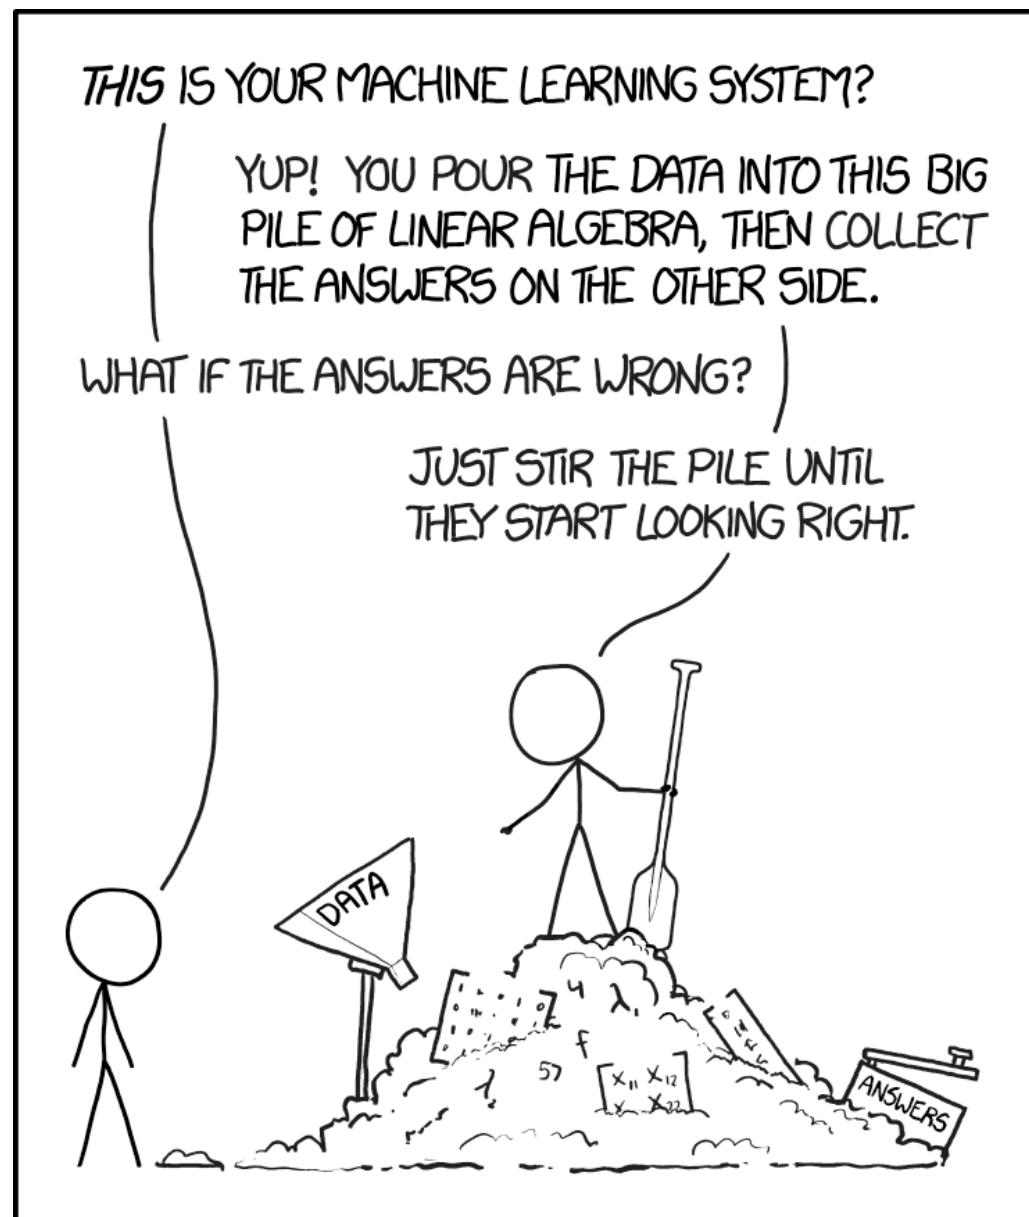
Can we predict CMEs?

Can we predict Geomagnetic storms?



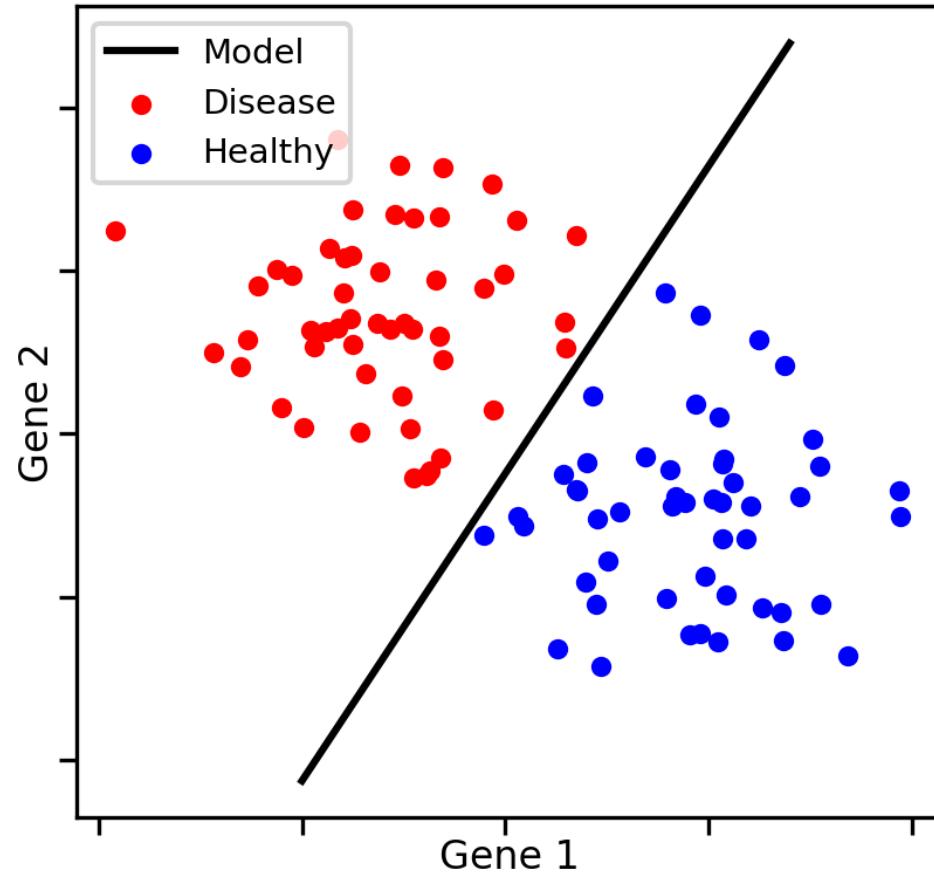
Supervised Learning using Neural Networks

TLDR version

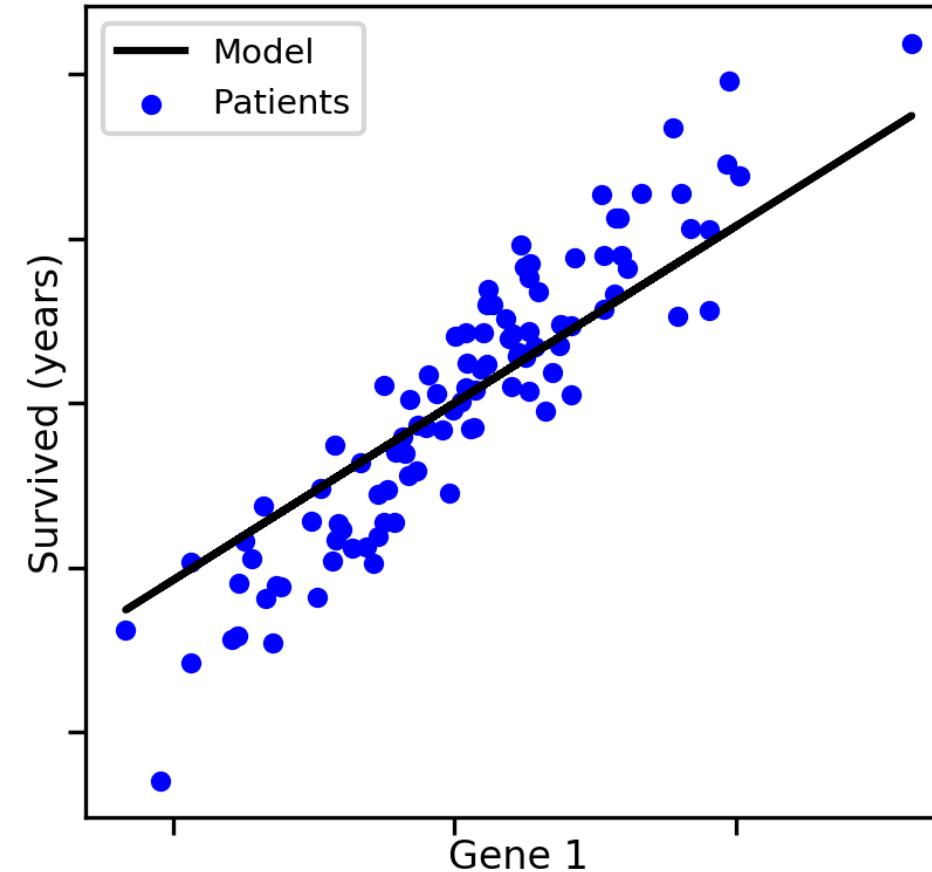


Type of Machine Learning Problems

Classification



Regression



Let's think of a problem (classification)



Data: Thousands of labeled numbers 1-9

Problem: Can we make a model that given the image can recognize the number ?

(one) Solution: Neural Networks!

Another problem to think of (regression)



Data: characteristic of a house (e.g., # of rooms)

Problem: What should the price be ?

(one) Solution: Neural Networks!

Outline

How to understand Neural Network modeling:

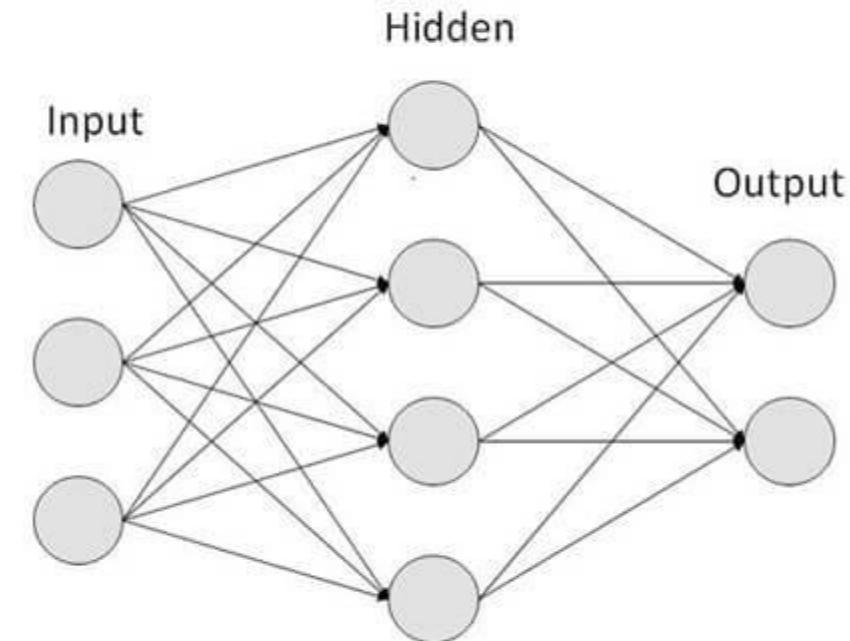
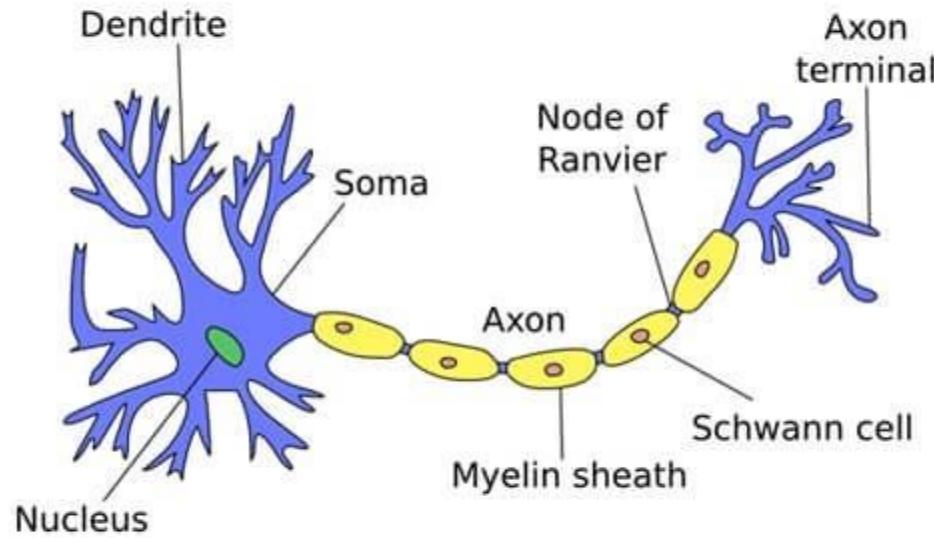
- Concept (basic principles) & Building (what are the components?)
- Train (How does it learn?)
- Evaluate (Did it really learn anything?)

Outline

How to understand Neural Network modeling:

- Concept (basic principles) & Building (what are the components?)
- Train (How does it learn?)
- Evaluate (Did it really learn anything?)

Biological vs Artificial Neural Networks



Neural Networks

Input/Output pairs

e.g.

Output = Price of house

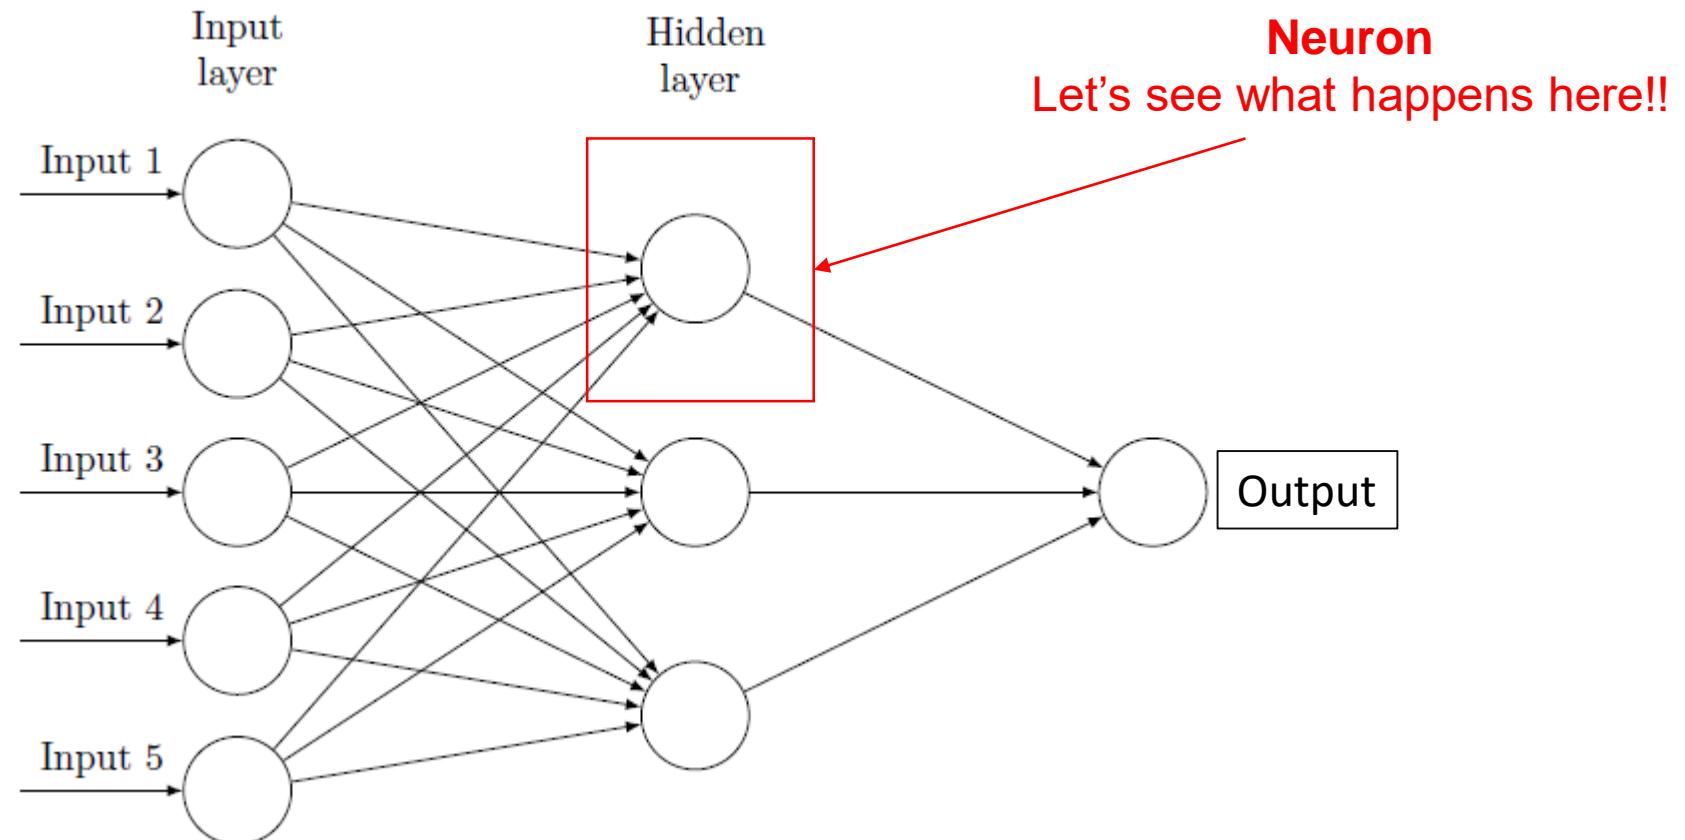
Input1 = Square meters

Input2 = # rooms

Input3 = Proximity to center

Input4 = Age it was built

...



Neuron

Let's see what happens here!!

A Neural Network Input and Output

$$y = f(x_1 w_1 + x_2 w_2 + \dots + x_n w_n + b)$$

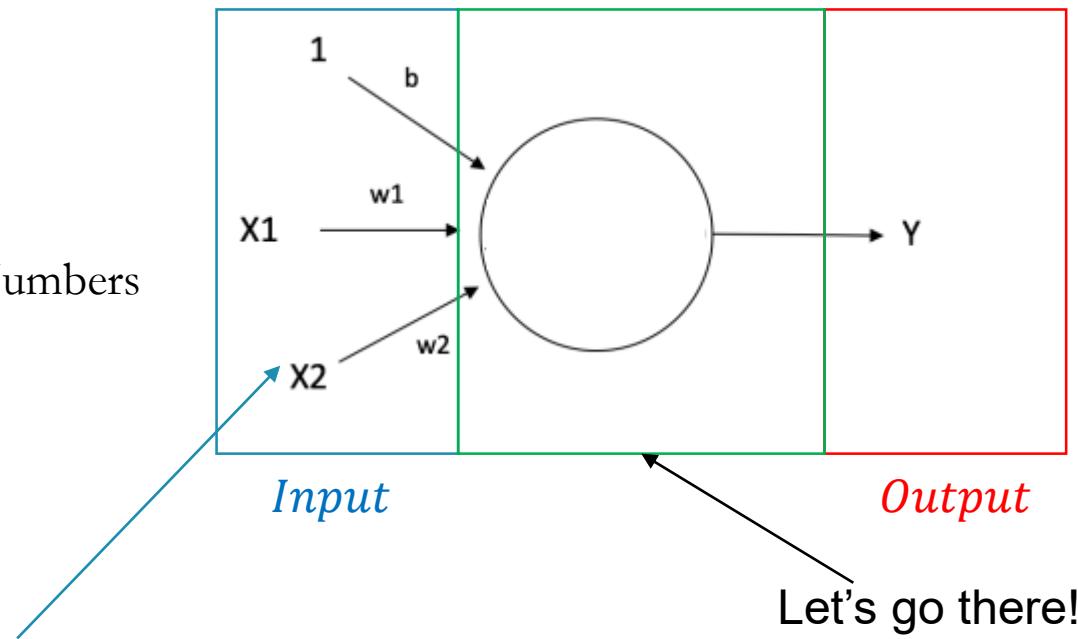
y : Output of Neuron

x_i : Inputs of Neuron

w_i : Weights of each Input

b : Bias for each neuron

Random Numbers

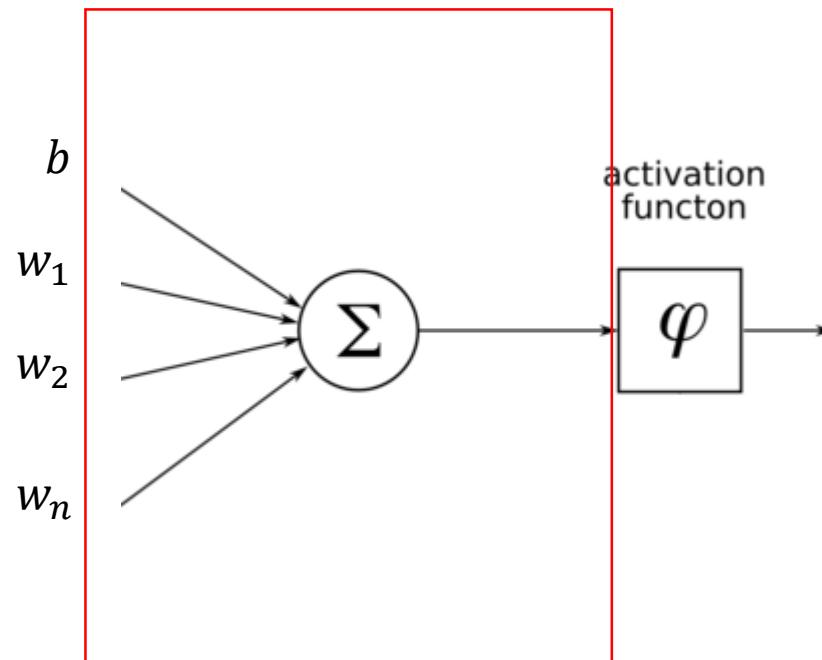


For example

x_1 = Square meters - $w_1 = 0.7$ (more important)

x_2 = Proximity to center - $w_2 = 0.3$ (less important)

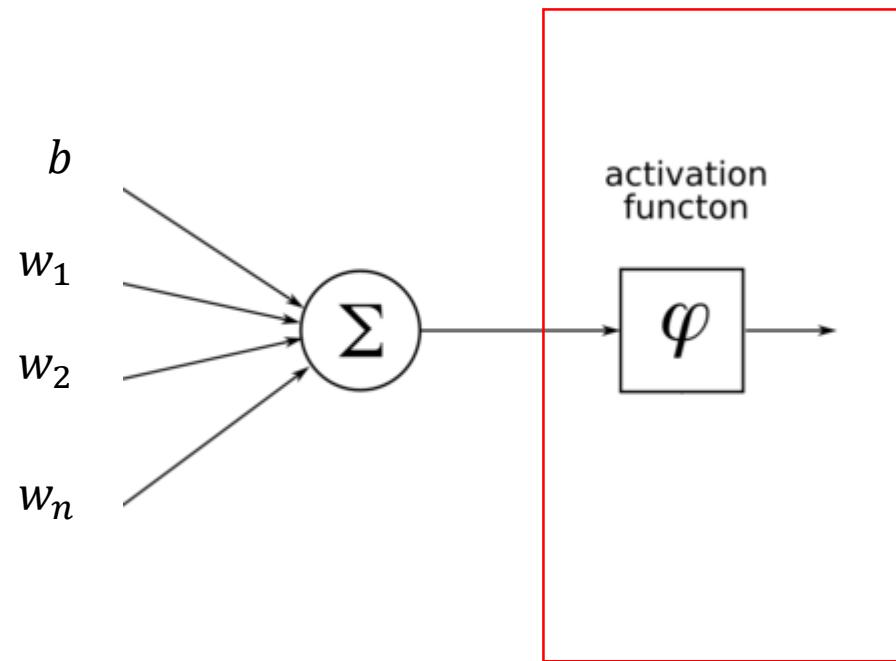
Inside the Neuron



Sum of all Data(x_i), Weights (w_i) and Biases (b)

$$y = \sum x_i w_i + b$$

Activation Function

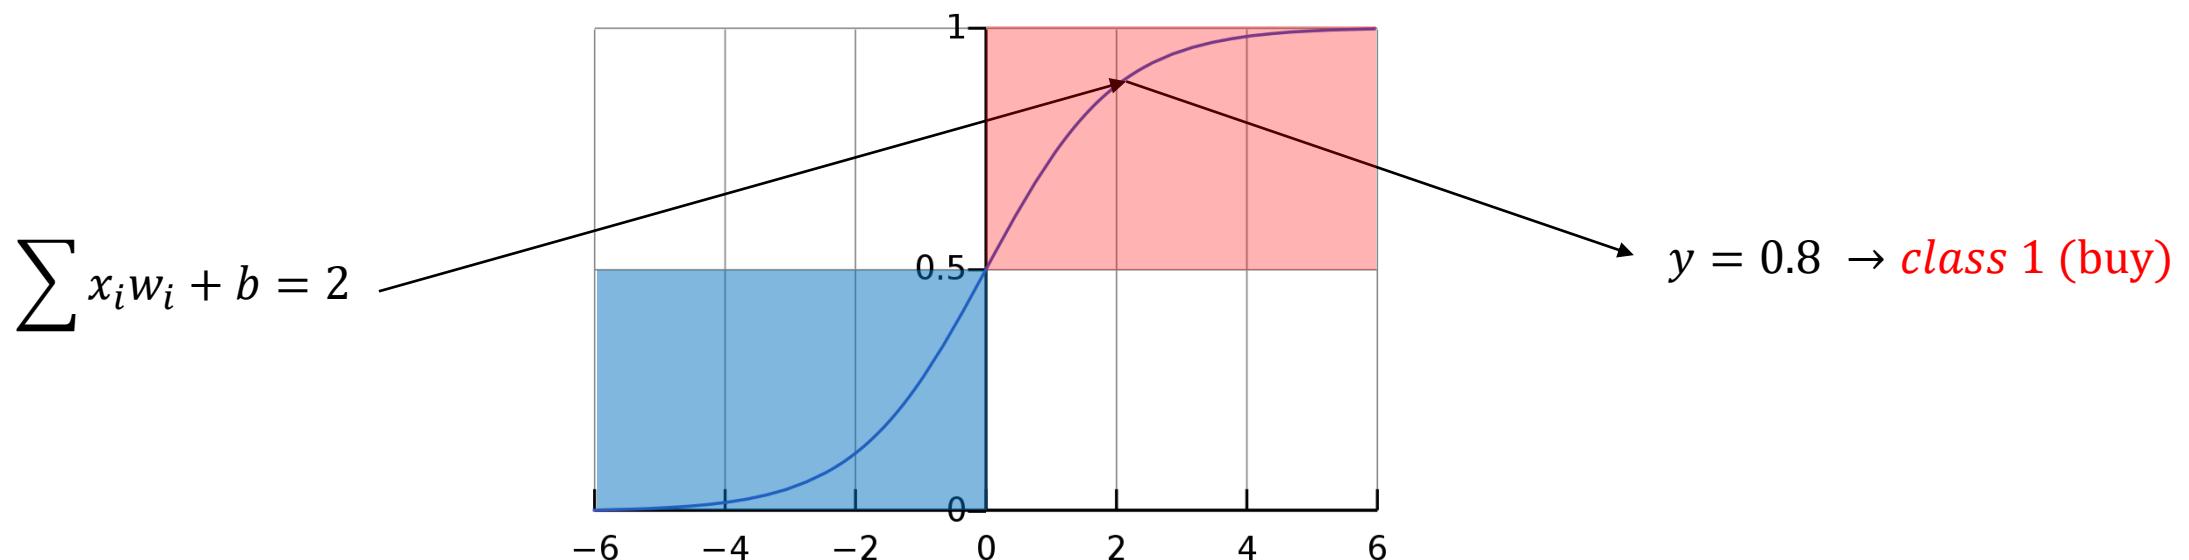


Apply $f(y)$ depending on goal

Activation function Examples

Goal: Classification (e.g., 0/1)

e.g. Should I buy the house?
Yes (class 1)
No (class 0)

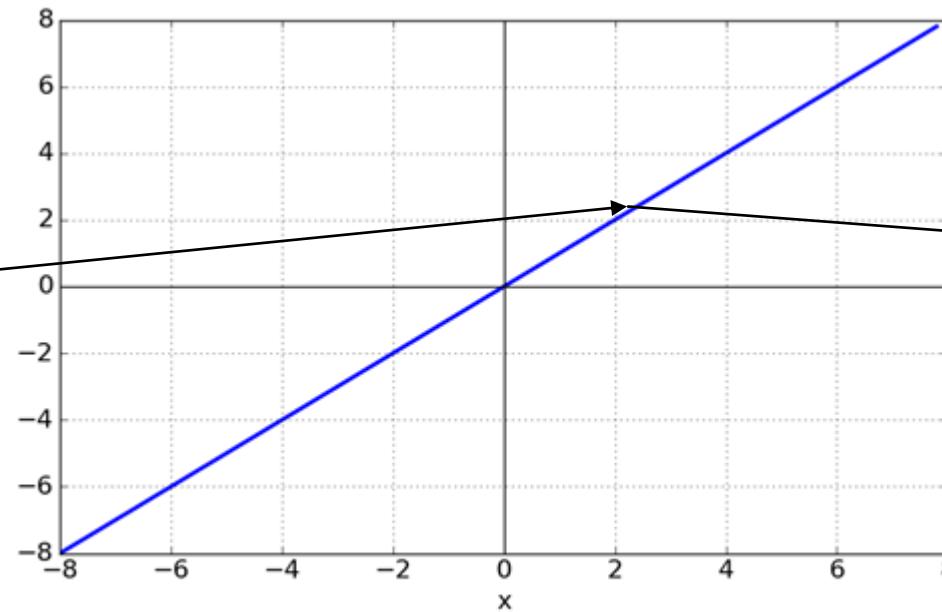


Activation function Examples

Goal: Regression (Continuous value)

e.g. How many million SEK will a house cost ?

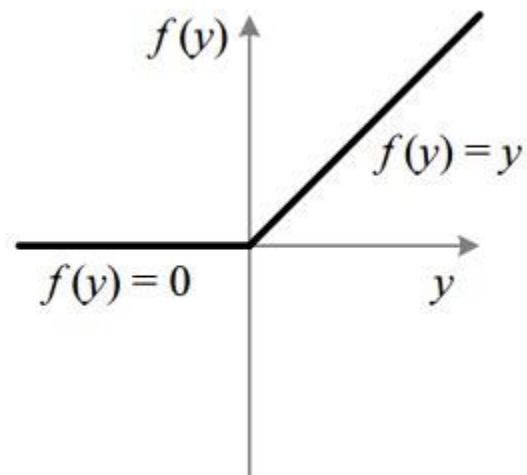
$$\sum x_i w_i + b = 2$$



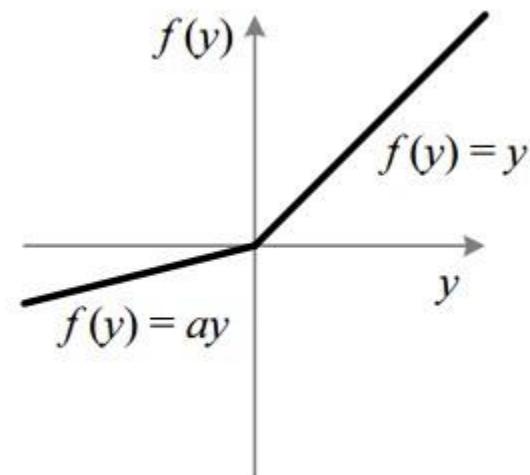
$y = 2$ million SEK price

Hidden layer activation functions

Goal: Add complexity to our mapping



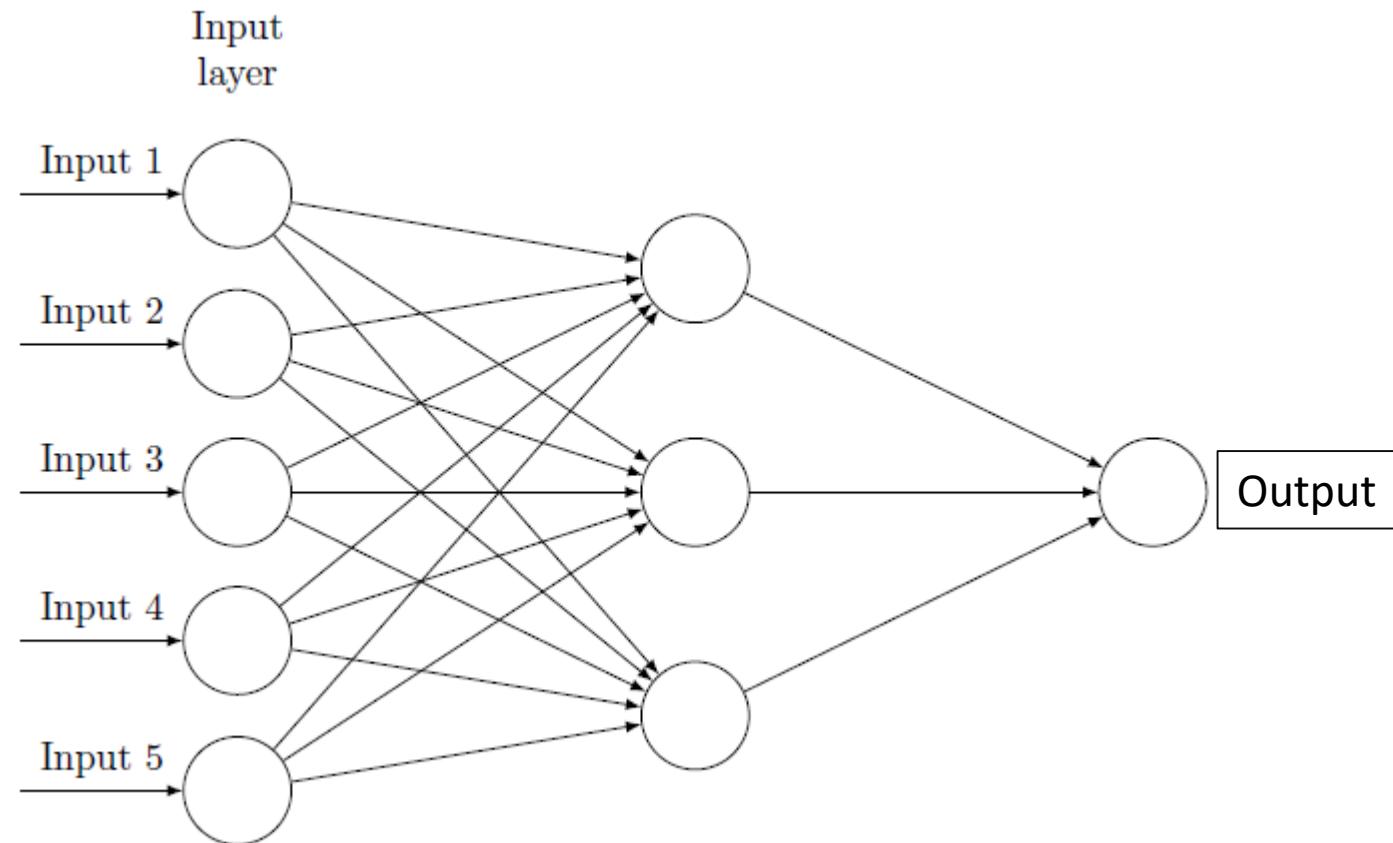
Rectified Linear Units (ReLU)



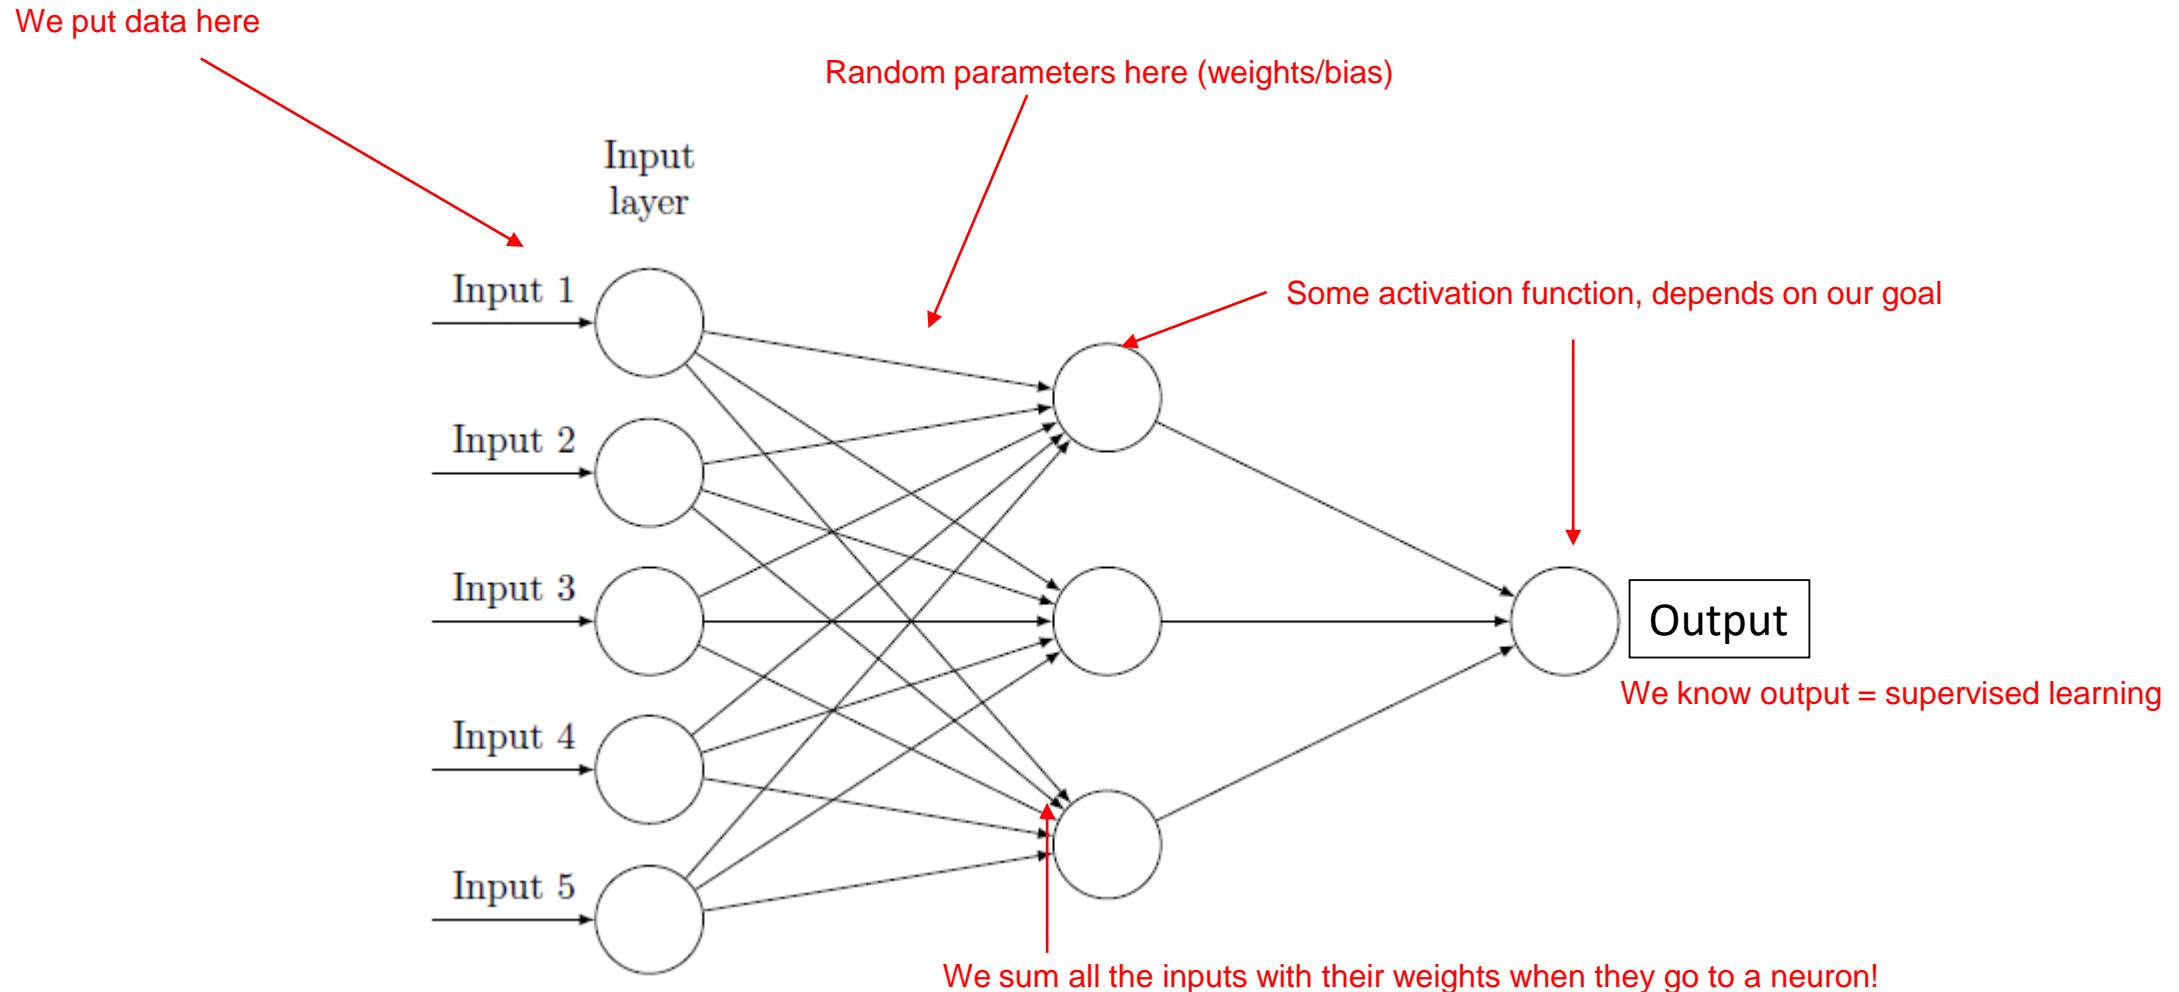
LeakyReLU

Question time!
What do we know so far ?

Neural Networks



Neural Networks



Outline

How to understand Neural Network modeling:

- Concept (basic principles) & Building (what are the components?)

Answer:

- Train (How does it learn?)
- Evaluate (Did it really learn anything?)

Outline

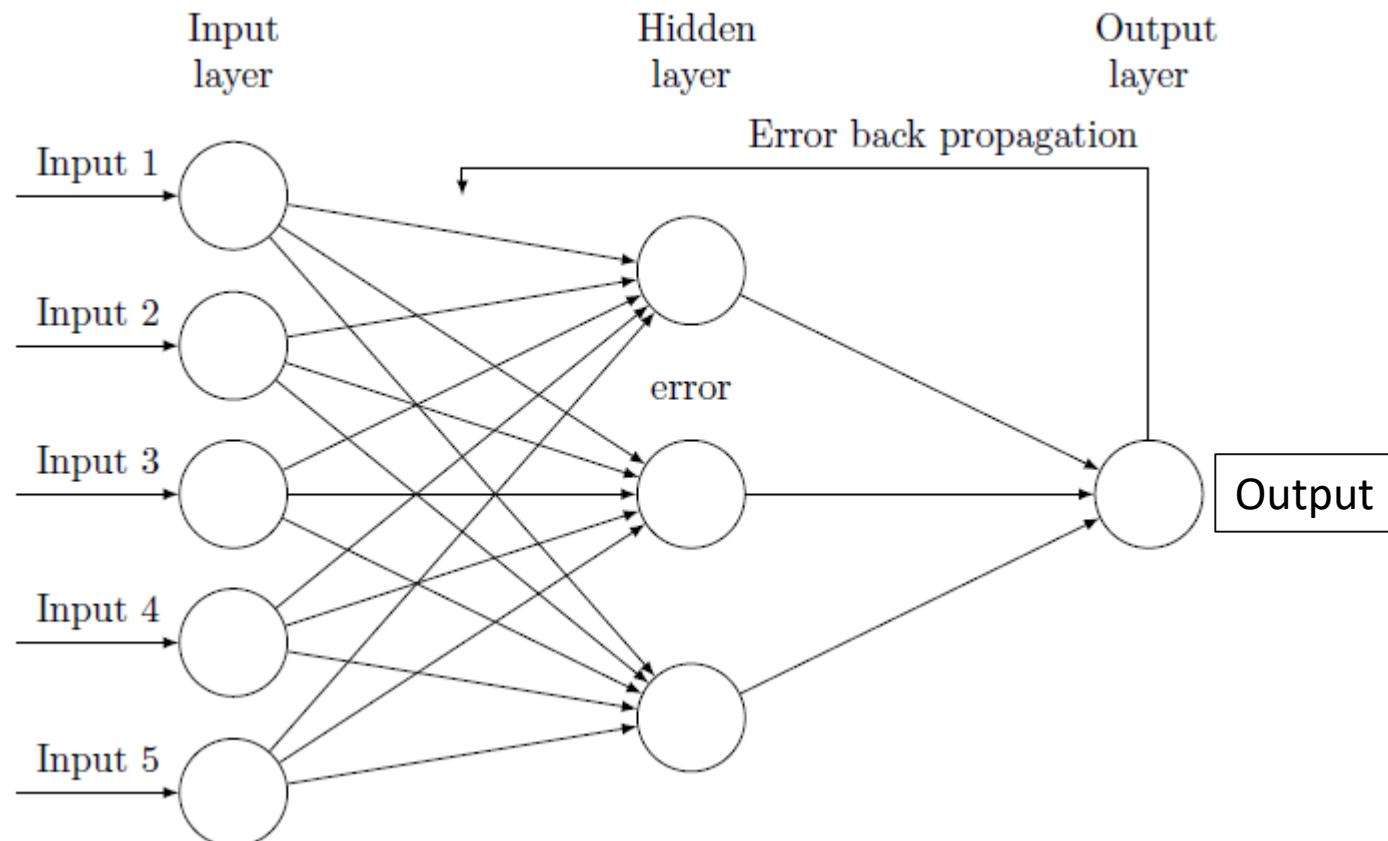
How to understand Neural Network modeling:

- Concept (basic principles) & Building (what are the components?)

Answer: Yes! We know that now! We have a randomly parametrized model..... But ok, then what ? It should just spit out random results now right ?

- Train (How does it learn?)
- Evaluate (Did it really learn anything?)

Neural Networks & Backpropagation



- *Random parameters = random result*
- *Error back propagation = change parameters based on data = modeling*

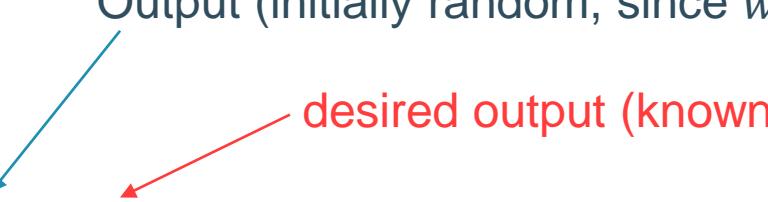
Cost (error) function

For example, mean squared error:

$$J = \frac{1}{N} \sum_{i=1}^N (Y' - Y)^2$$

Output (initially random, since w, b are random)

desired output (known)



The goal is to minimize/maximize a cost function.

Question: How do we minimize this? What do we have to do when we are asked to find the minimum of a function ?

With an algorithm that uses the minima of the function: **Gradient Descent**.

2 things we care about

Where to go and how fast to go there ?

Where = Opposite of derivative/slope of the error

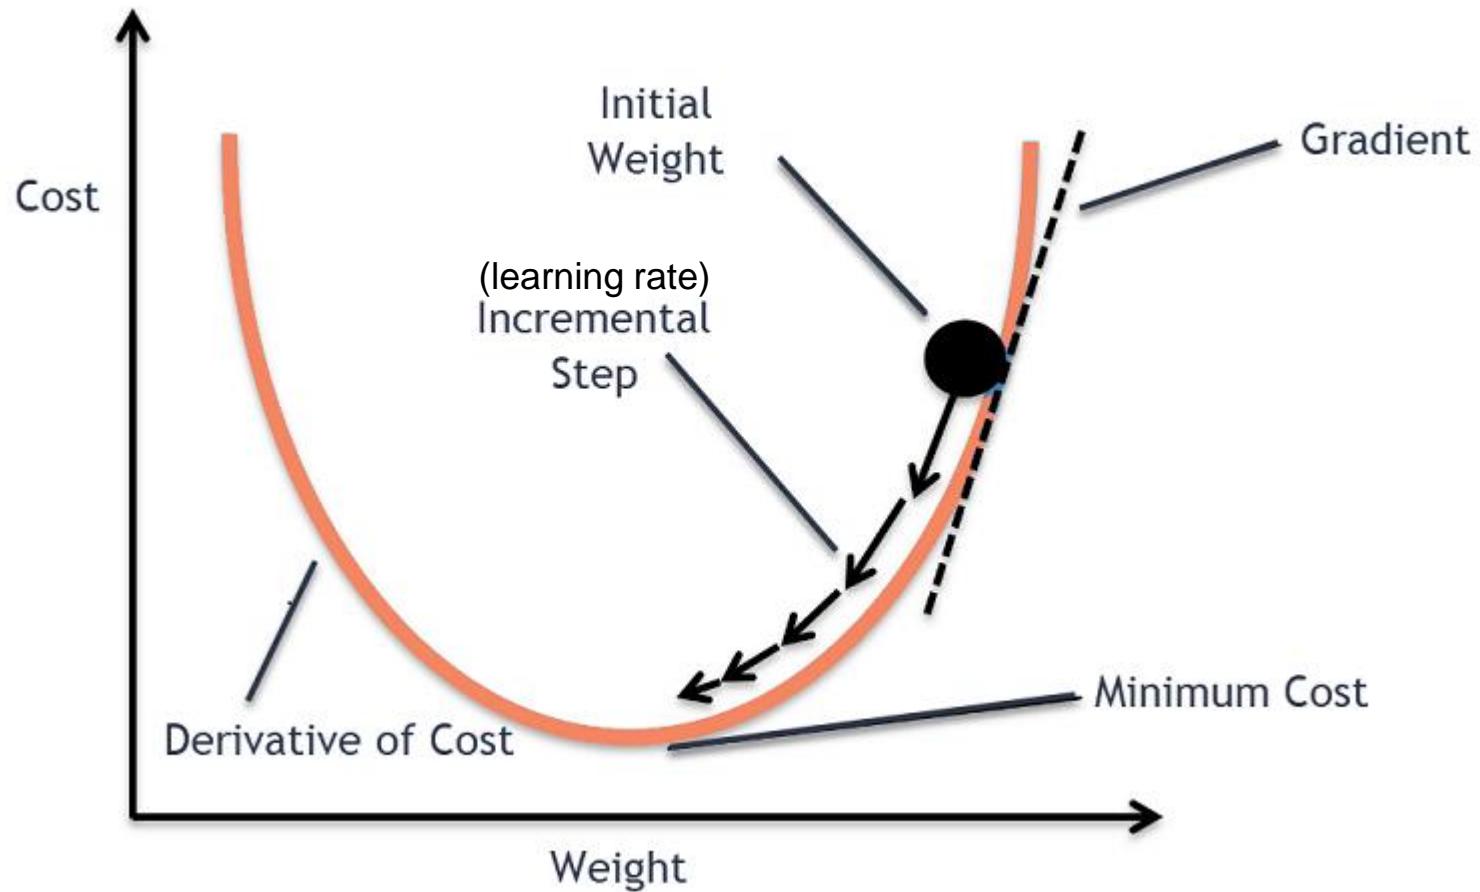
$$\theta_j = \theta_j - \alpha \frac{\partial}{\partial \theta_j} J$$

Where, j = cost function

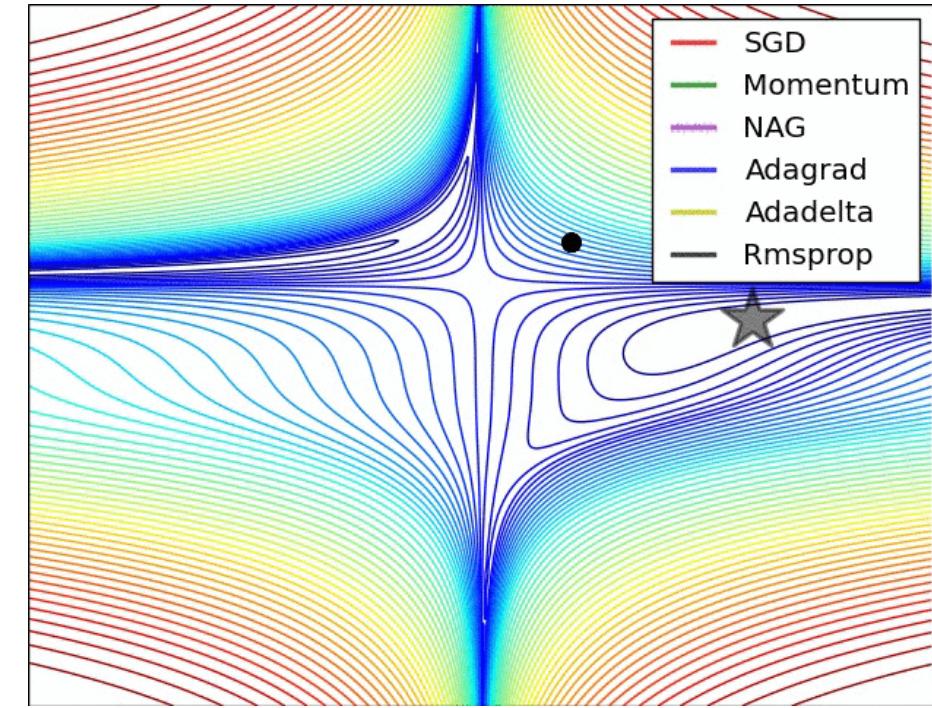
How fast = Learning rate (numerical factor)

<https://developers.google.com/machine-learning/crash-course/fitter/graph>

Error Back Propagation (Gradient Descent)



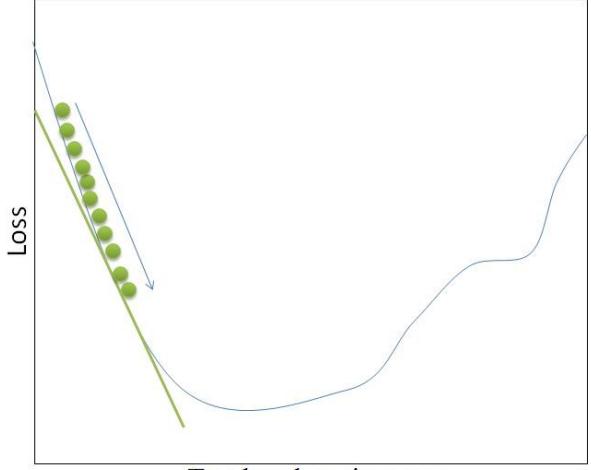
Automatically done by
“optimizers”



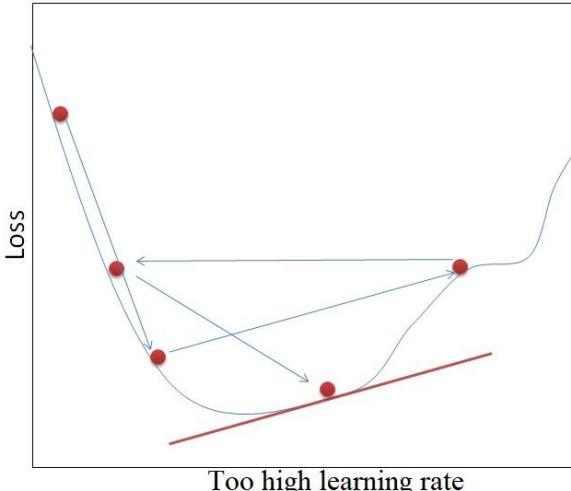
Back-propagation = calculating the derivatives.

Gradient descent = Descending through the gradient to change parameters.

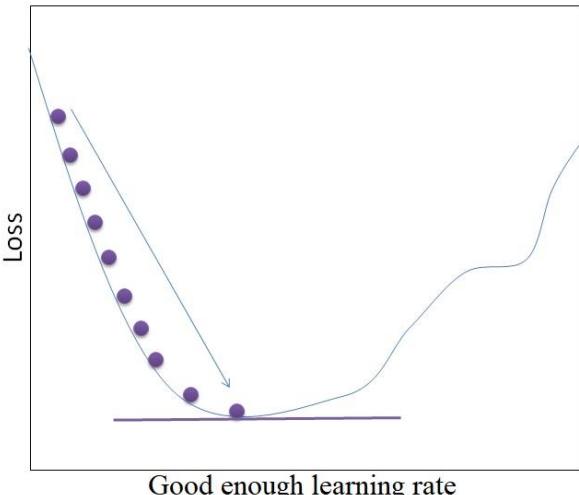
Why Optimizers? (learning rate + momentum)



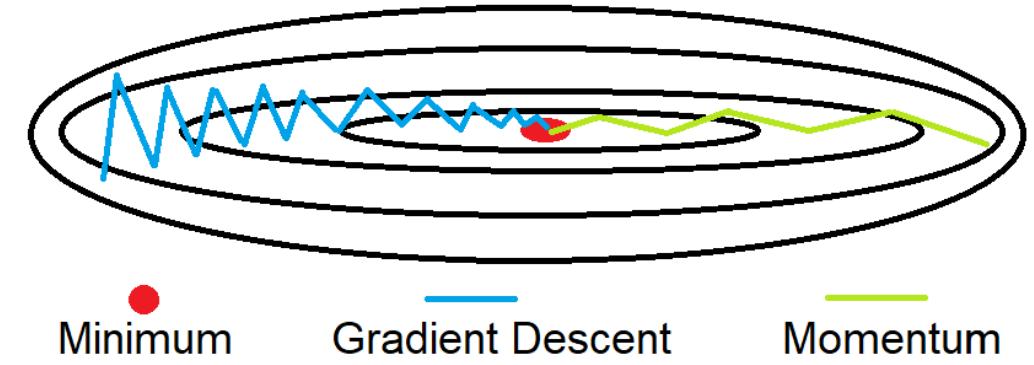
Too low learning rate



Too high learning rate



Good enough learning rate



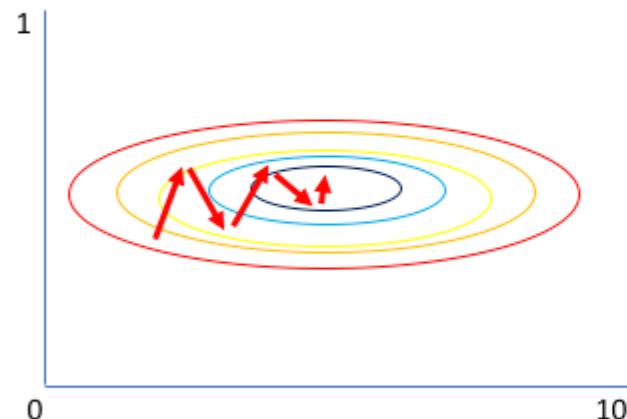
Minimum

Gradient Descent

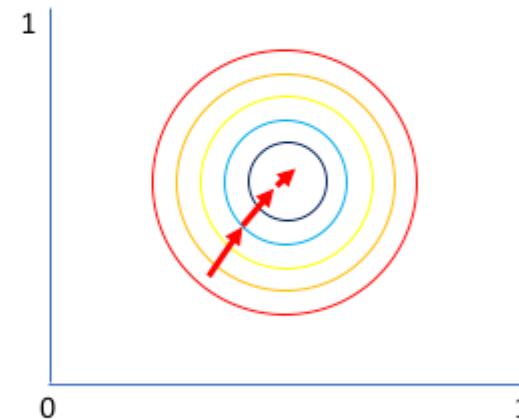
Momentum

Normalization of data – Features

Why normalize?



Gradient of larger parameter
dominates the update



Both parameters can be
updated in equal proportions

e.g.,

Number of rooms
[0 ... 5]

Square meters
[0 ... 300]

Outline

How to understand Neural Network modeling:

- ~~Concept (basic principles) & Building (what are the components?)~~

Answer: Yes! We know that now! We have a randomly parametrized model..... But ok, then what ? It should just spit out random results now right ?

- Train (How does it learn?)

Answer: ...

- Evaluate (Did it really learn anything?)

Outline

How to understand Neural Network modeling:

- ~~Concept (basic principles) & Building (what are the components?)~~

Answer: Yes! We know that now! We have a randomly parametrized model..... But ok, then what ? It should just spit out random results now right ?

- Train (How does it learn?)

Answer: Now we calculate the derivative of the error and we change the parameters of the network. Sure... But how do we test that thing ? I mean... It learned something that we explicitly trained. But everyone can do that... is it smart enough to deal with brand new information ?

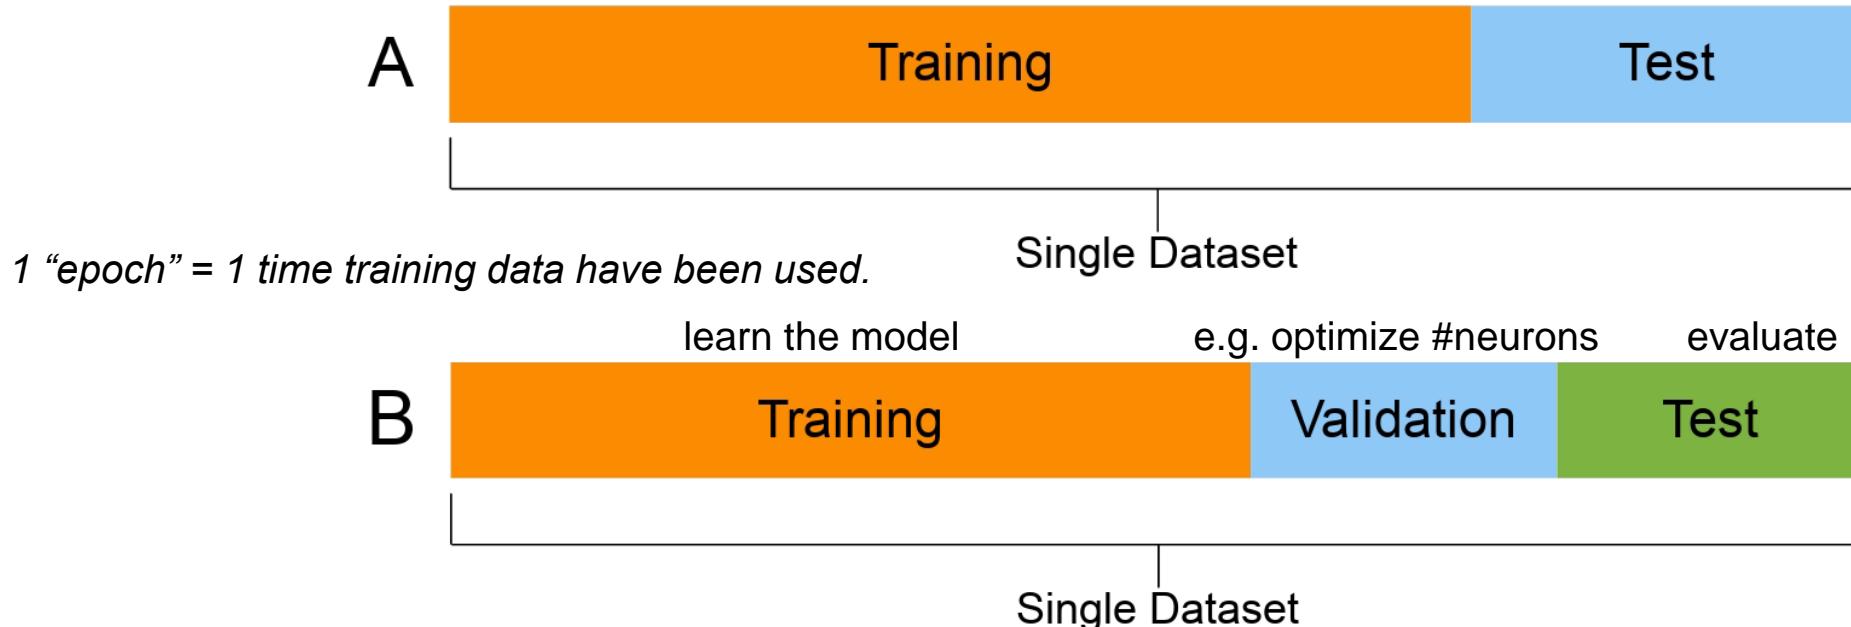
- Evaluate (Did it really learn anything?)

Train – Test – Validation sets

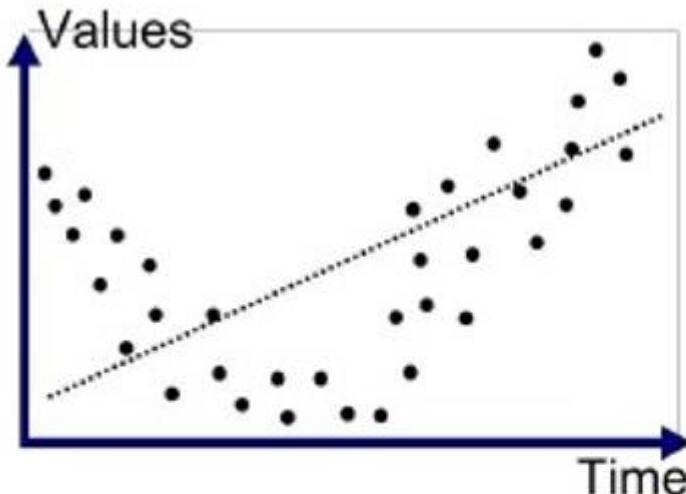
3	4	2	1	9	5	6	2	1	8
8	9	1	2	5	0	0	6	6	4
6	7	0	1	6	3	6	3	7	0
3	7	7	9	4	6	6	1	8	2
2	9	3	4	3	9	8	7	2	5
1	5	9	8	3	6	5	7	2	3
9	3	1	9	1	5	8	0	8	4
5	6	2	6	8	5	8	8	9	9
3	7	7	0	9	4	8	5	4	3
7	9	6	4	7	0	6	9	2	3

Question

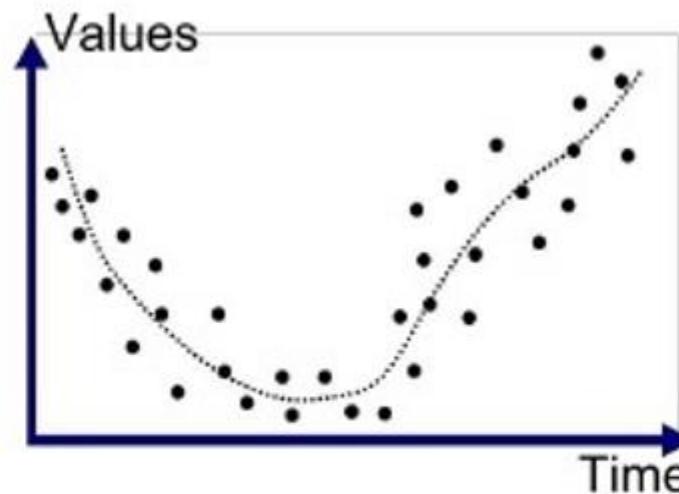
How do we evaluate the model now ??



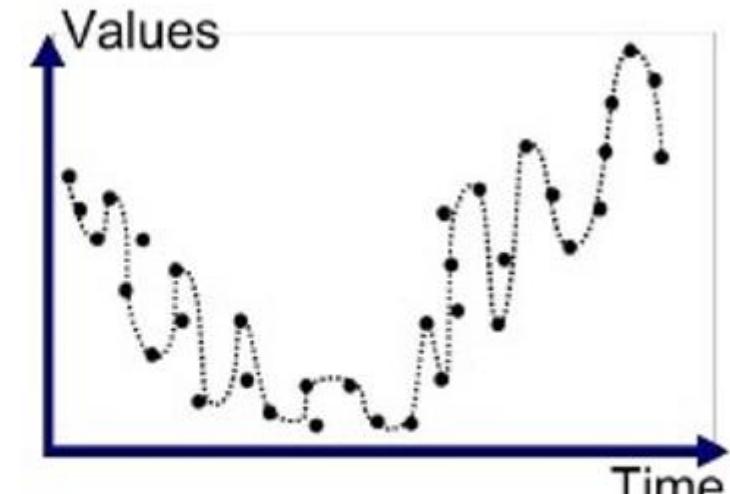
Overfitting



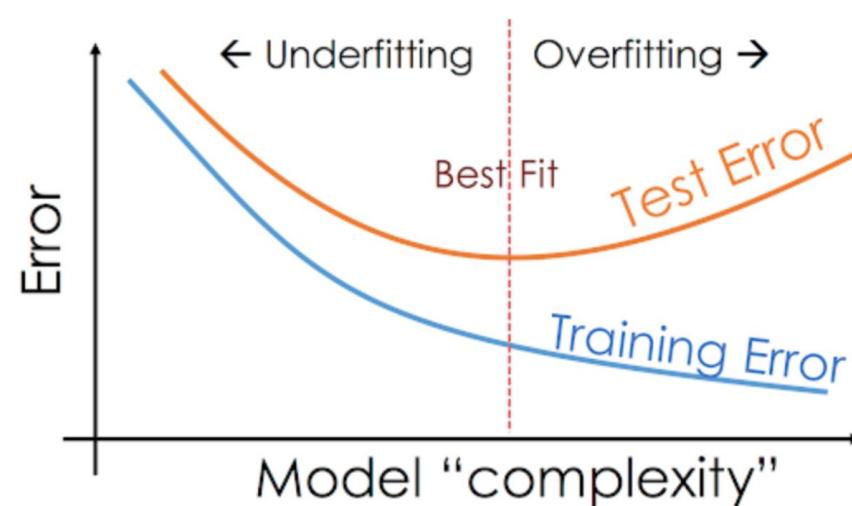
Underfitted



Good Fit/Robust

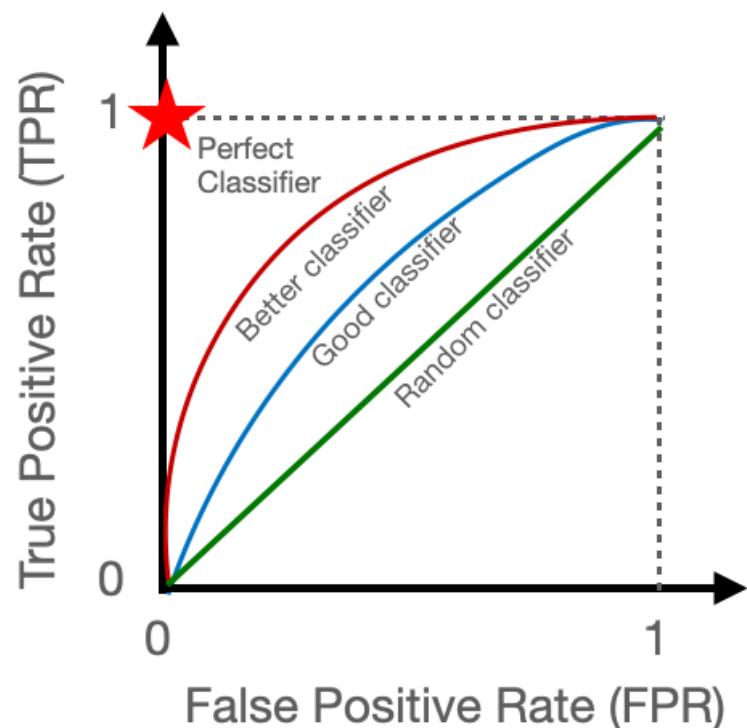


Overfitted



Evaluation Metrics (Classification)

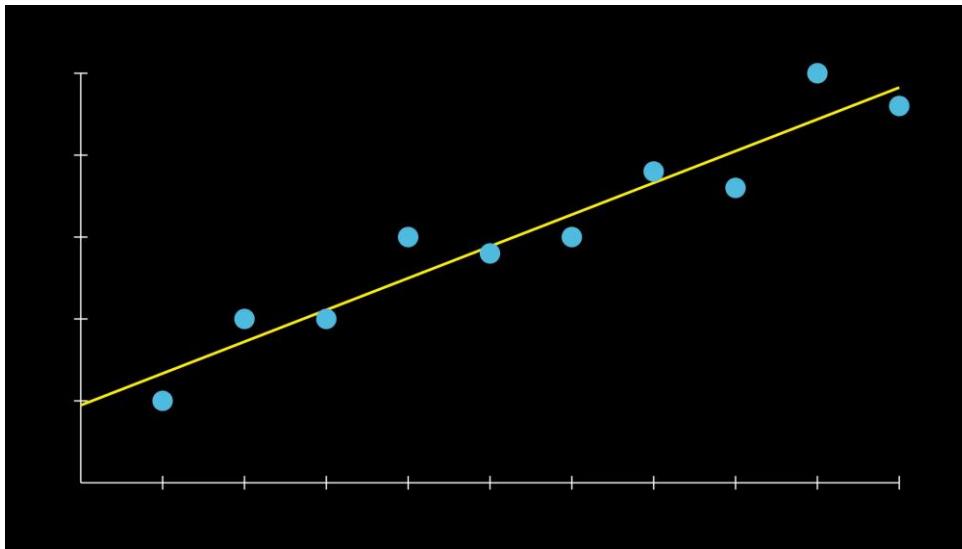
	Predicted 0	Predicted 1
Actual 0	TN	FP
Actual 1	FN	TP



Metric	Formula
True positive rate, recall	$\frac{TP}{TP+FN}$
False positive rate	$\frac{FP}{FP+TN}$
Precision	$\frac{TP}{TP+FP}$
Accuracy	$\frac{TP+TN}{TP+TN+FP+FN}$
F-measure	$\frac{2 \cdot \text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$

Evaluation Metrics (Regression)

$$MSE = \frac{1}{n} \sum \left(\underbrace{y - \hat{y}}_{\text{The square of the difference between actual and predicted}} \right)^2$$



Mean squared error

$$MSE = \frac{1}{n} \sum_{t=1}^n e_t^2$$

Root mean squared error

$$RMSE = \sqrt{\frac{1}{n} \sum_{t=1}^n e_t^2}$$

Mean absolute error

$$MAE = \frac{1}{n} \sum_{t=1}^n |e_t|$$

Mean absolute percentage error

$$MAPE = \frac{100\%}{n} \sum_{t=1}^n \left| \frac{e_t}{y_t} \right|$$

Outline

How to understand Neural Network modeling:

- ~~Concept (basic principles) & Building (what are the components?)~~

Answer: Yes! We know that now! We add a lot of stuff, to a randomly parametrized model.....
But ok, then what ? It should just spit out random result now right ?

- ~~Train (How does it learn?)~~

Answer: Ok Ok, we got it. Now we calculate the derivative and by comparing the initially random output to the one we wanted, we change the parameters of the network. Sure... But ok, how do we test that thing ? I mean... It learned something that was taught. But everyone can do that... is it smart enough to deal with brand new information ?

- Evaluate (Did it really learn anything?)

Answer: ...

Outline

How to understand Neural Network modeling:

- ~~Concept (basic principles) & Building (what are the components?)~~

Answer: Yes! We know that now! We add a lot of stuff, to a randomly parametrized model.....
But ok, then what ? It should just spit out random result now right ?

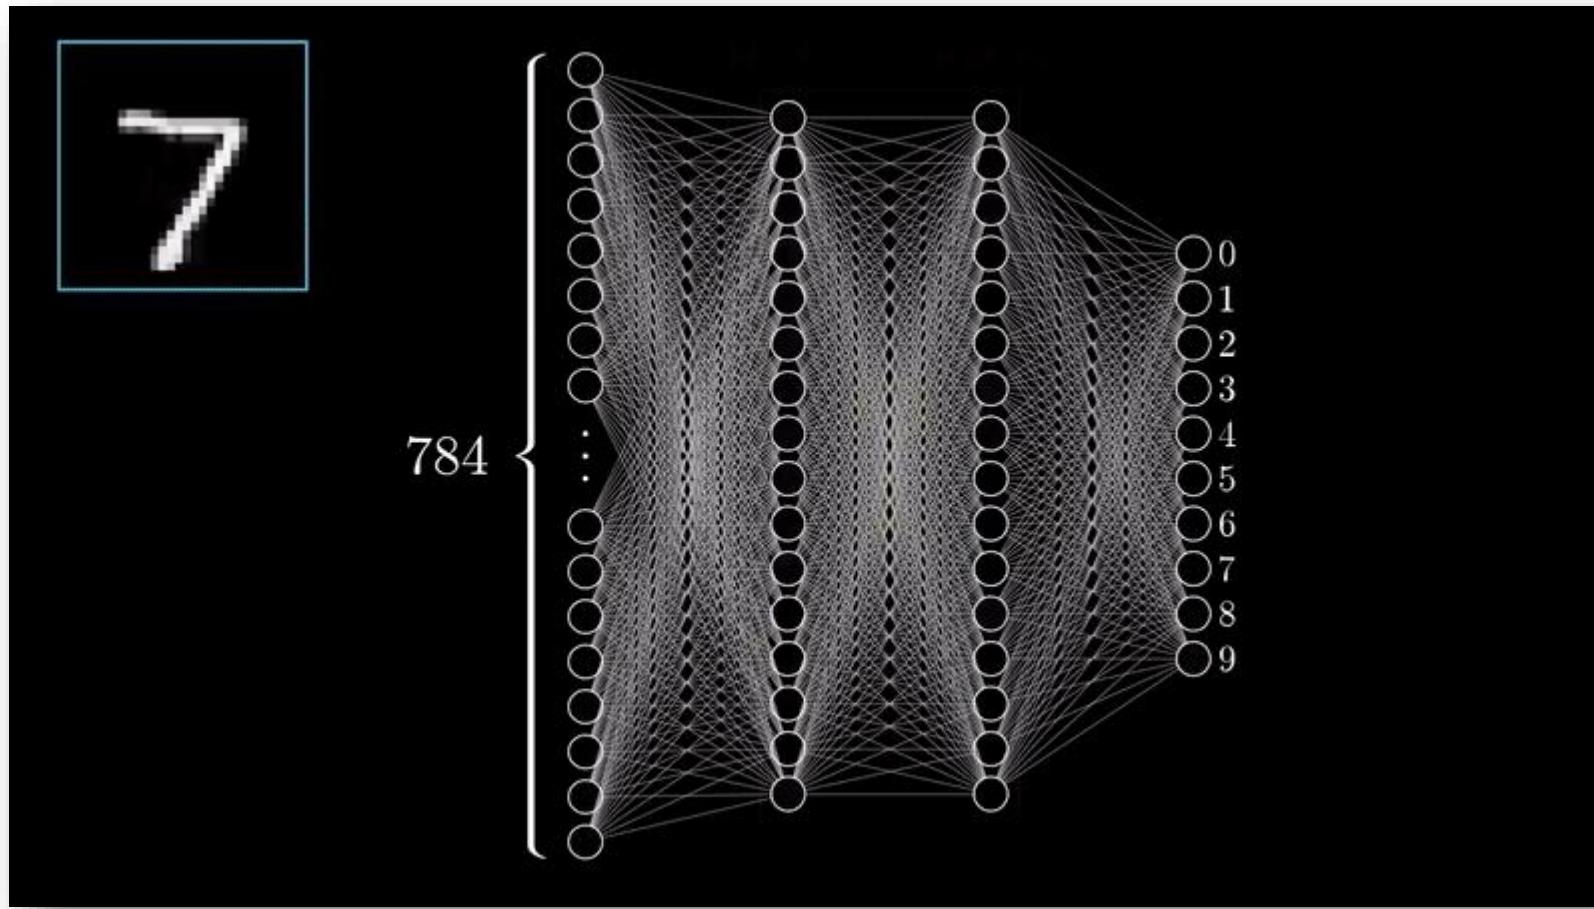
- ~~Train (How does it learn?)~~

Answer: Ok Ok, we got it. Now we calculate the derivative and by comparing the initially random output to the one we wanted, we change the parameters of the network. Sure... But ok, how do we test that thing ? I mean... It learned something that was taught. But everyone can do that... is it smart enough to deal with brand new information ?

- Evaluate (Did it really learn anything?)

Answer: Ok, so we put data, we change the parameters based on these and then we test with different data to see if the model is smart enough to pick the correct answer. Cool! Then we use different metrics to quantify what happened.

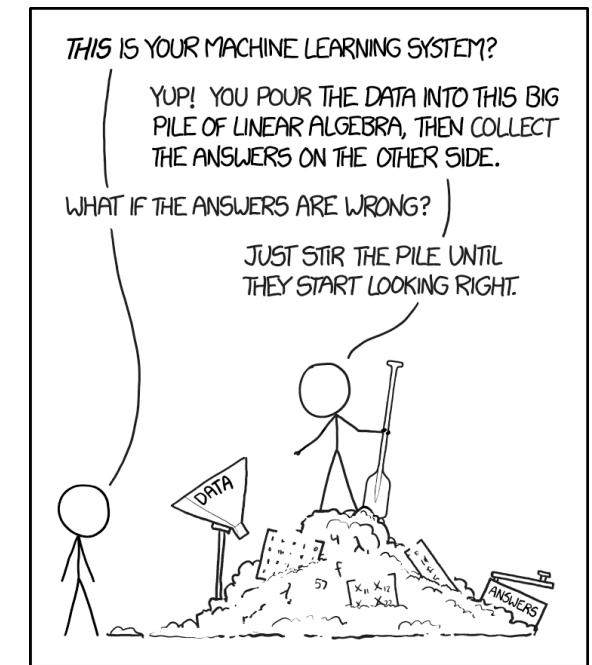
A Trained Neural Network



*Video Courtesy: **3Blue1Brown** (Check him on YouTube!)

Play time

<http://playground.tensorflow.org>

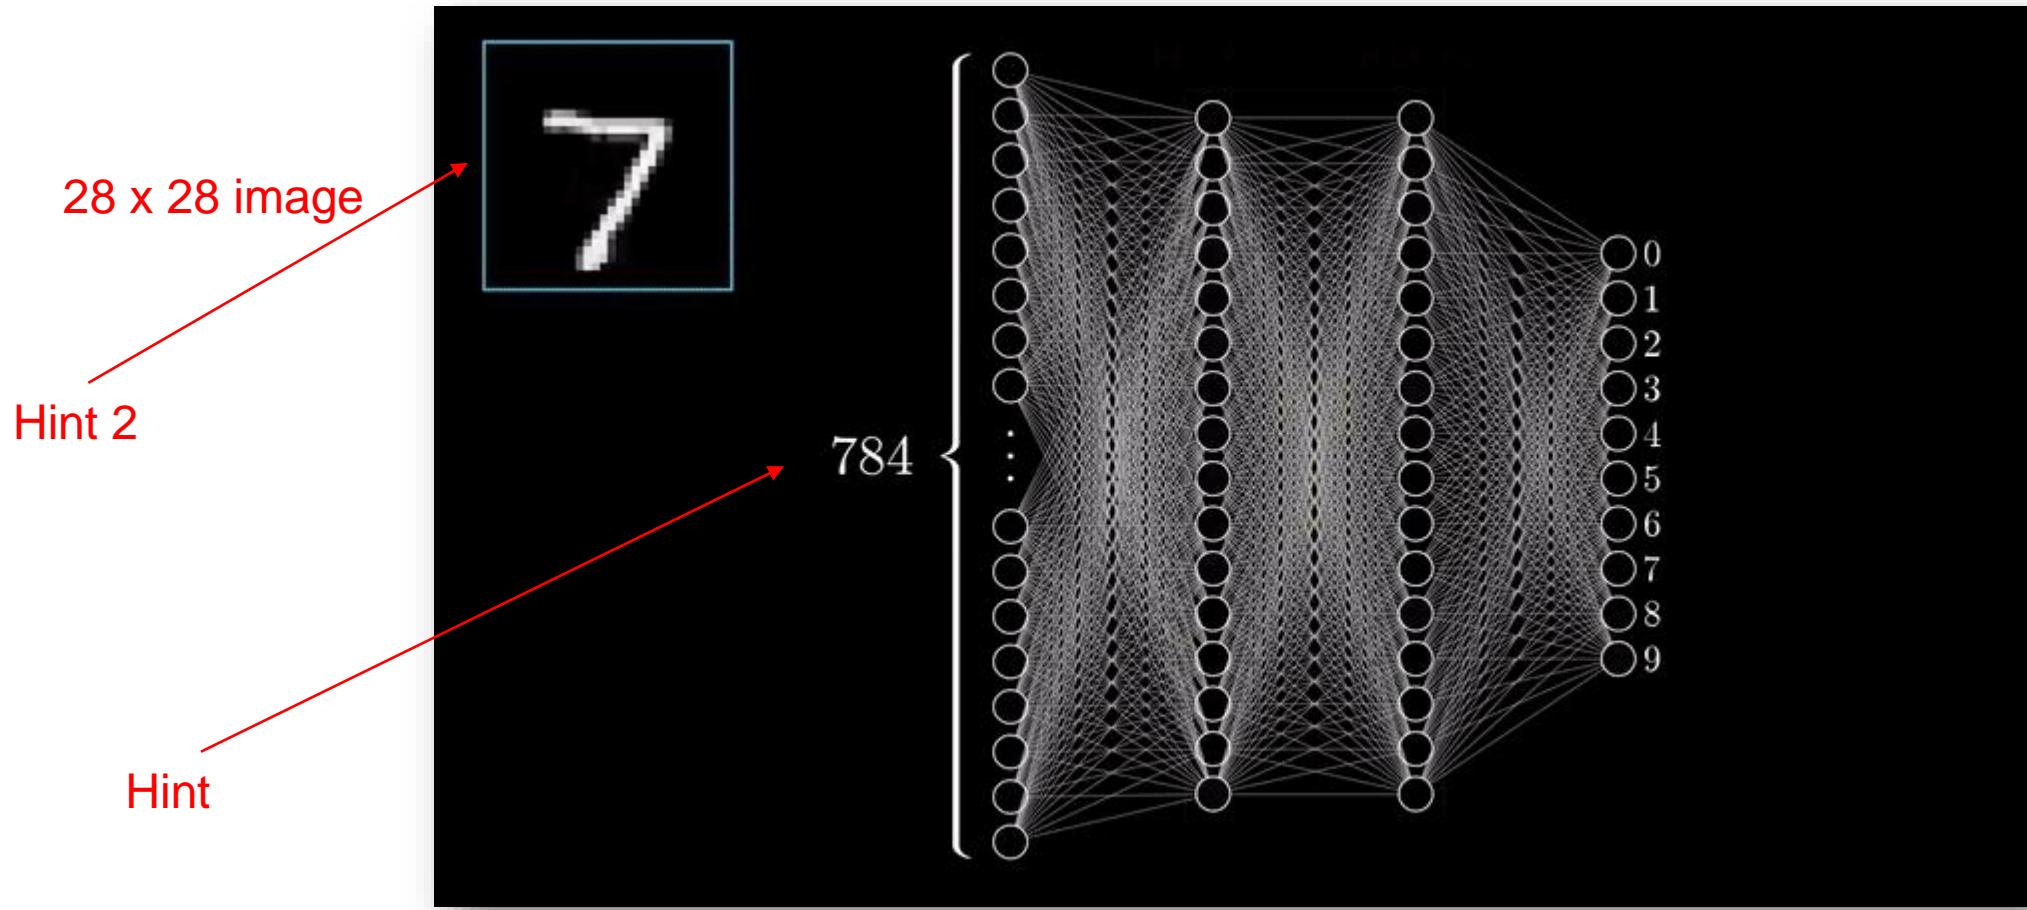


We can try to “stir the pile” with:

- Extra input
- Extra neurons
- Extra layers

But wait, there is more !

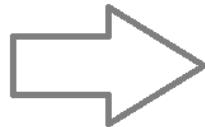
See anything weird ?



*Video Courtesy: **3Blue1Brown** (Check him on YouTube!)

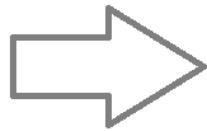
Neural Networks with Images

1	1	0
4	2	1
0	2	1



1
1
0
4
2
1
0
2
1

Neural Networks with Images – Dog example



“slightly” exaggerated

Convolutional Neural Networks

Convolutional Neural Network (CNN) Layers

Convolution
Extract features & Keep spatial relationship

1 <small>x1</small>	1 <small>x0</small>	1 <small>x1</small>	0	0
0 <small>x0</small>	1 <small>x1</small>	1 <small>x0</small>	1	0
0 <small>x1</small>	0 <small>x0</small>	1 <small>x1</small>	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved Feature

Pooling/Subsampling
Reduce dimensionality & retain information

12	20	30	0
8	12	2	0
34	70	37	4
112	100	25	12

2×2 Max-Pool

20	30
112	37

Convolution and max pooling layer on image

Convolution

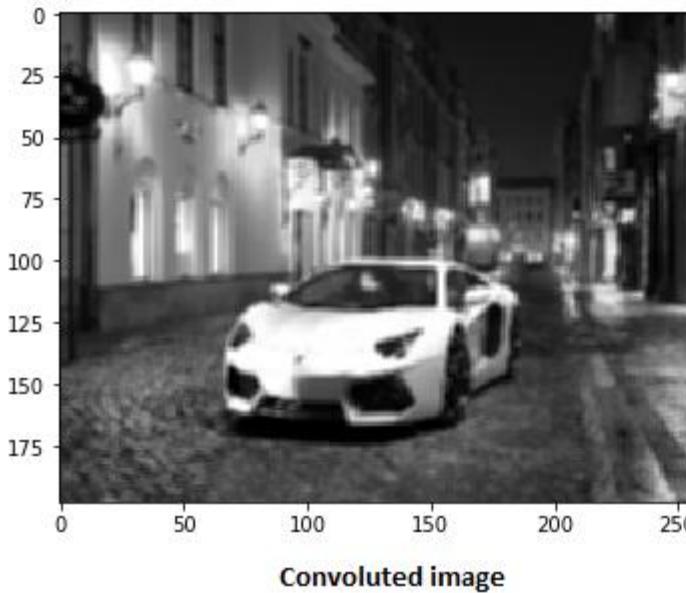


*

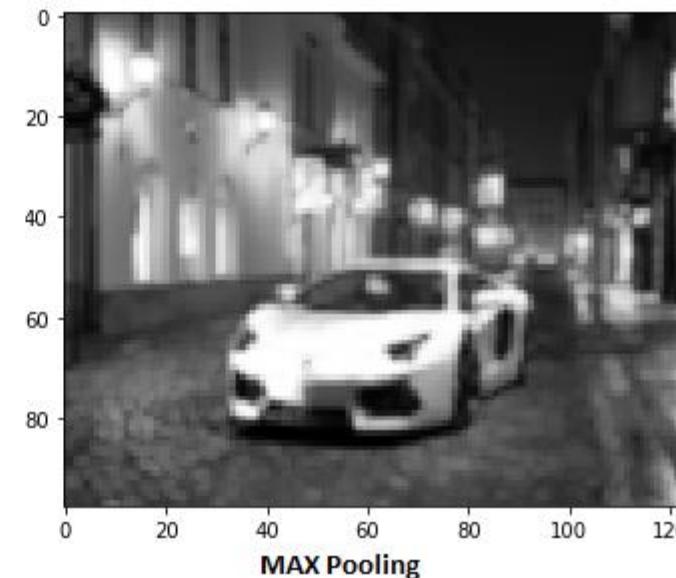
1	0	-1
2	0	-2
1	0	-1



Max Pooling

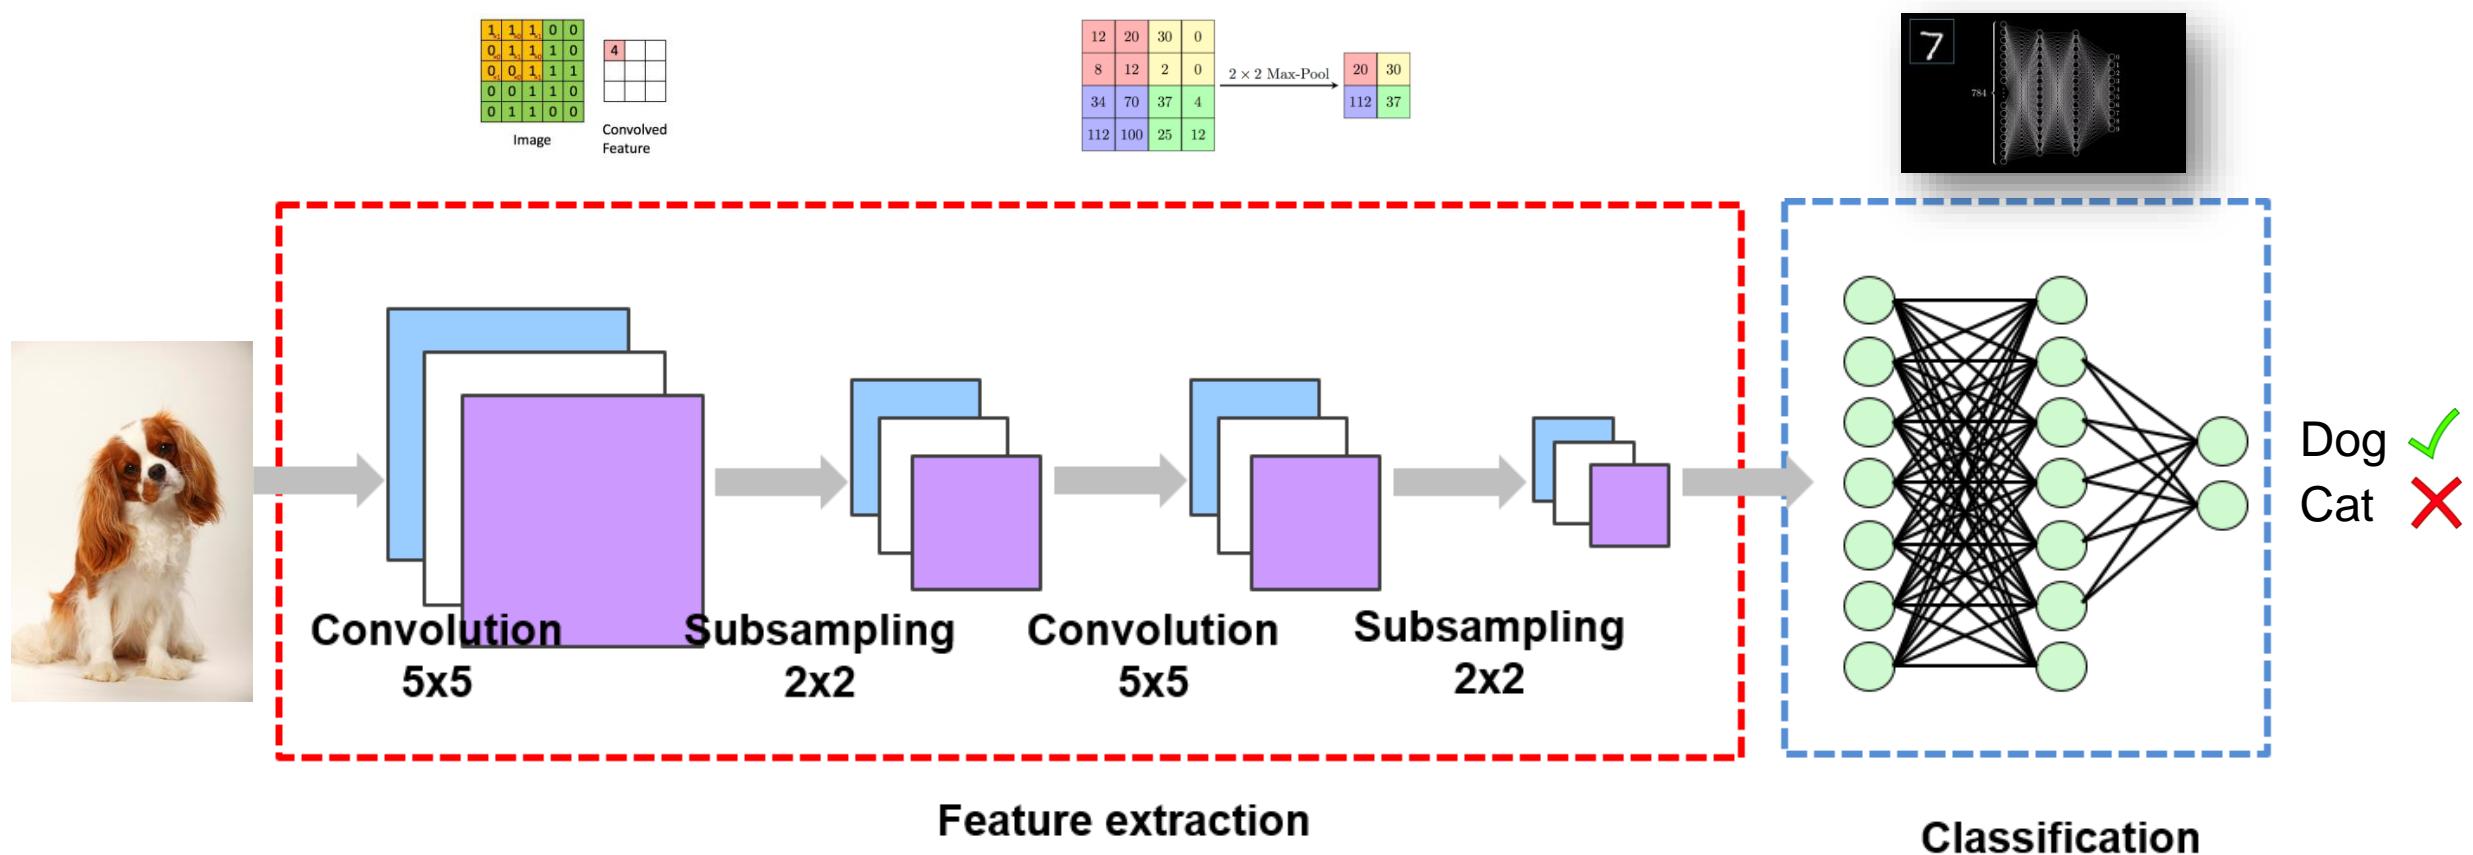


Convoluted image



MAX Pooling

Example of CNN

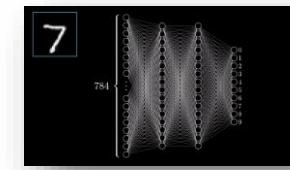


*Figure Courtesy: Suhyun Kim iSystems Design Labs

NN vs CNN

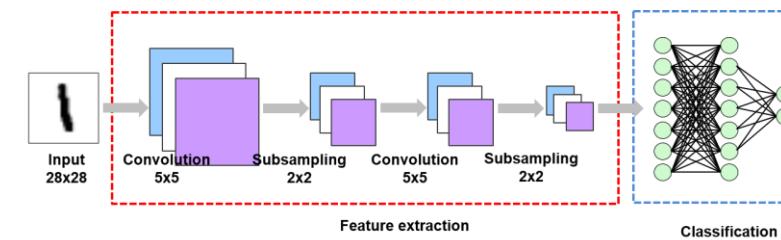
Input: MNIST database

0 0 0 0 0 0 0 0 0 0 0 0
1 1 1 1 1 1 1 1 1 1 1 1
2 2 2 2 2 2 2 2 2 2 2 2
3 3 3 3 3 3 3 3 3 3 3 3
4 4 4 4 4 4 4 4 4 4 4 4
5 5 5 5 5 5 5 5 5 5 5 5
6 6 6 6 6 6 6 6 6 6 6 6
7 7 7 7 7 7 7 7 7 7 7 7
8 8 8 8 8 8 8 8 8 8 8 8
9 9 9 9 9 9 9 9 9 9 9 9



Neural Network Result:

97.8%



Convolution Neural Network Result:

99.07%

Hands on examples of Supervised Learning

Hands on examples with MNIST

- Neural Network (MNIST)
- Convolution Neural Network (MNIST)

Follow this link & install yourself : <https://github.com/SavvasRaptis/machine-learning-examples>

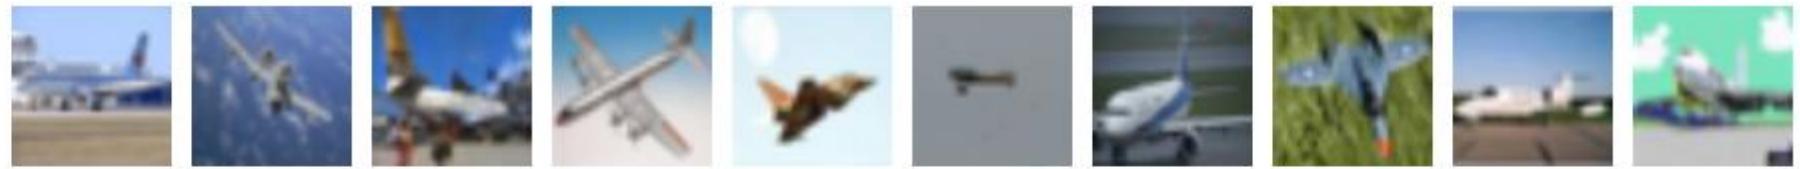
Or directly use: <https://mybinder.org/v2/gh/SavvasRaptis/machine-learning-examples/HEAD>





If MNIST is easy for you, check CIFAR-10

airplane



automobile



bird



cat



deer



dog



Can get harder!



@teenybiscuit

<https://www.kaggle.com/datasets/returnofspputnik/chihuahua-or-muffin>

Supervised Learning in Space Physics

EXAMPLE

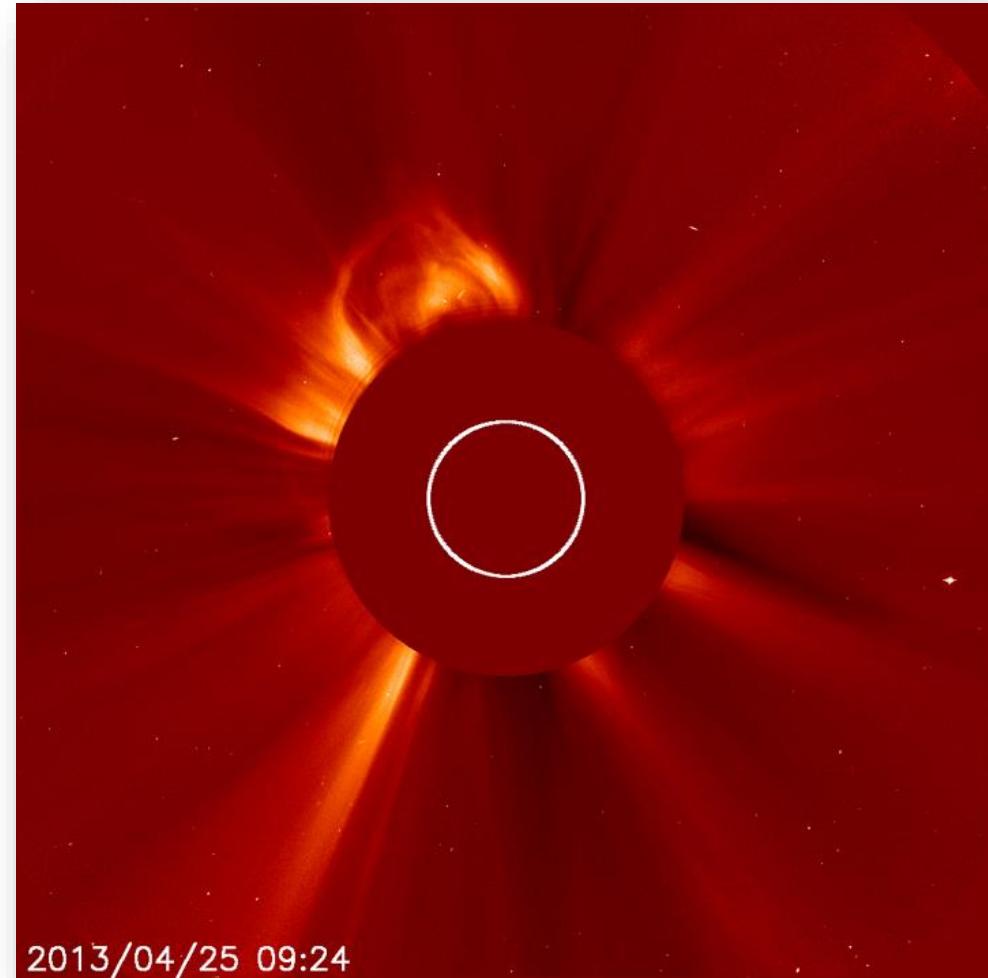
CME prediction example (Classification problem)

Coronal Mass Ejections (CMEs)

CMEs
Energetic particles from the Sun

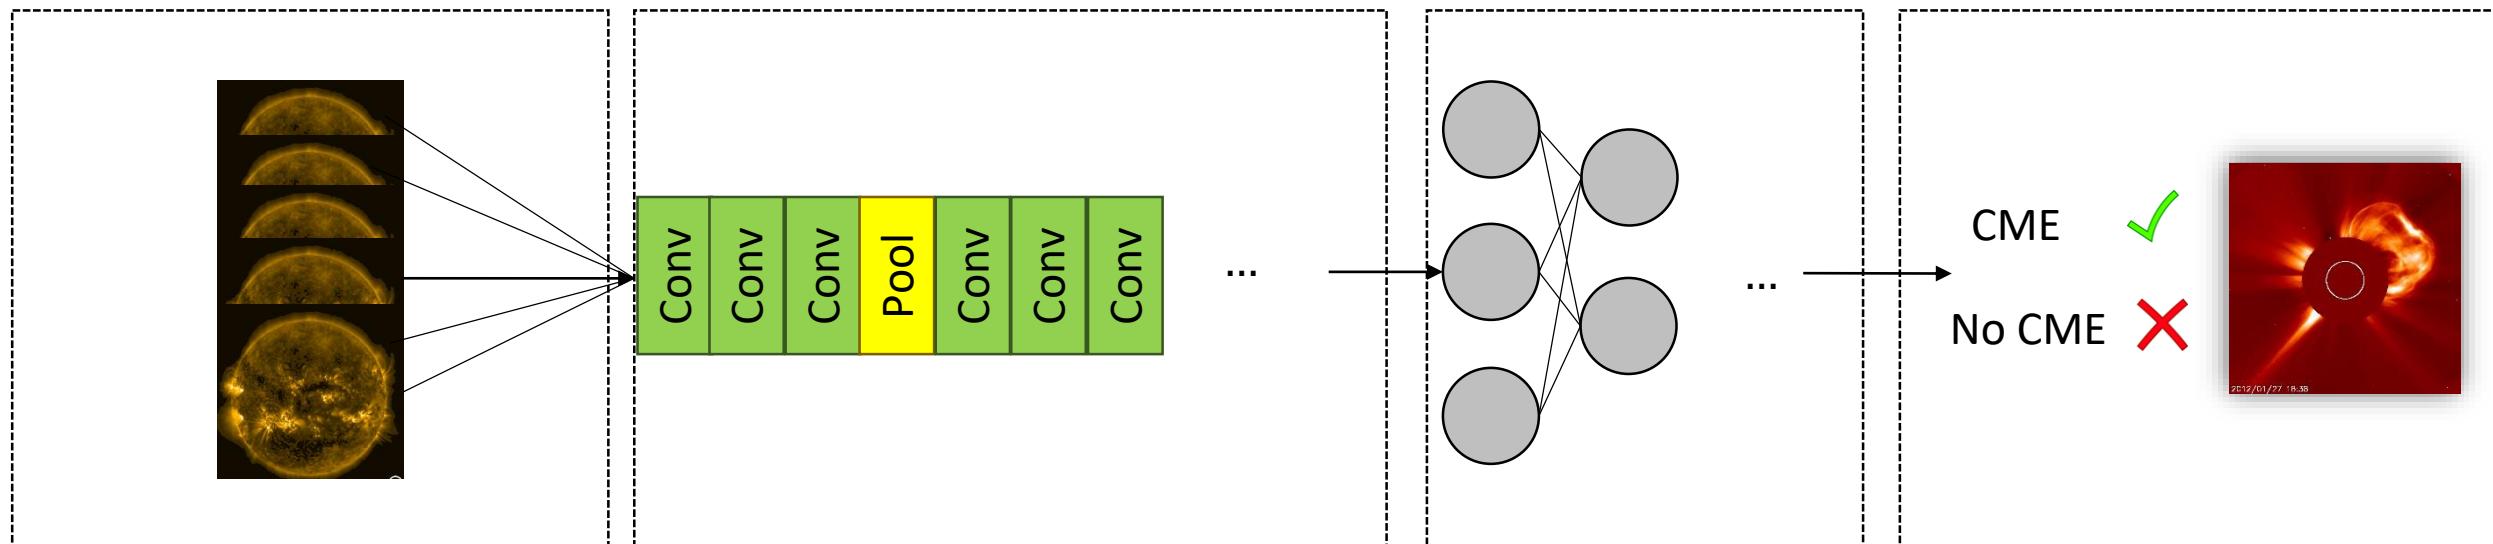


Problems in
Satellites/Communications/Grids/infrastructures etc.



Project description

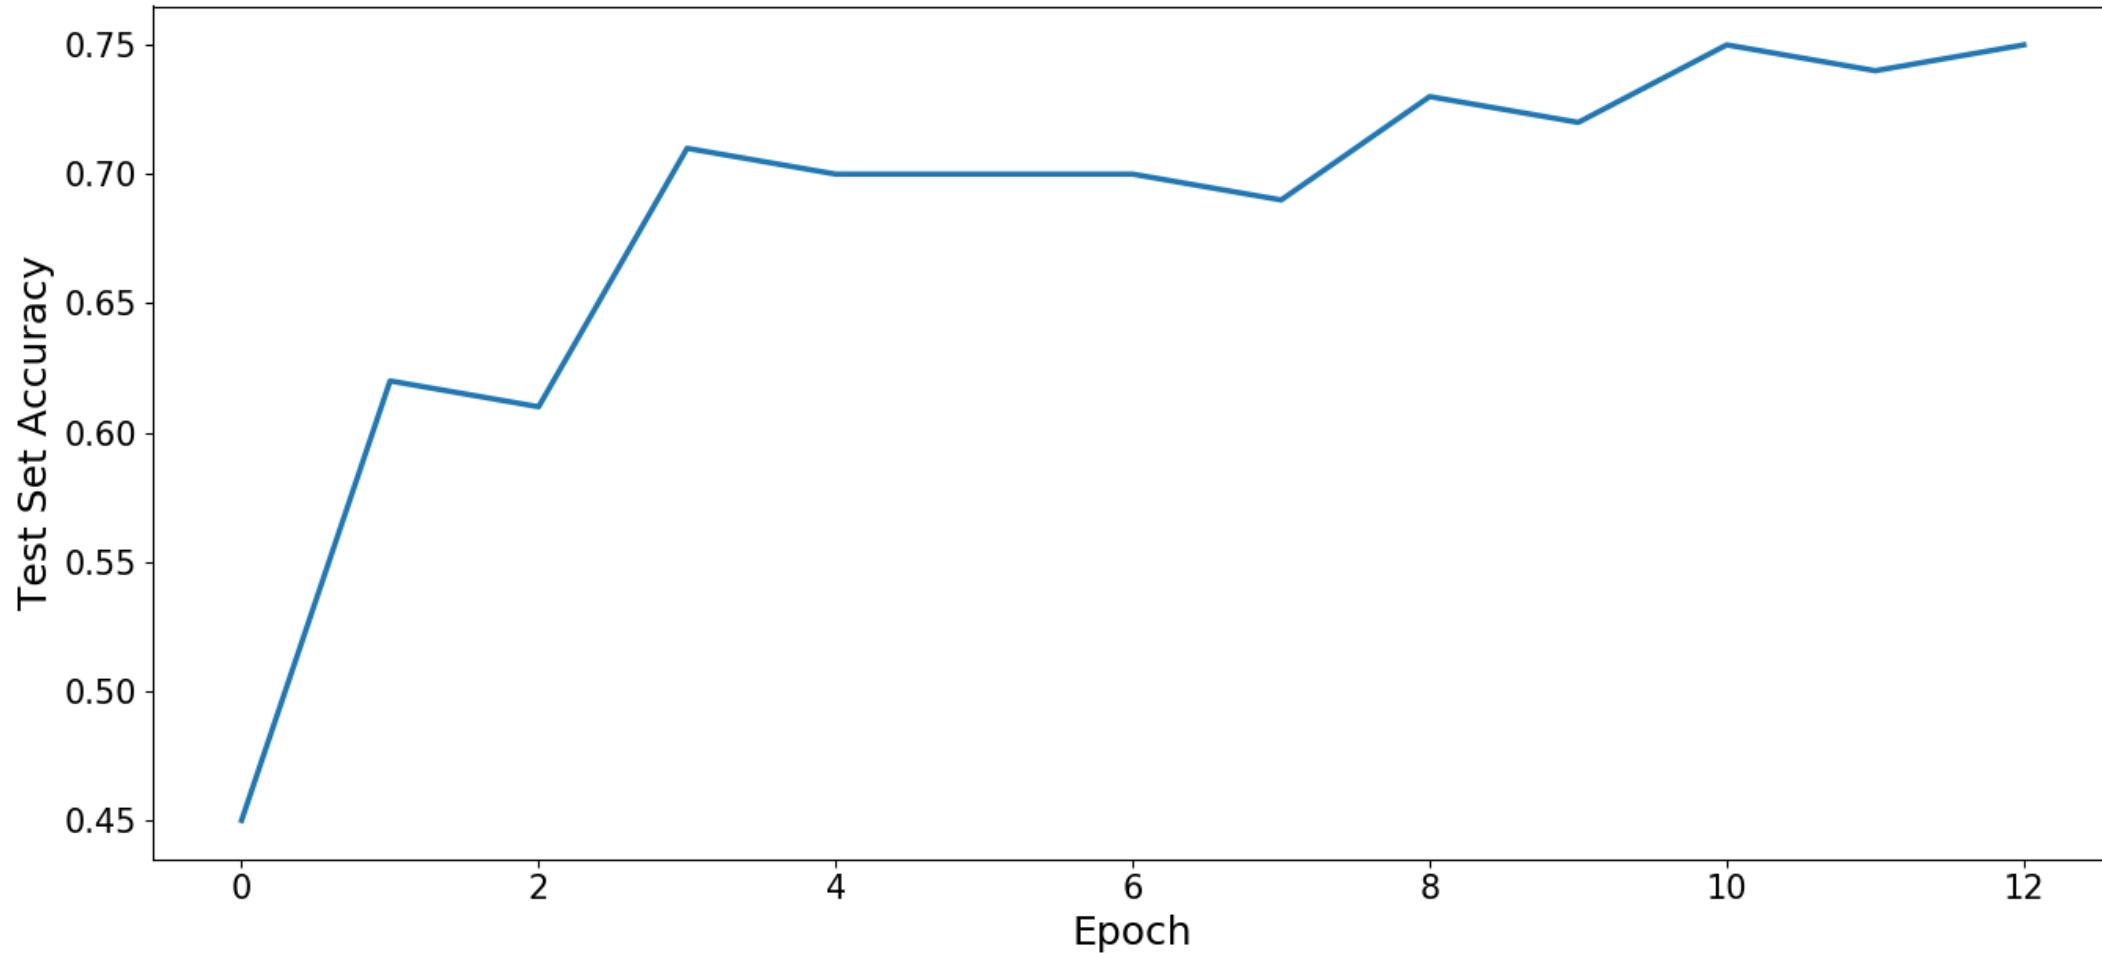
Input [13,512,512] Convolution Neural Network Dense Neural Network Output
Feature Extraction Classification CME / No CME



Input = 13 SDO images, 2 [h] history before the event.

Output = 1/0

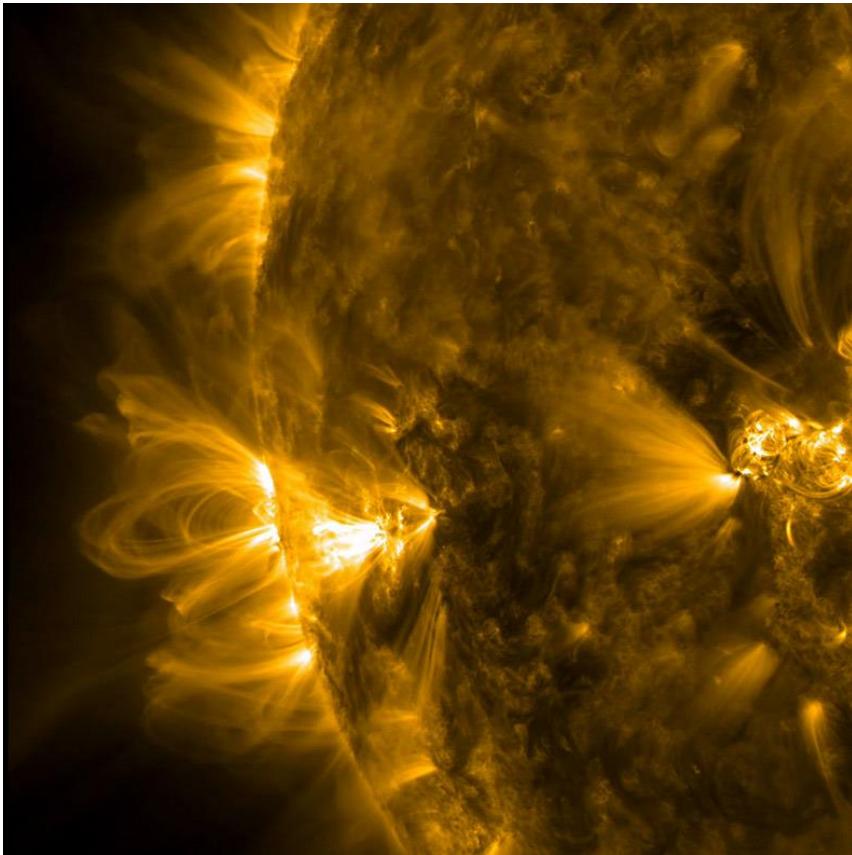
Result of CNN



EXAMPLE

Coronal Loop reconstruction example (Regression problem)

Project description



Basic Properties

- More easily seen in EUV & X-ray
- More common in solar max
- May be relevant to the coronal heating problem

Information

- We see projections, therefore **2D can be observationally derived** but **3D needs to be computed.**
- **3D = vital for magnetic field extrapolation**
- There are many methods, but have limitations. [Aschwanden (2011), Chifu et al. (2017)]

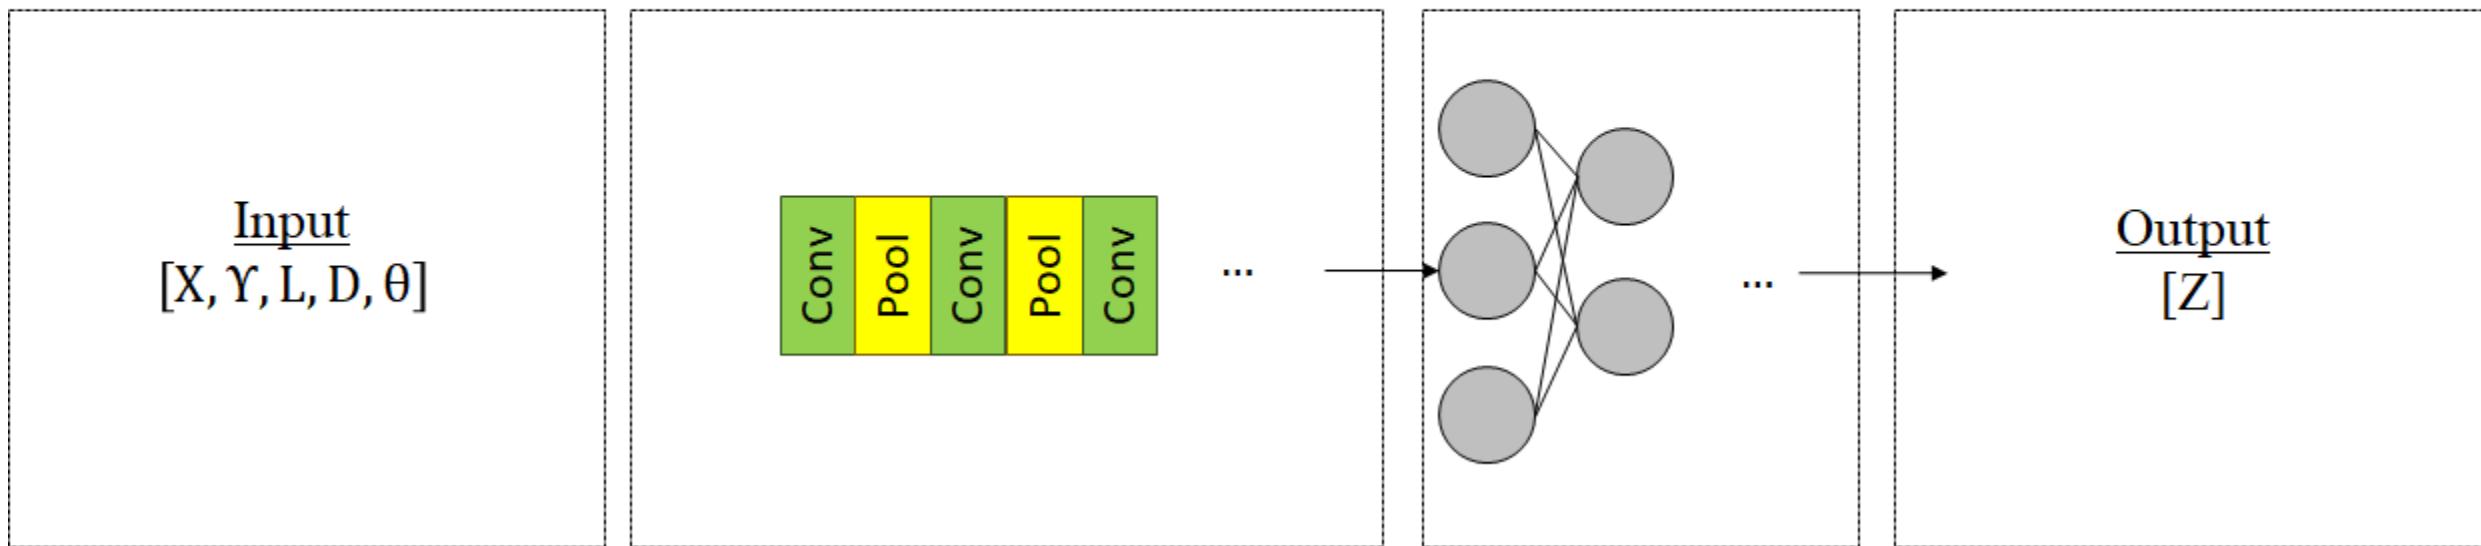
Neural Network Architecture

Convolution Neural Network

Feature Extraction

Dense Neural Network

Regression



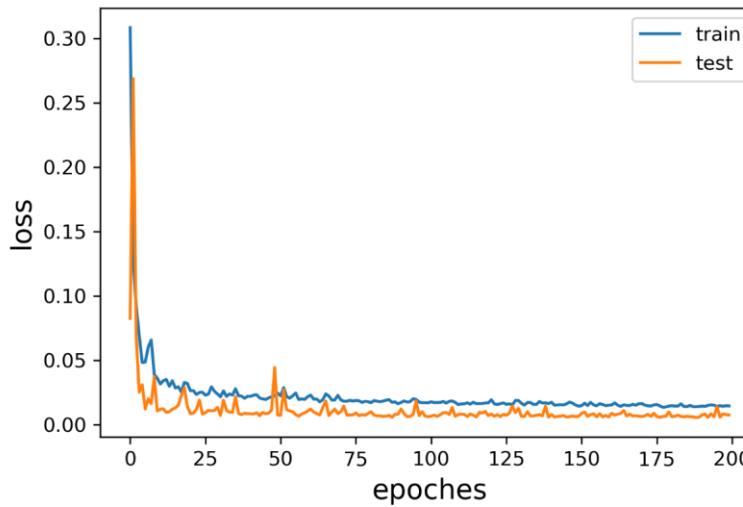
X, Y, Z : Coordinates (1500 points each)

L : Length of the coronal Loop

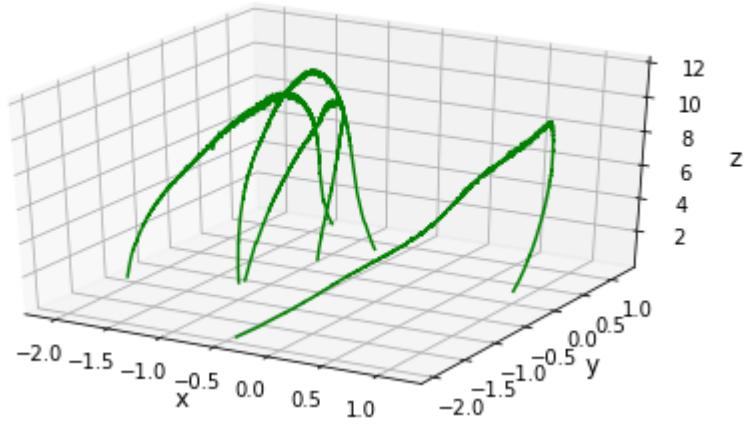
D : Distance between footpoints

Θ : Angle defined by the two footpoints with the top of the loops

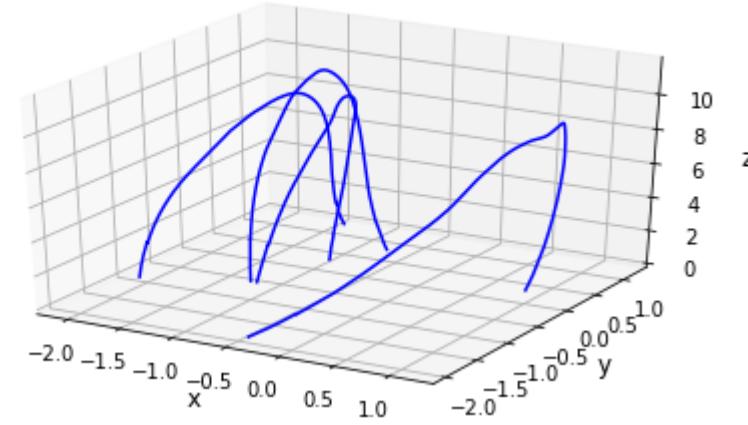
Results of the reconstruction



Predicted Coronal Loops



Actual Coronal Loops

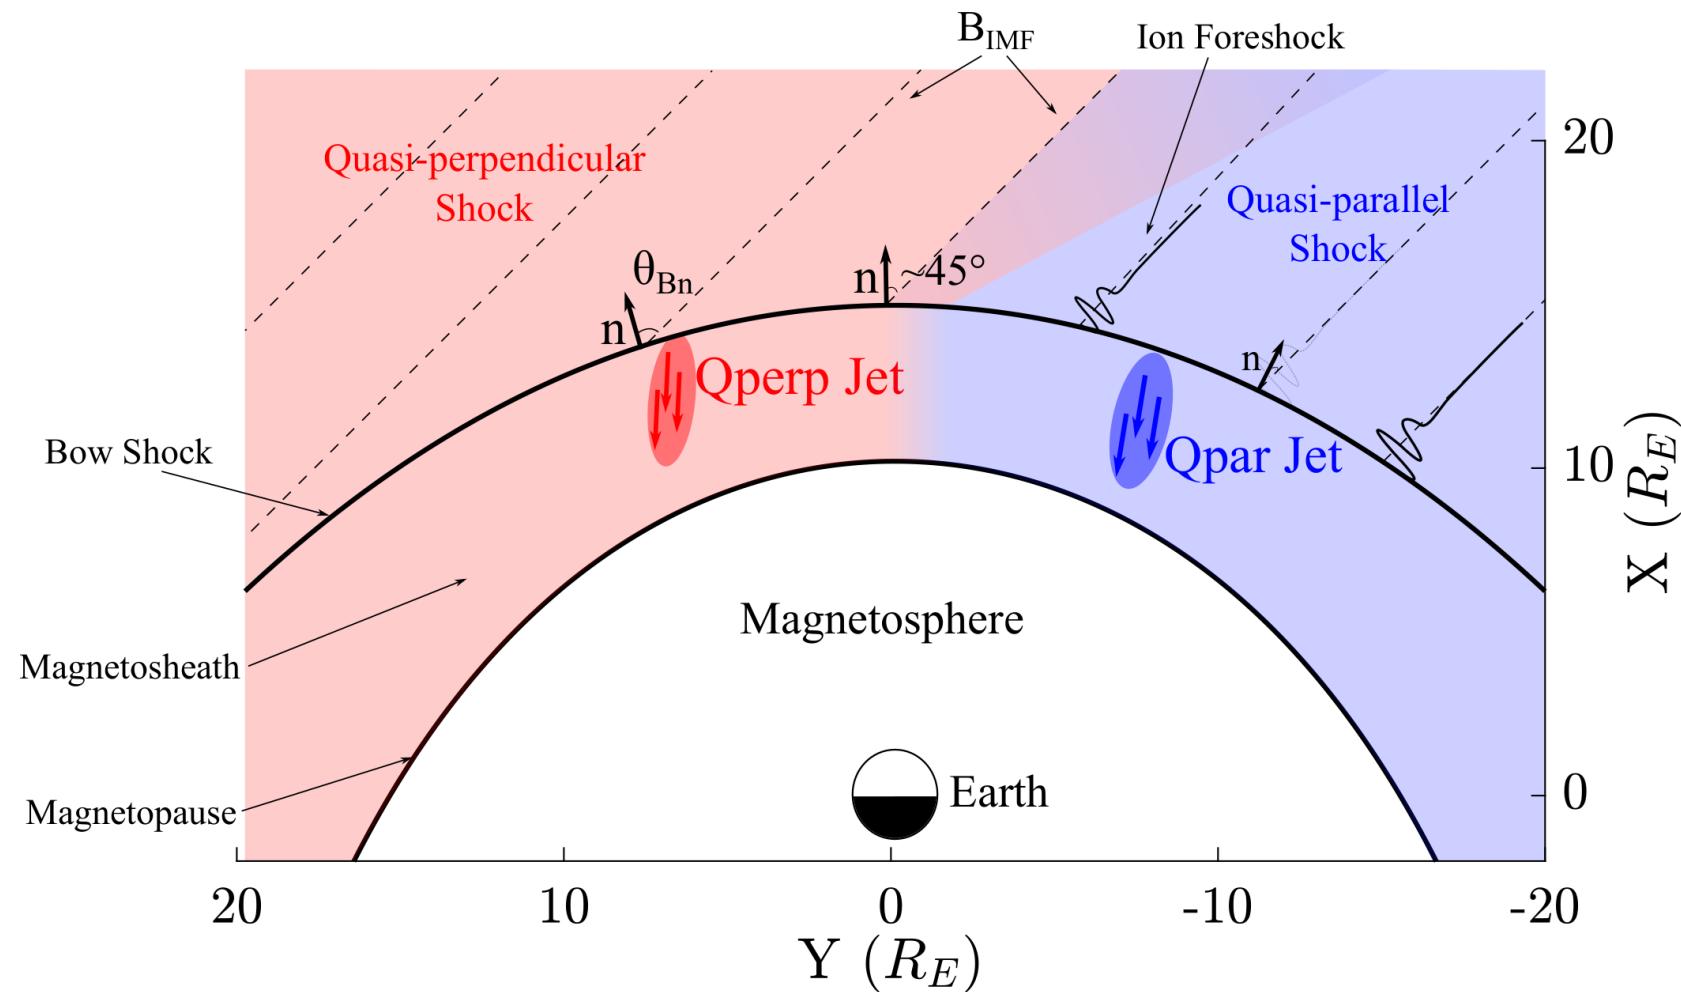


See original results (Chifu & Gafeira 2021)

EXAMPLE

Magnetosheath jet classification (Classification problem)

Magnetosheath jet



Definition

Magnetosheath jets are **transient localized enhancements** of **dynamic pressure** (density and/or velocity increase)

Classification

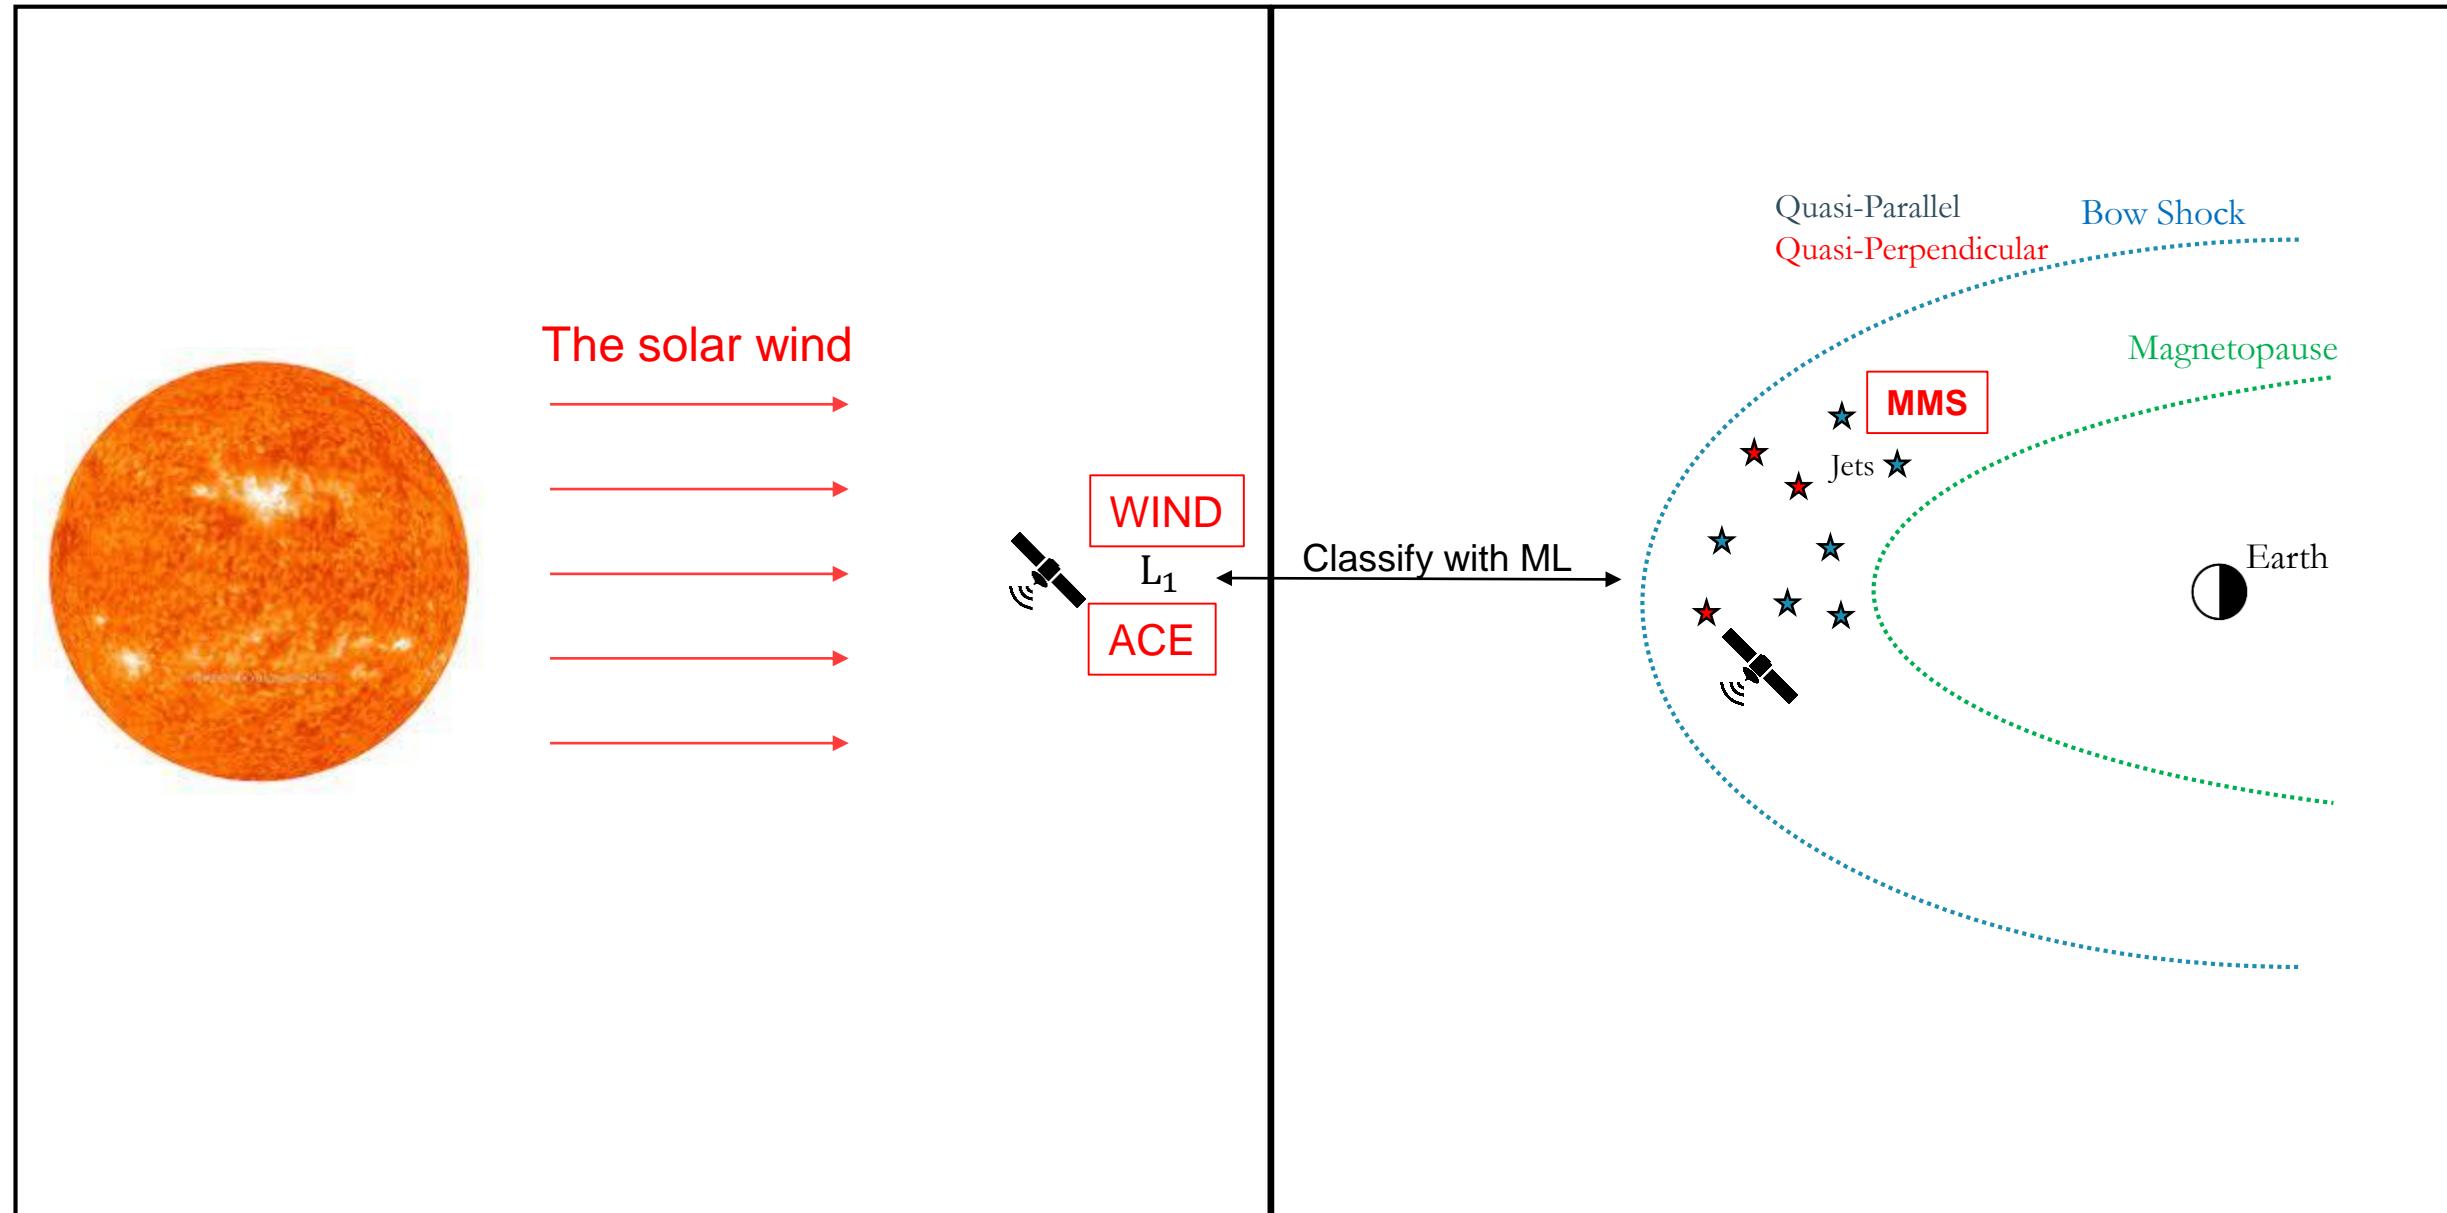
“ Θ_{Bn} is the angle between the IMF and the shock’s normal vector”

$$Qpar = \Theta_{Bn} \lesssim 45^\circ$$

$$Qperp = \Theta_{Bn} \gtrsim 45^\circ$$

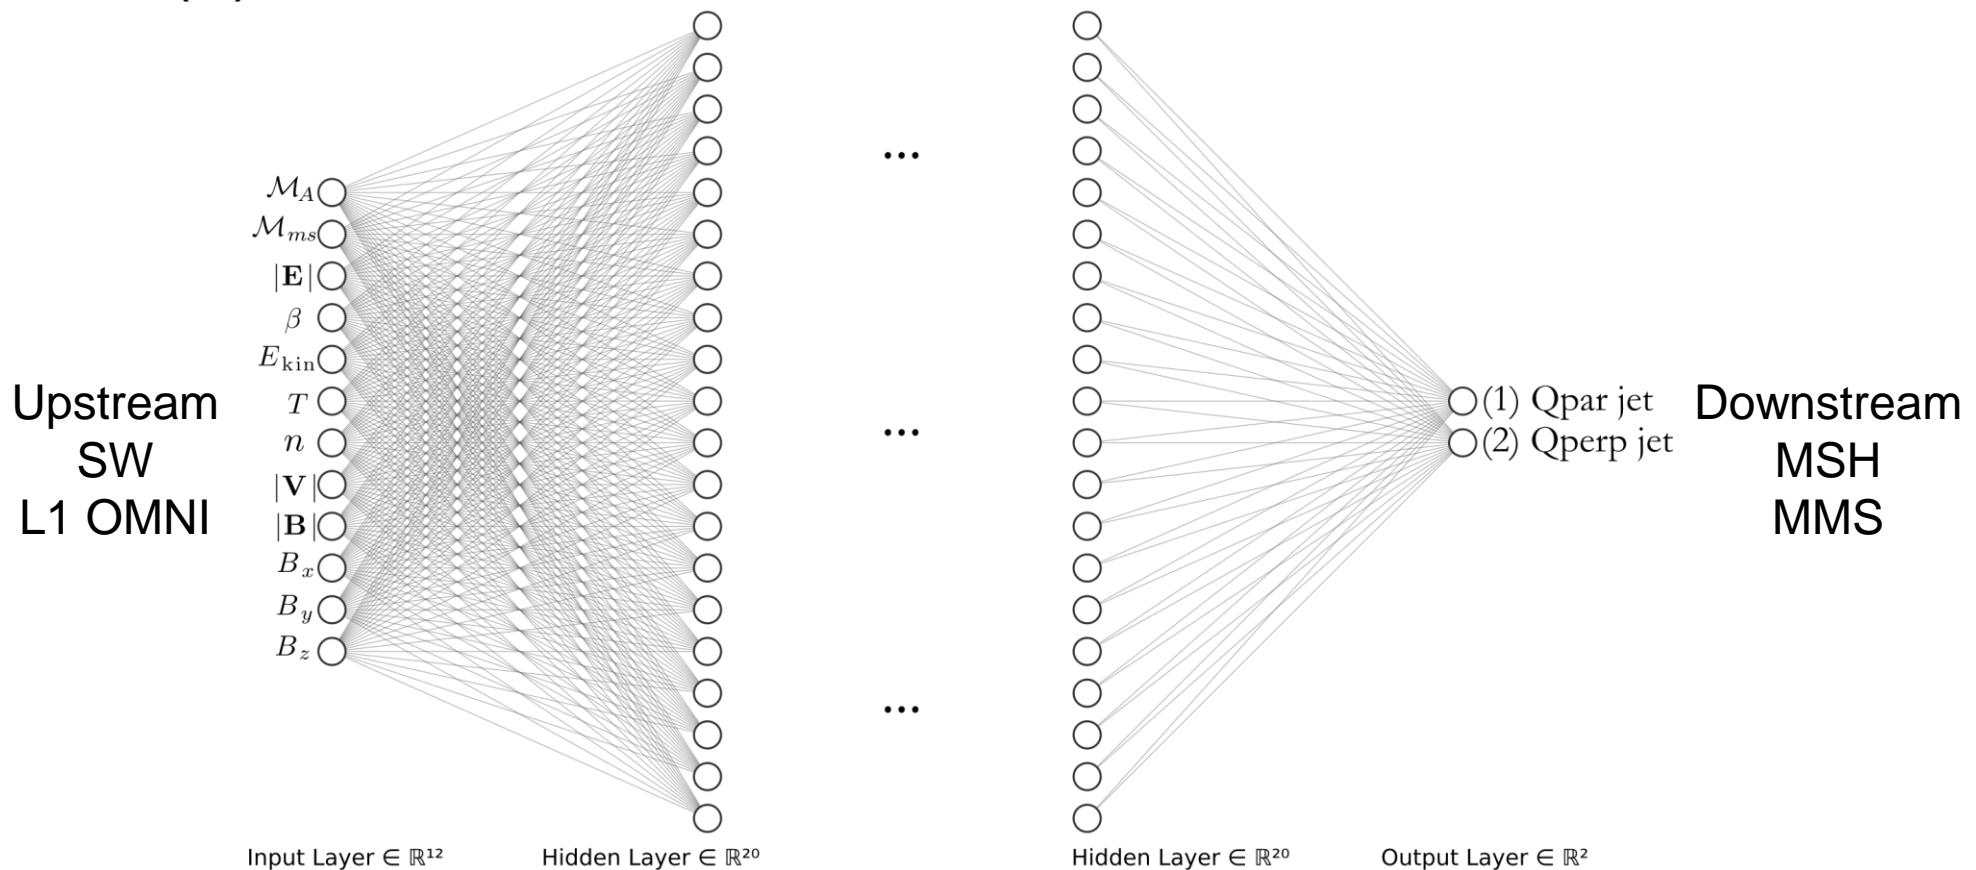
“~10 times more often in Qpar MSH”

Classification task



Classifying jets with Neural Networks

(a)



(b)

Results

Class	Neural network (B) (%)	Neural network (No - B) (%)	θ_{cone} (%)	Coplanarity (%)	Modeling (%)
Qpar	98	95	61	81	74
Qperp	88	87	94	79	86



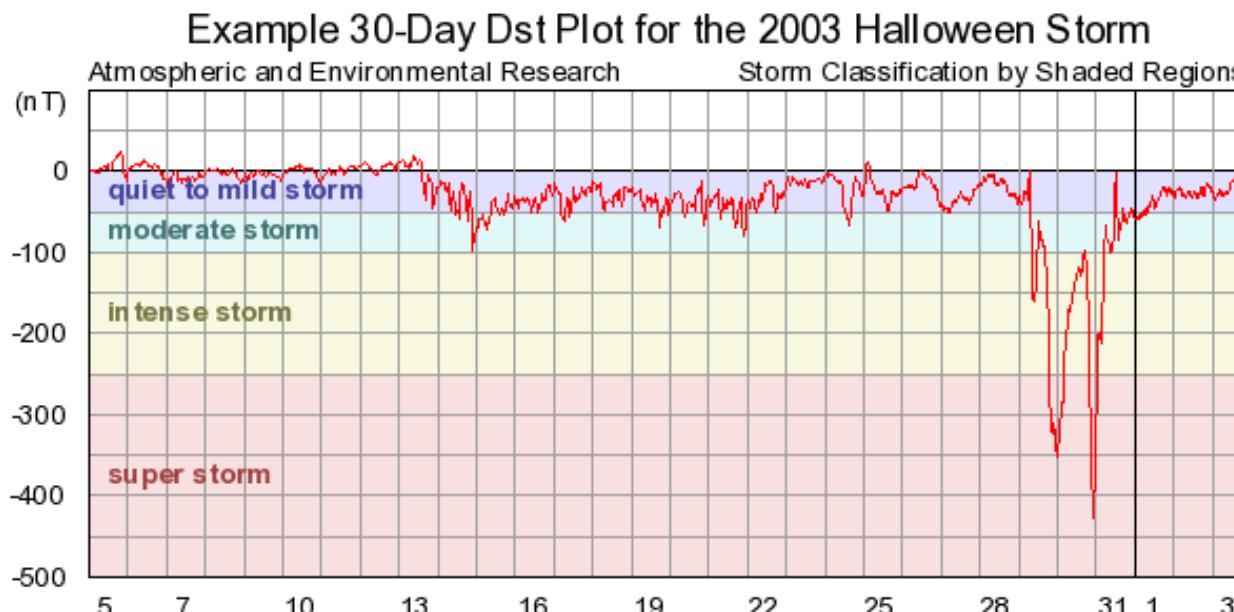
**HANDS-ON
LEARNING**

Forecasting of DST index (Regression problem)

Prediction of the Dst index

“average disturbance field at the Earth's equator from four low-latitude magnetic observatories”

The value of Dst along with its pattern (sharpness, duration, etc.) provide a profile for the **geomagnetic storm**



Recap questions

Why we call this a regression problem ??

What would a classification problem be here ?



<https://github.com/SavvasRaptis/machine-learning-examples>

Try for example:

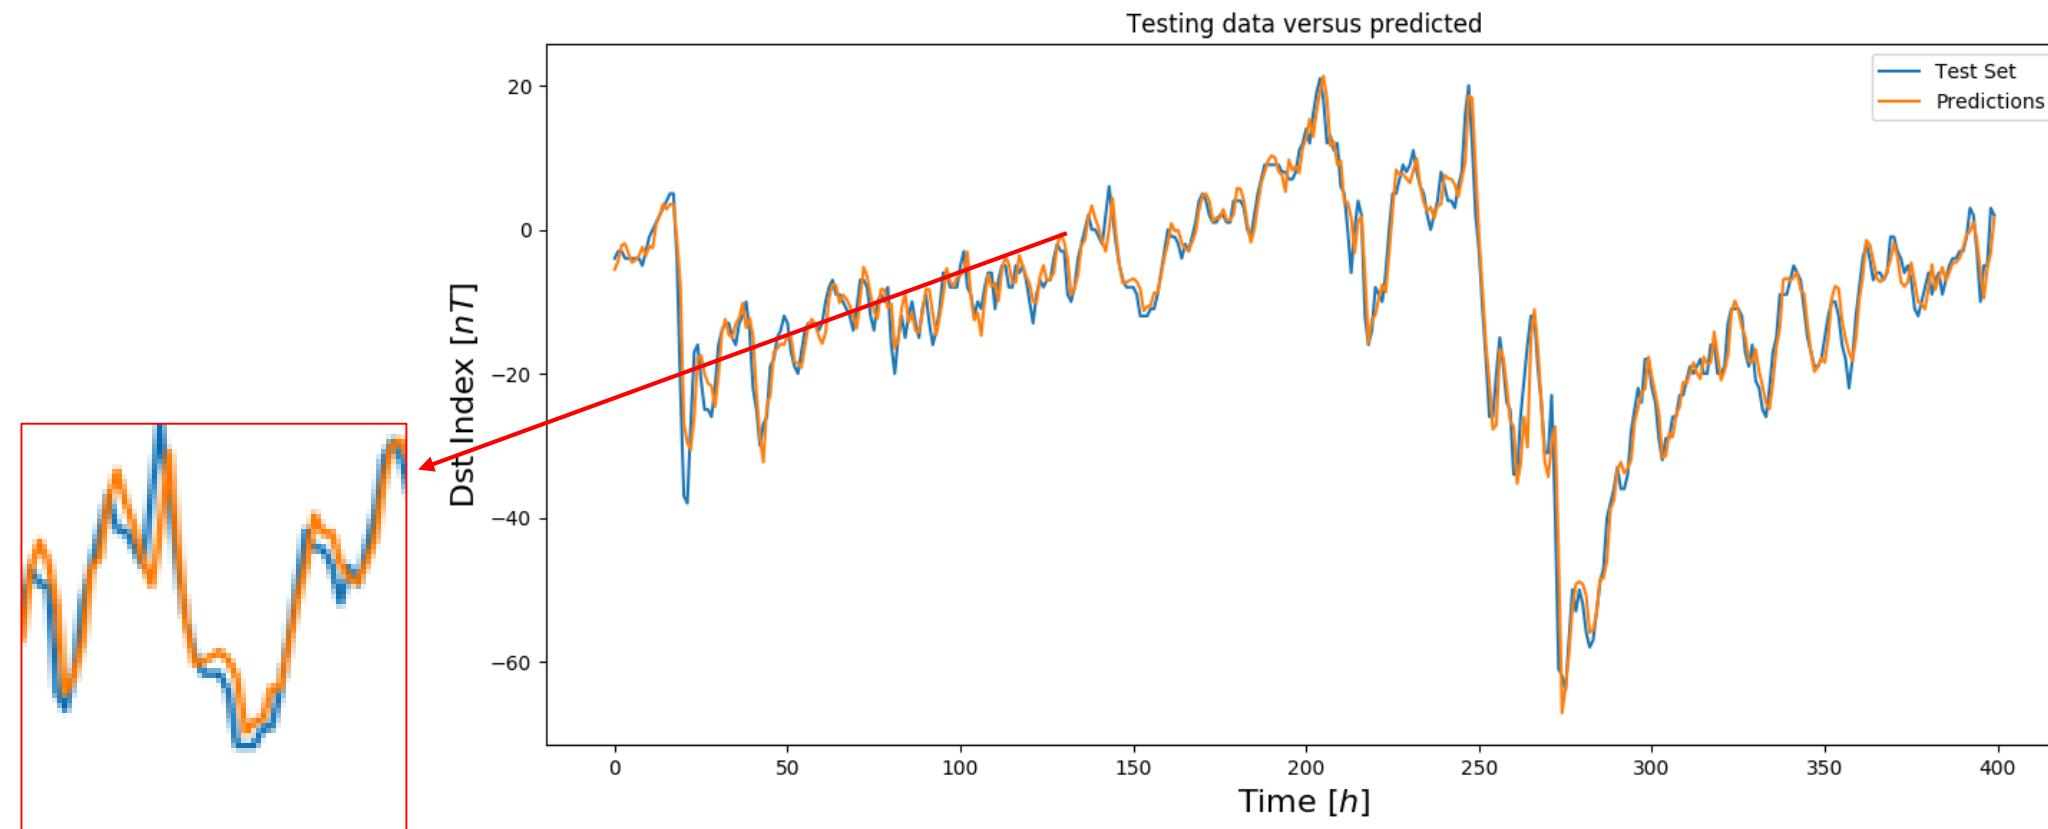
- `t_fcast = 1`
- `t_fcast = 3`

Prediction of the Dst index

Input: Quantities
from satellite
ACE (L1)

4 Layers of (20, 10, 5, 1) Neurons
500 Epochs

Output: DST index
(1 Hour later)



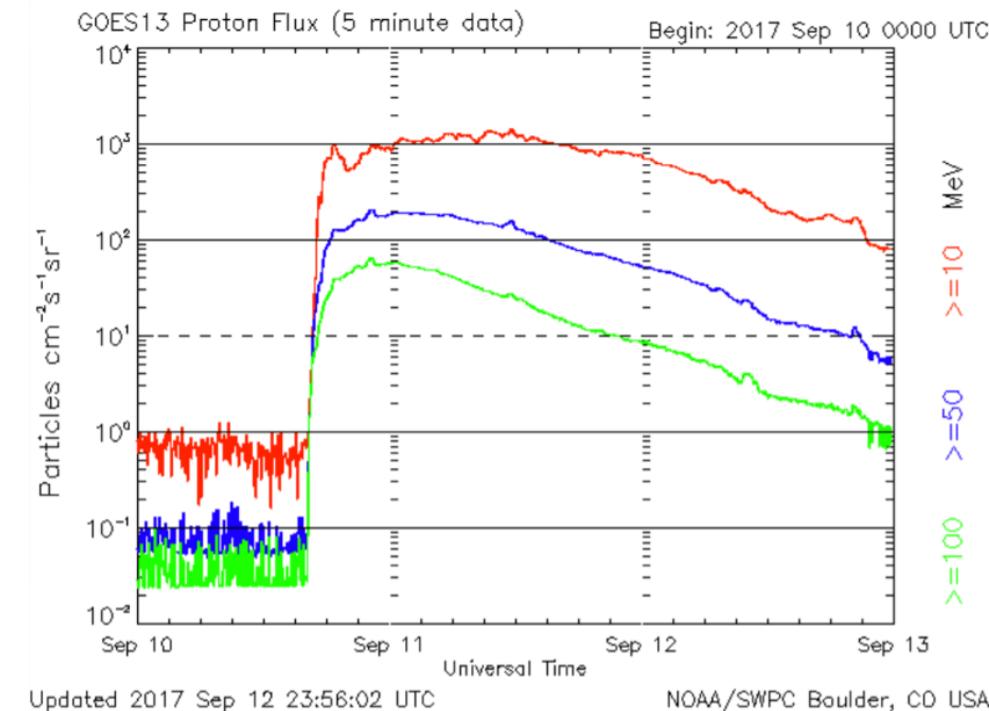
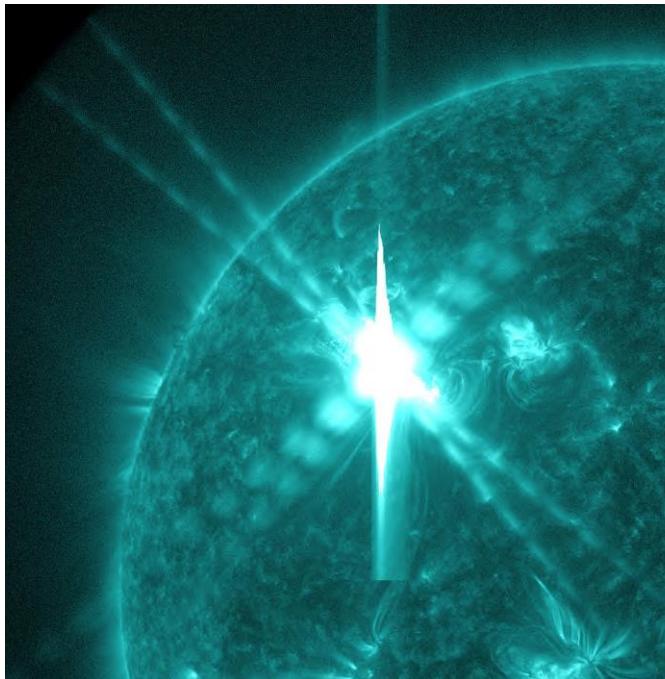
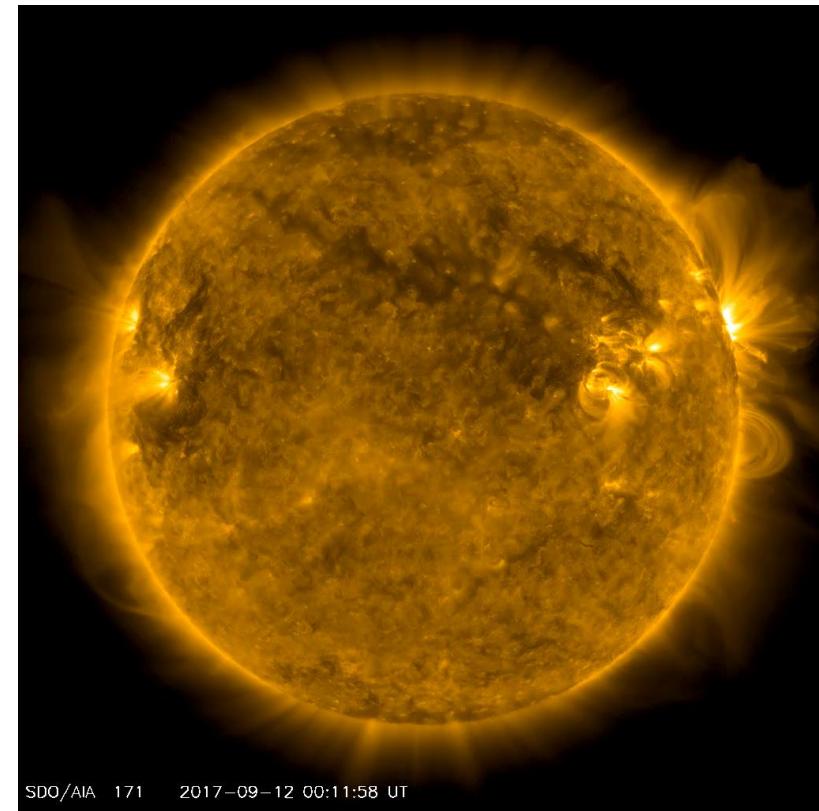
Anything fishy ?? Is this a good prediction ??



**HANDS-ON
LEARNING**

Solar Energetic Particle Prediction (Classification problem)

Solar Energetic Particles (SEPs) & Flares



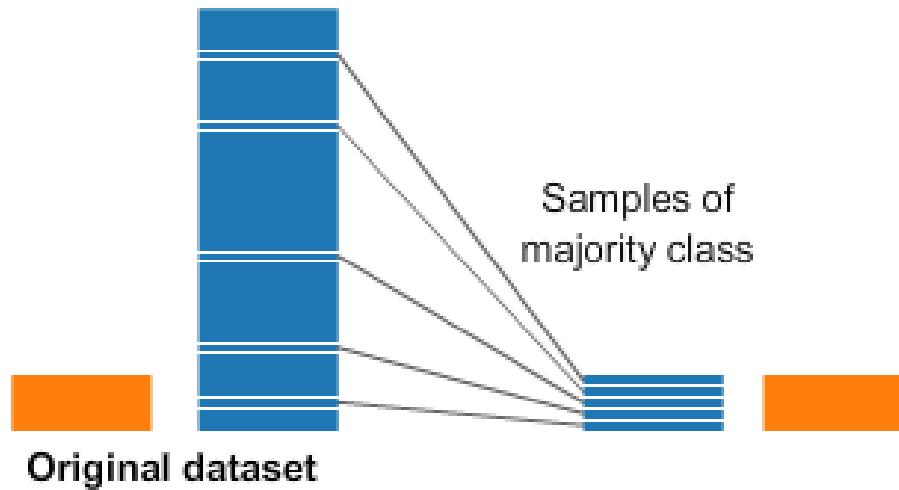
Flares = X-ray emission

High energy proton flux
SEP

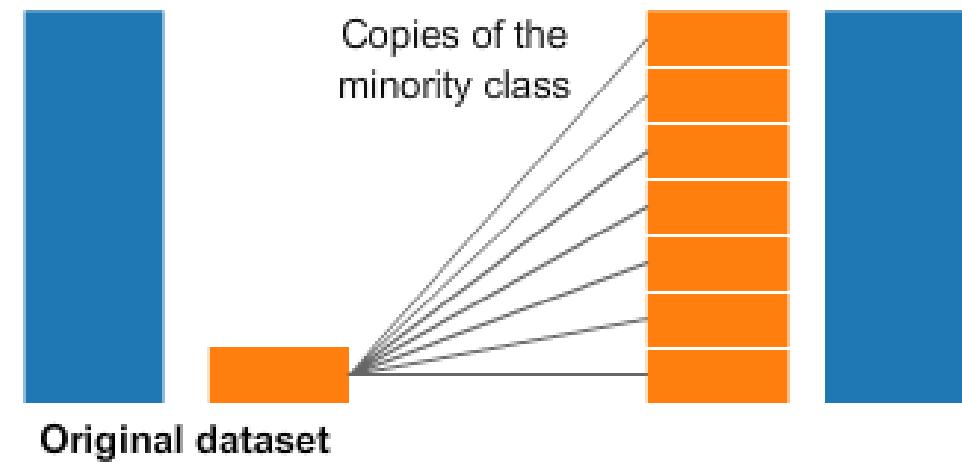
Classes in real datasets

Question: In the examples we show before... how “balanced” was the dataset?

Undersampling

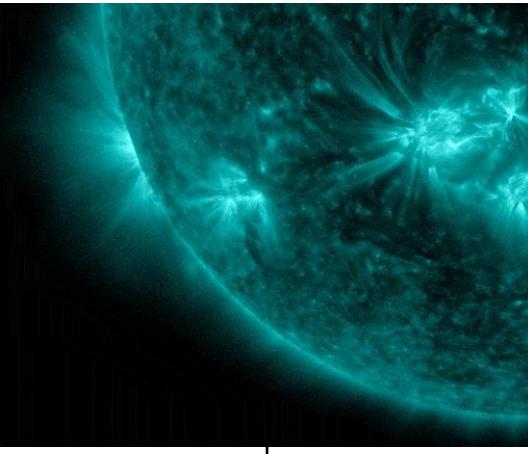


Oversampling



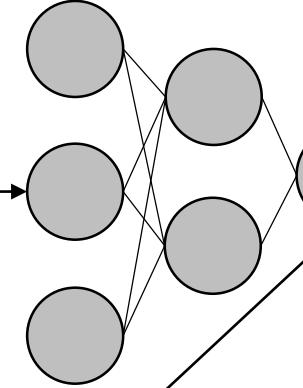
Application on forecasting SEPs

Problem description



24 features

Modeling



Validation & Results

	Predicted Yes	Predicted No
Actual Yes	TN	FP
Actual No	FN	TP

Occurrence of SEP based on X-rays of flares

Comparisons to other models recently published

Review of Solar Energetic Particle Models (Advances in Space Research)



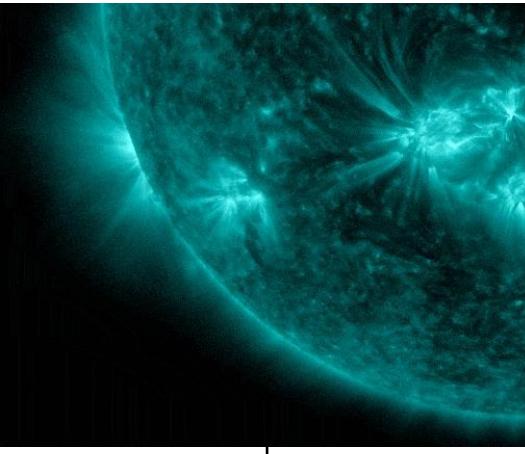
<https://github.com/SavvasRaptis/machine-learning-examples>

Try for example:

- 50 epochs
- 300 epochs

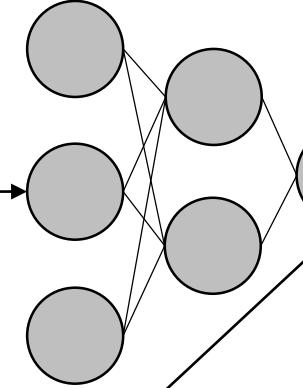
Application on forecasting SEPs

Problem description



24 features

Modeling



Validation & Results

	Predicted Yes	Predicted No
Actual Yes	TN	FP
Actual No	FN	TP

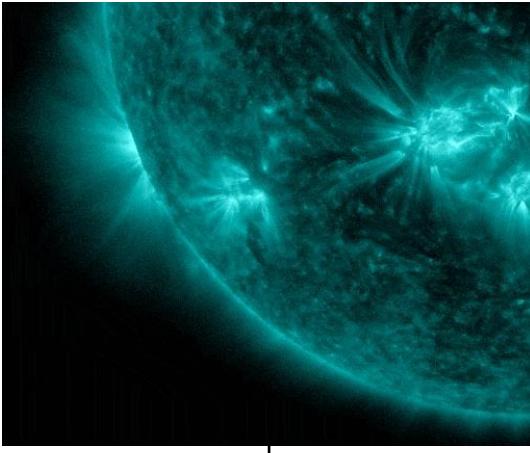
Occurrence of SEP based on X-rays of flares

Comparisons to other models recently published

Review of Solar Energetic Particle Models (Advances in Space Research)

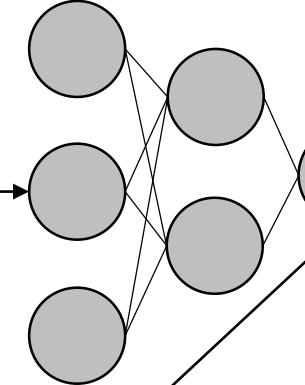
Application on forecasting SEPs

Problem description



24 features

Modeling



Validation & Results

	SEP predicted always YES	SEP predicted always NO
SEP occurred YES	191/220 [86.81 %]	19/220 [8.63 %]
SEP occurred NO	[7.77 %]	[92.23 %]

Occurrence of SEP based on X-rays of flares

Comparisons to other models recently published

Review of Solar Energetic Particle Models (Advances in Space Research)

Learn more on machine learning & space physics

EF2245 Space Physics II 7.5 credits

Before course selection



[Administer](#) [About course](#)

The course is a continuation of Space Physics I, and aims to give the student a more quantitative knowledge of space plasma phenomena and research methods.

DD2424 Deep Learning in Data Science 7.5 credits

Before course selection



[Administer](#) [About course](#)

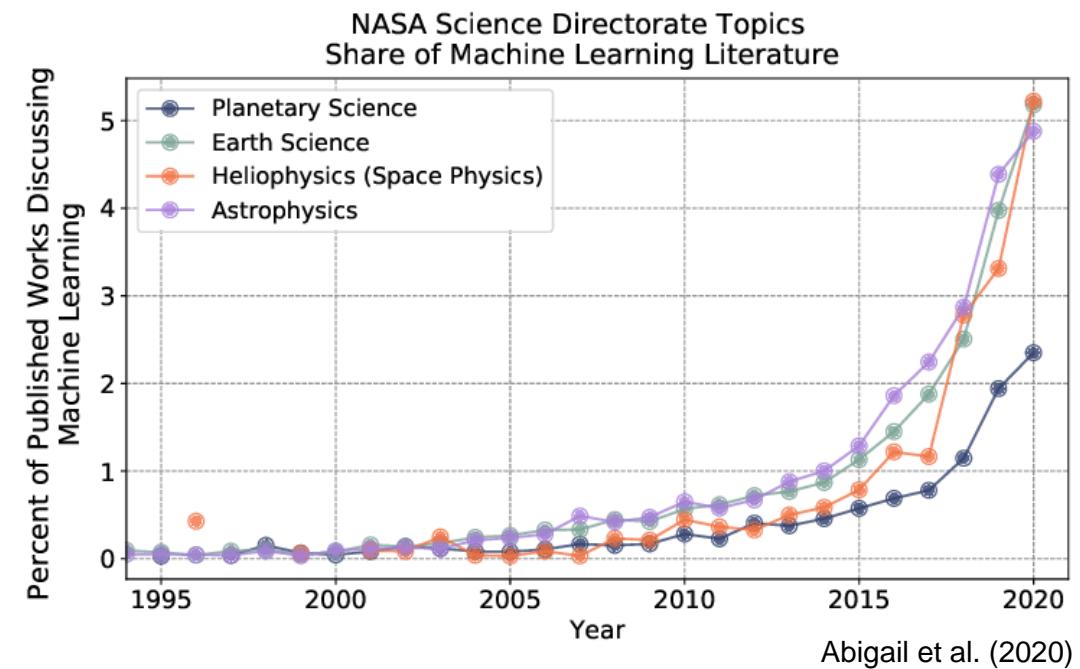
Final notes

- We only saw a small piece of the puzzle today. Applications in physics include:
 - *Unsupervised* learning (clustering of magnetospheric environments)
 - *Reinforcement* learning (optimize when and where to make observations with telescopes)
 - Physics – Informed Neural Networks (PINNS). (e.g., diffusion in radiation belts)

I will be more than glad to share information about these topics if you want to learn more & provide material.

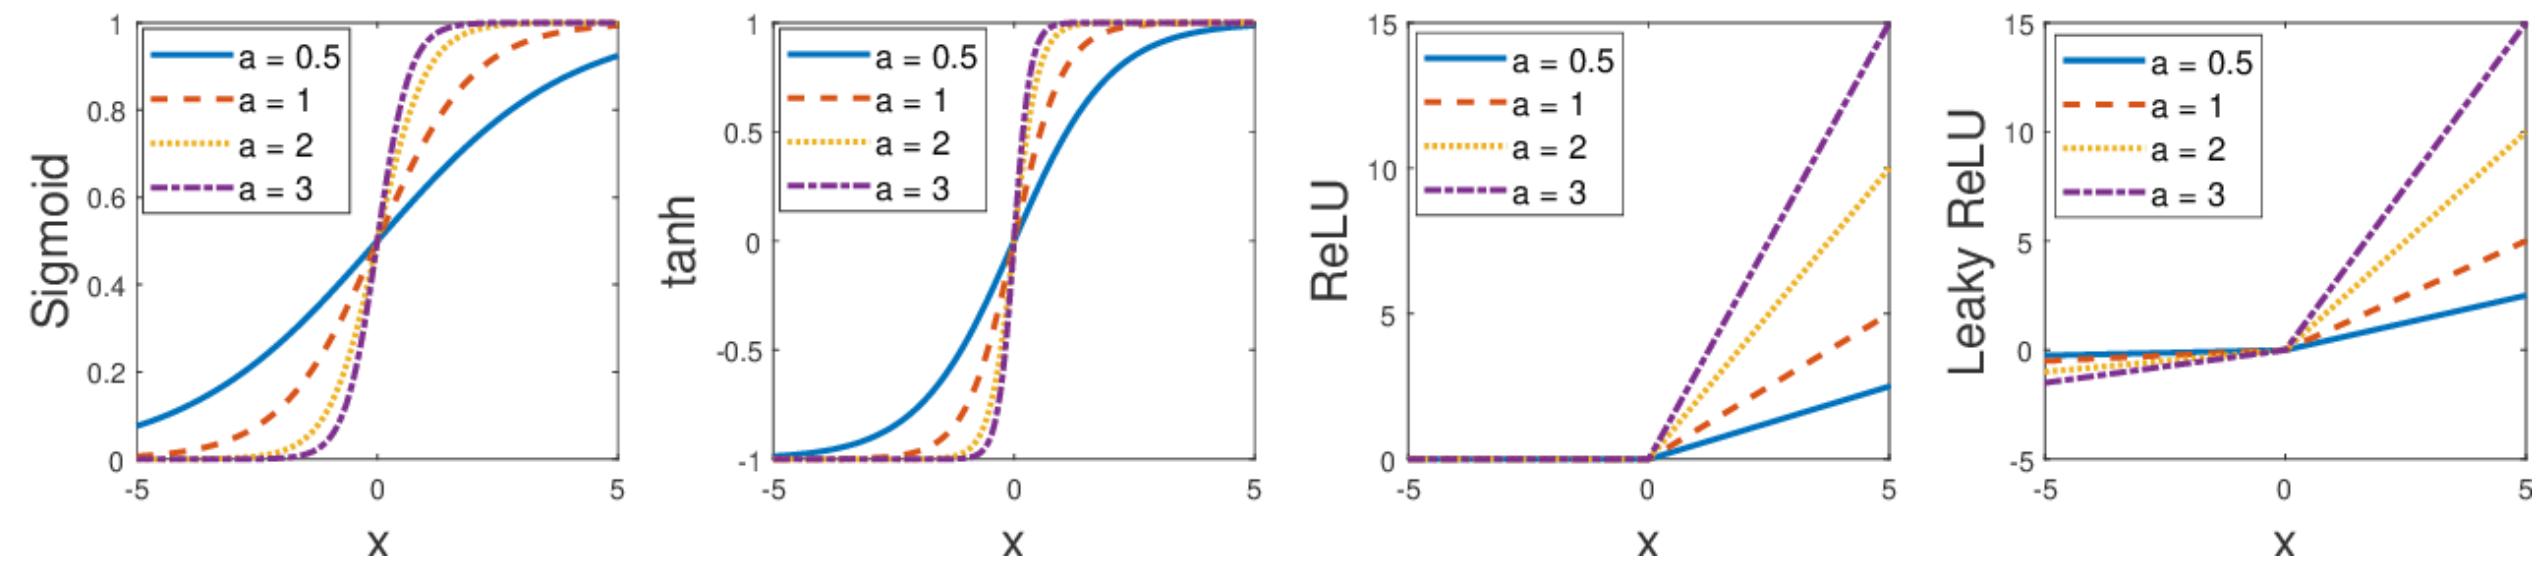
Summary & Take-Home Message

- Machine learning techniques are already used a lot in research, both in academia and in industry.
- Things move very quickly, PINNs that I mentioned in the previous slide are just here for the last ~5 years and there are hundreds of articles on them already.



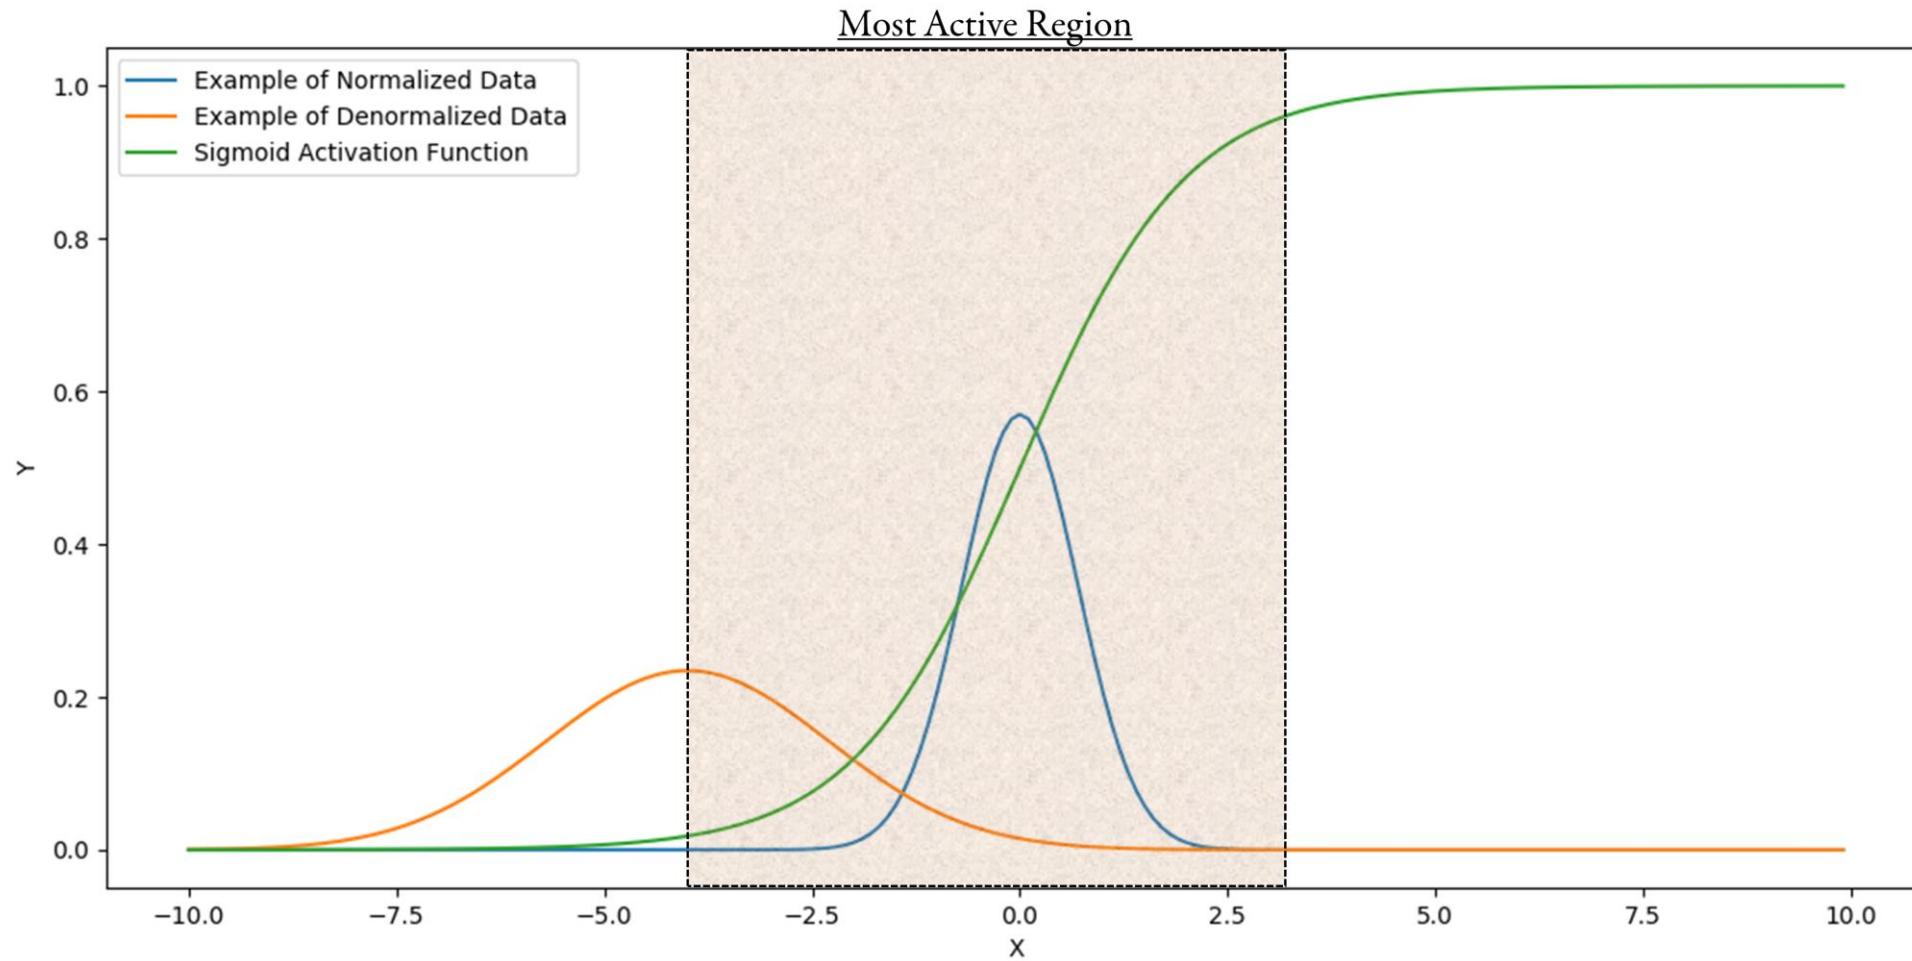
EXTRAS

Adaptive activation functions

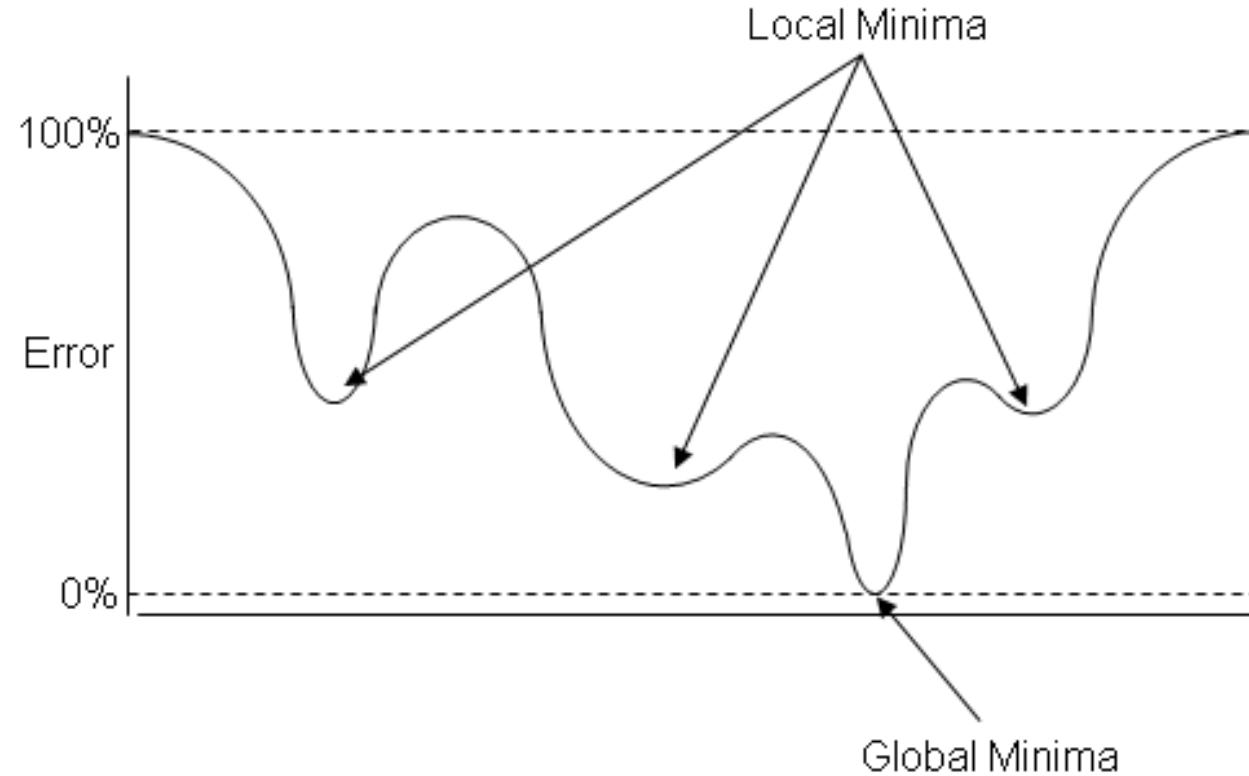


TLDR: optimize the activation function \forall neurons during training = no hardcoded a

Why normalization is vital?



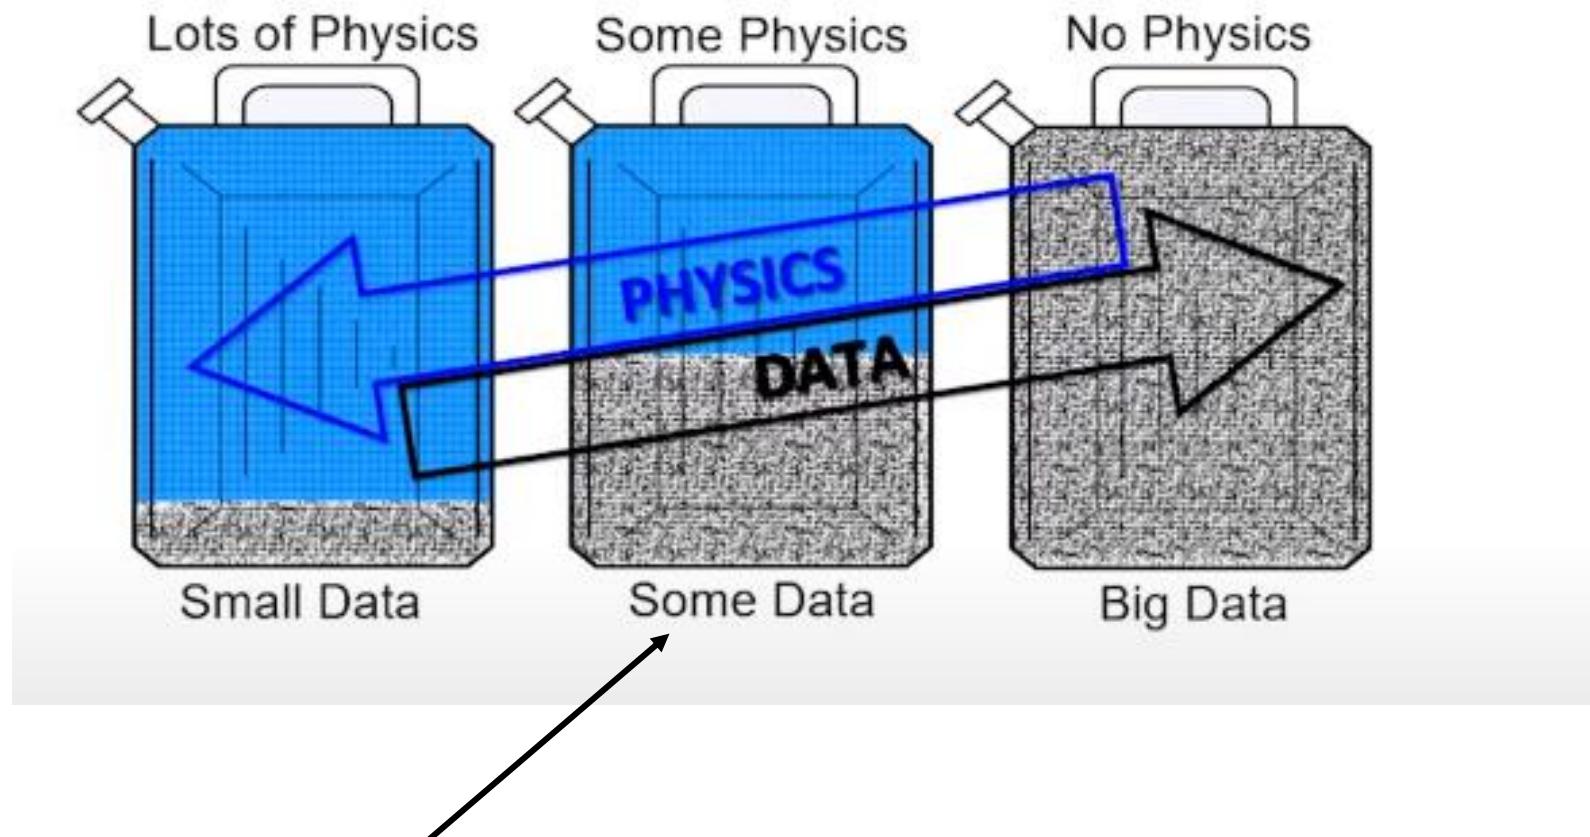
“Proven” to solve a vital problem in optimization



TLDR: Huge problem in traditional machine learning a few years ago = can be solved now
Kenji Kawaguchi et. al (2019)

Physics Informed Machine Learning

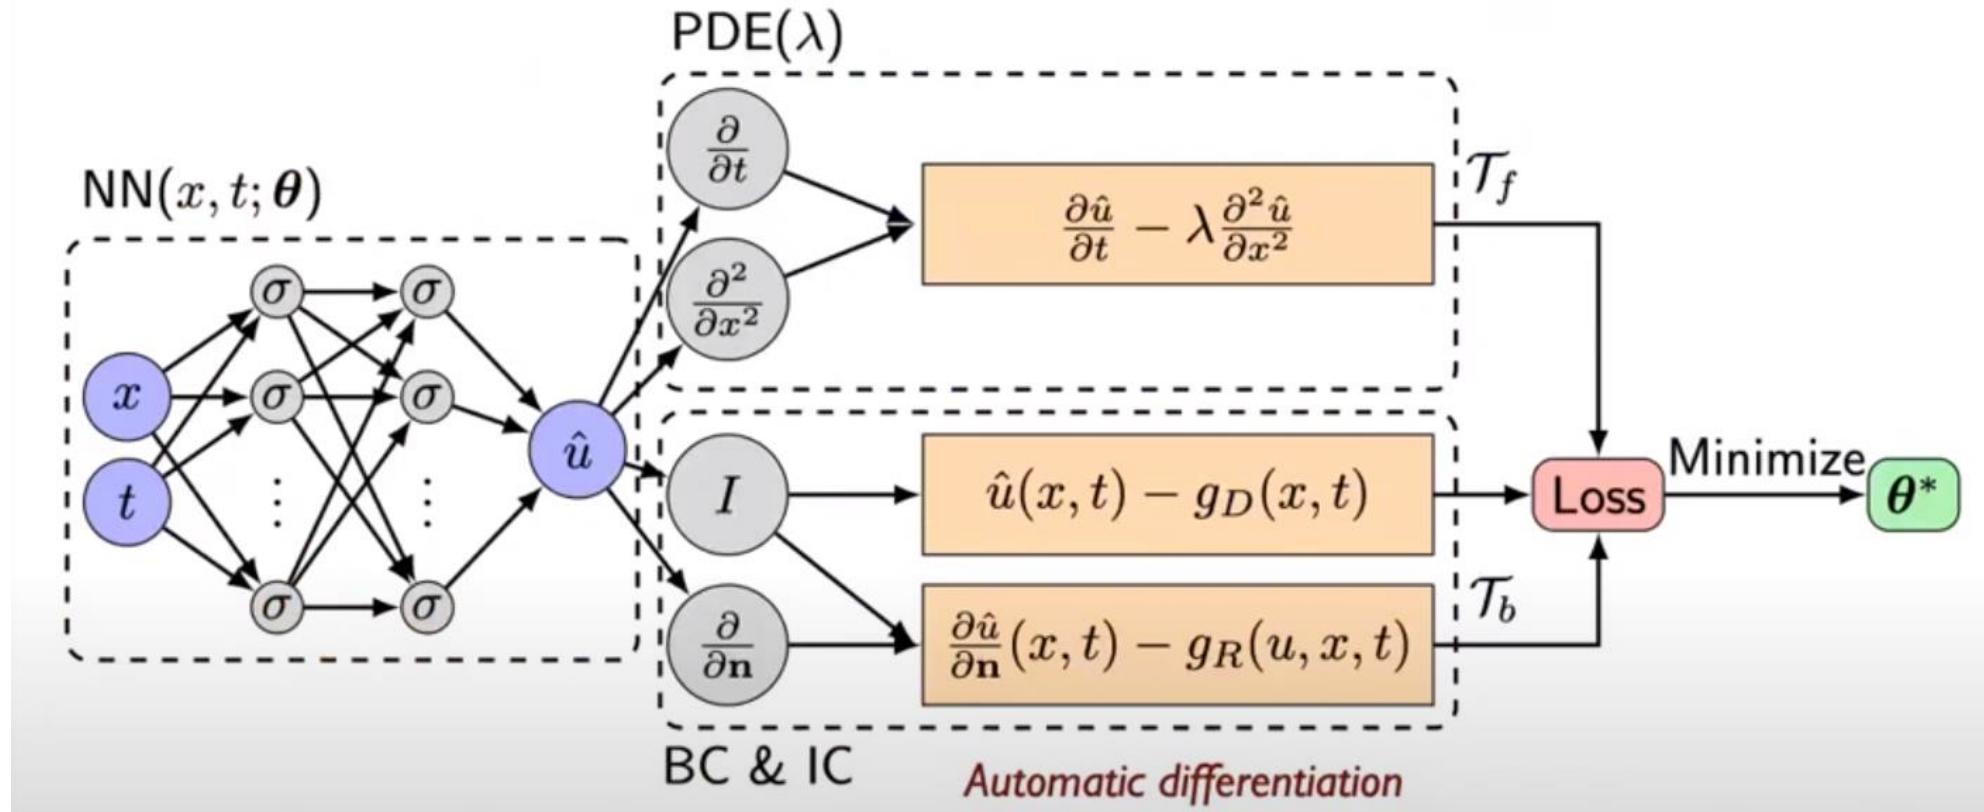
Three scenarios of Physics-Informed Learning Machines



Usual case in certain fields (fluid mechanics, space, plasma, medicine etc.)

What exactly is a PINN then ?

$$f \left(\mathbf{x}; \frac{\partial u}{\partial x_1}, \dots, \frac{\partial u}{\partial x_d}; \frac{\partial^2 u}{\partial x_1 \partial x_1}, \dots, \frac{\partial^2 u}{\partial x_1 \partial x_d}; \dots; \boldsymbol{\lambda} \right) = 0, \quad \mathbf{x} \in \Omega \quad \mathcal{B}(u, \mathbf{x}) = 0 \quad \text{on} \quad \partial\Omega.$$



$$\text{loss} = \text{data fit} + \text{PDE residual} + \text{ICs fit} + \text{BCs fit}$$

Idea = If output is a differentiable quantity with respect to an input = can compute PDEs = combined loss functions

2 type of problems

Data-driven solutions to partial differential equations

- 1) Burger's equation
- 2) Schrodinger Equation

Data-driven discovery of nonlinear partial differential equations

- 3) Navier-Stokes Equation

Data-driven solutions (forward problem)

Example 1 – Burger's equation

$$u_t + uu_x - (0.01/\pi)u_{xx} = 0, \quad x \in [-1, 1], \quad t \in [0, 1],$$

$$u(0, x) = -\sin(\pi x),$$

$$u(t, -1) = u(t, 1) = 0.$$

Let us define $f(t, x)$ to be given by

$$f := u_t + uu_x - (0.01/\pi)u_{xx},$$

$$MSE = MSE_u + MSE_f,$$

where

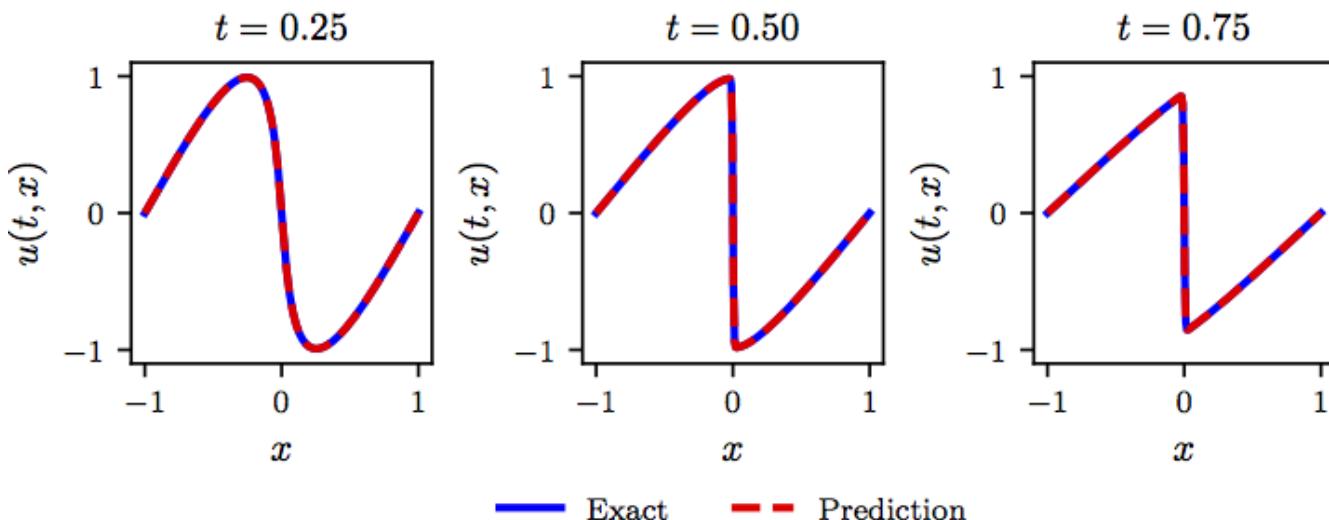
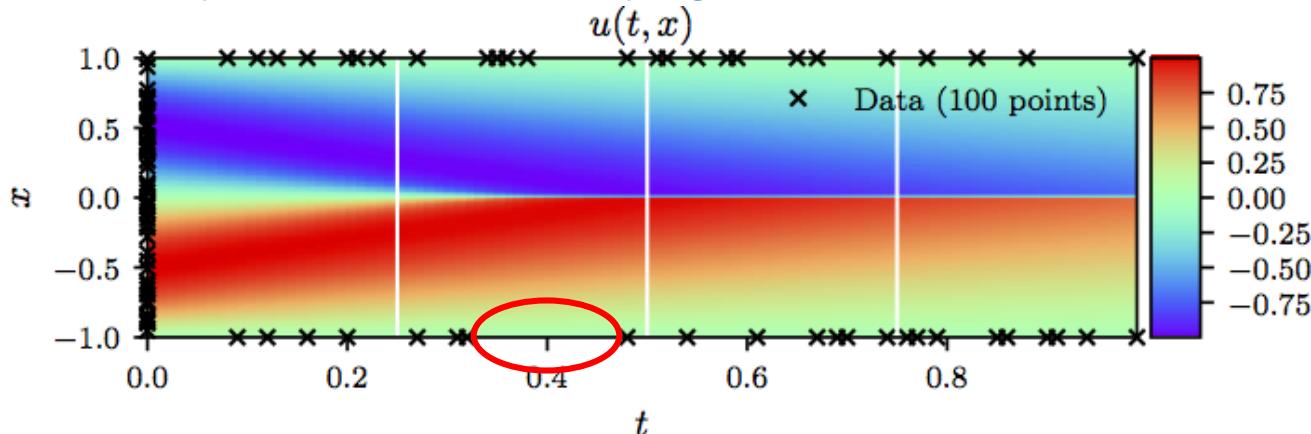
$$MSE_u = \frac{1}{N_u} \sum_{i=1}^{N_u} |u(t_u^i, x_u^i) - u^i|^2,$$

and

$$MSE_f = \frac{1}{N_f} \sum_{i=1}^{N_f} |f(t_f^i, x_f^i)|^2.$$

Here we basically enforce initial and boundary conditions.

Penalty of points not satisfying f / residual loss



Example 2 – Schrodinger equation

$$ih_t + 0.5h_{xx} + |h|^2h = 0, \quad x \in [-5, 5], \quad t \in [0, \pi/2],$$

$$h(0, x) = 2 \operatorname{sech}(x),$$

$$h(t, -5) = h(t, 5),$$

$$h_x(t, -5) = h_x(t, 5),$$

$$f := ih_t + 0.5h_{xx} + |h|^2h,$$

$$MSE = MSE_0 + MSE_b + MSE_f,$$

where

$$MSE_0 = \frac{1}{N_0} \sum_{i=1}^{N_0} |h(0, x_0^i) - h_0^i|^2,$$

$$MSE_b = \frac{1}{N_b} \sum_{i=1}^{N_b} \left(|h^i(t_b^i, -5) - h^i(t_b^i, 5)|^2 + |h_x^i(t_b^i, -5) - h_x^i(t_b^i, 5)|^2 \right)$$

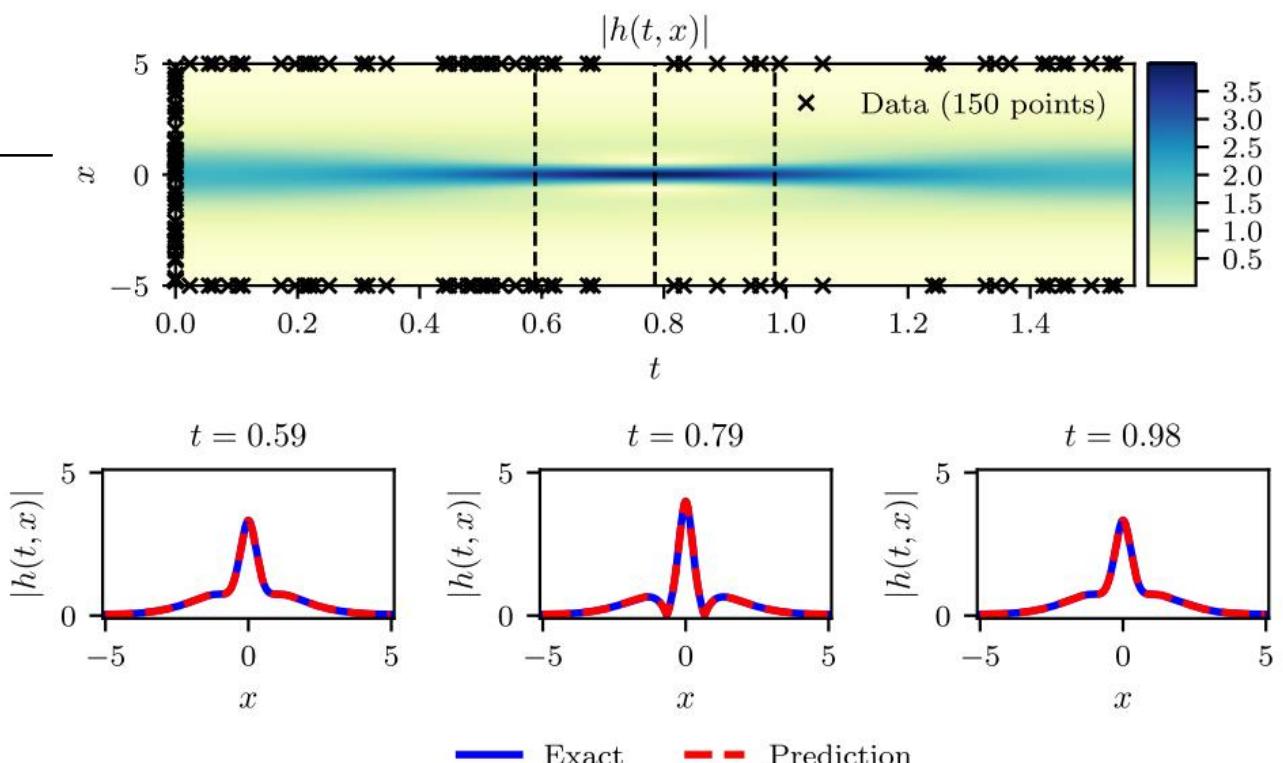
and

$$MSE_f = \frac{1}{N_f} \sum_{i=1}^{N_f} |f(t_f^i, x_f^i)|^2.$$

Same thing as we saw before (initial)

Basically enforcing the above conditions (boundary)

Penalty of points not satisfying f (residuals)



Data-driven discoveries (inverse problem)

Example 3 – Navier–Stokes equation (1)

$$\begin{aligned} u_t + \lambda_1(uu_x + vu_y) &= -p_x + \lambda_2(u_{xx} + u_{yy}), \\ v_t + \lambda_1(uv_x + vv_y) &= -p_y + \lambda_2(v_{xx} + v_{yy}), \end{aligned}$$

Basically imposing boundary conditions

The system of differential equations

Also, continuity equation/divergence-free result:

$$u_x + v_y = 0.$$

Goal: Estimate $\lambda_{1,2}$ and $p(t, x, y)$

So basically we define the following system (f,g):

$$\begin{aligned} f &:= u_t + \lambda_1(uu_x + vu_y) + p_x - \lambda_2(u_{xx} + u_{yy}), \\ g &:= v_t + \lambda_1(uv_x + vv_y) + p_y - \lambda_2(v_{xx} + v_{yy}), \end{aligned}$$

First step: Generate some data for steady state solution

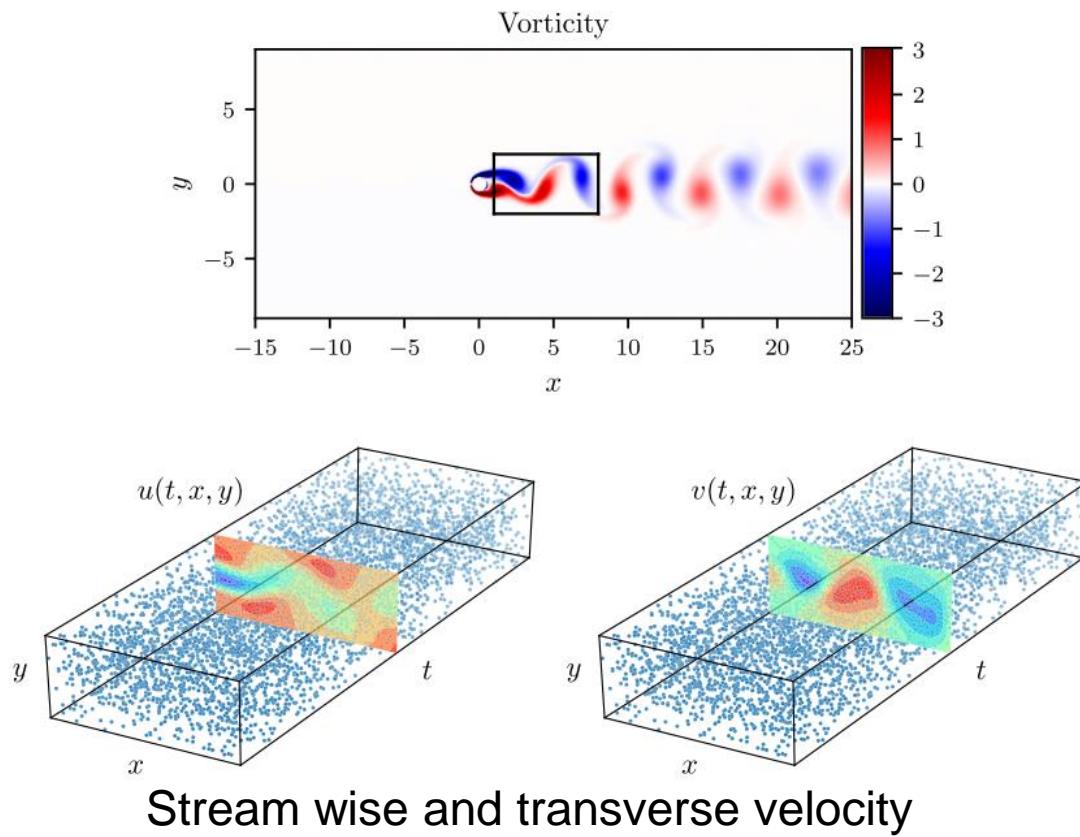
Second step: minimize MSE to satisfy boundary conditions for $t \neq 0$ and the general PDE.

$$MSE := \boxed{\frac{1}{N} \sum_{i=1}^N (|u(t^i, x^i, y^i) - u^i|^2 + |v(t^i, x^i, y^i) - v^i|^2)} + \boxed{\frac{1}{N} \sum_{i=1}^N (|f(t^i, x^i, y^i)|^2 + |g(t^i, x^i, y^i)|^2)}.$$

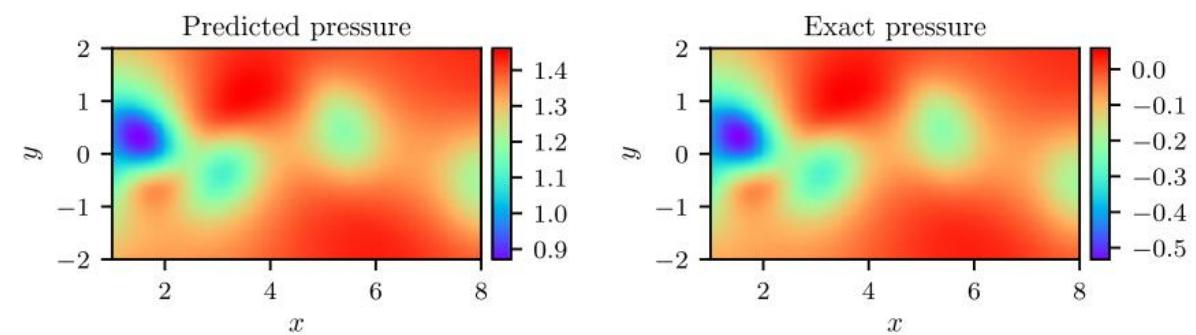
Example 3 – Navier–Stokes equation (2)

5000 scattered data points

Training Data

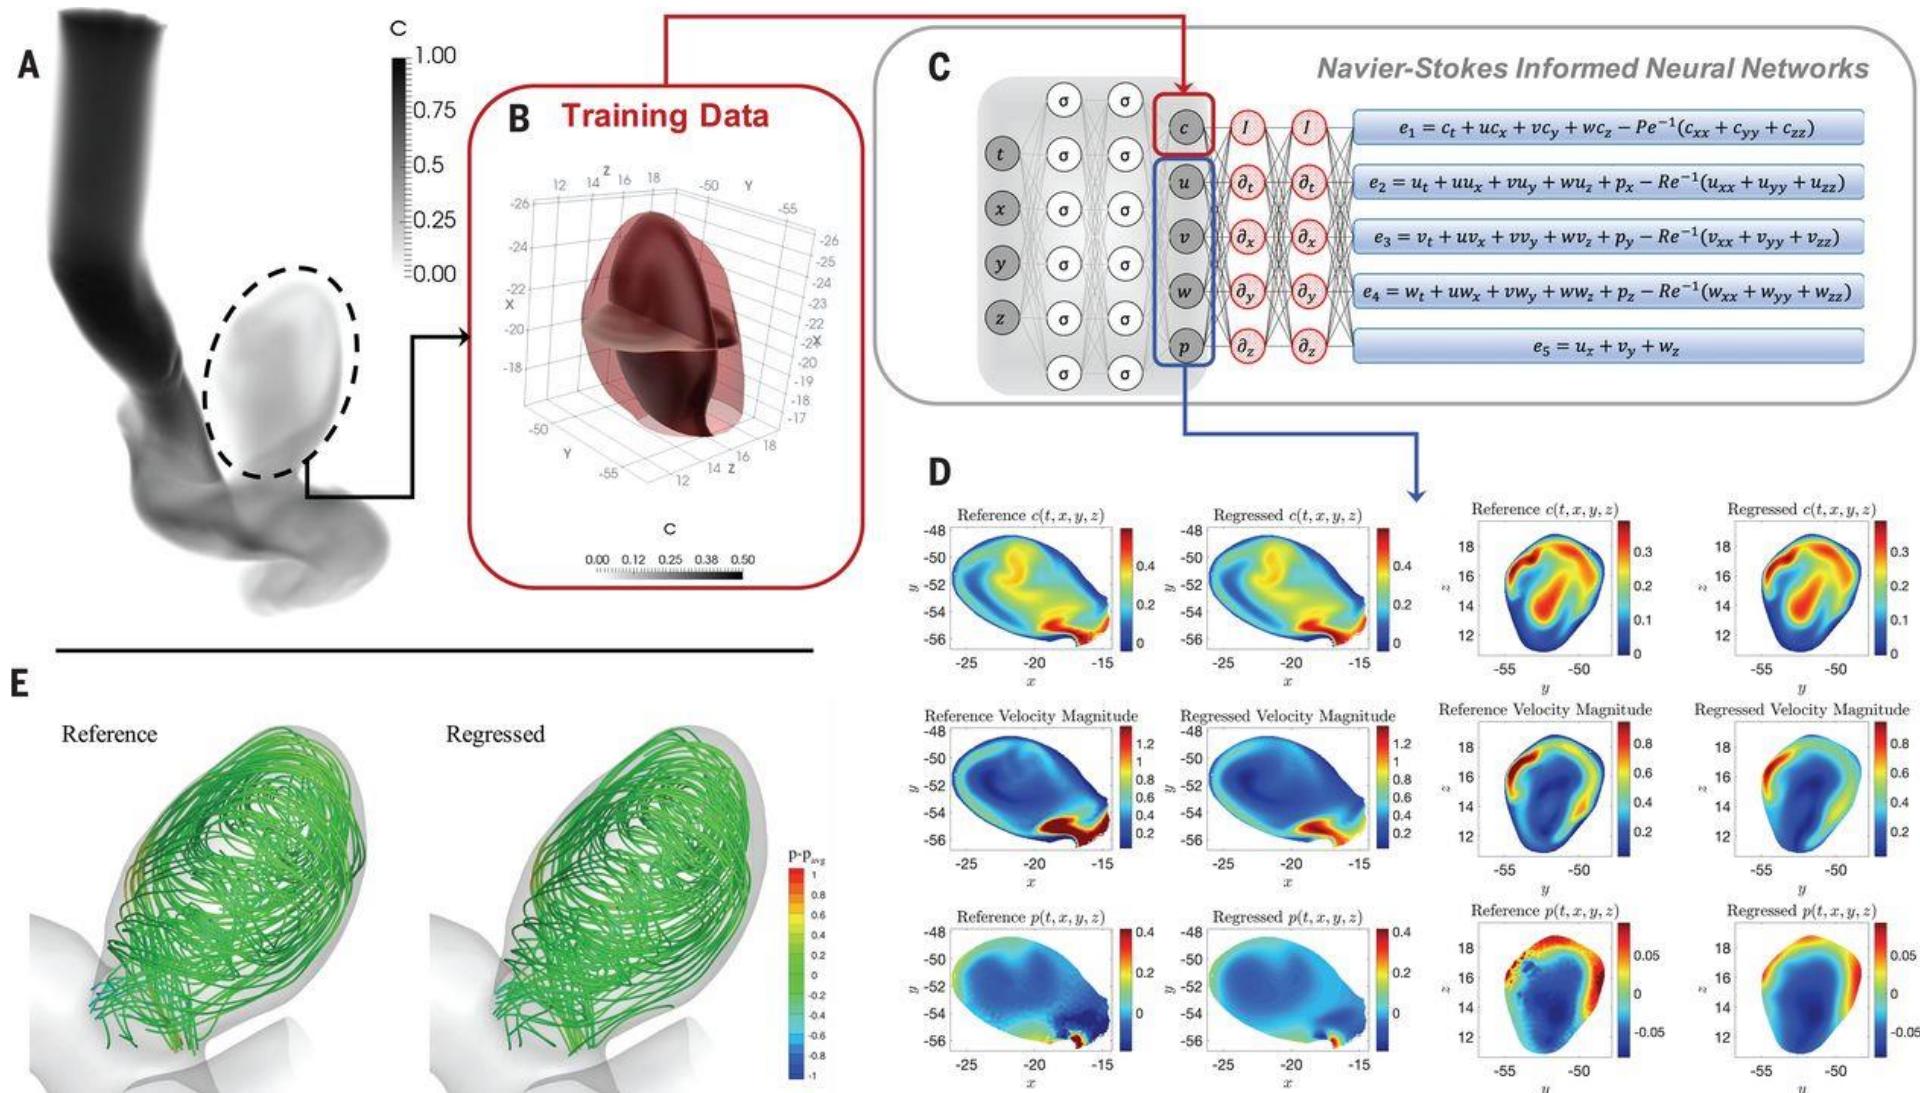


2 Results



Correct PDE	$u_t + (uu_x + vu_y) = -p_x + 0.01(u_{xx} + u_{yy})$ $v_t + (uv_x + vv_y) = -p_y + 0.01(v_{xx} + v_{yy})$
Identified PDE (clean data)	$u_t + 0.999(uu_x + vu_y) = -p_x + 0.01047(u_{xx} + u_{yy})$ $v_t + 0.999(uv_x + vv_y) = -p_y + 0.01047(v_{xx} + v_{yy})$
Identified PDE (1% noise)	$u_t + 0.998(uu_x + vu_y) = -p_x + 0.01057(u_{xx} + u_{yy})$ $v_t + 0.998(uv_x + vv_y) = -p_y + 0.01057(v_{xx} + v_{yy})$

Example 4 – Hidden Fluid Mechanics



Maziar Raissi et. al (2020) Science (<https://science.science.org/content/367/6481/1026>)

https://maziarraissi.github.io/research/09_hidden_fluid_mechanics/

Bayesian Physics Informed Neural Networks (B-PINNs)

Statistics > Machine Learning

[Submitted on 13 Mar 2020]

B-PINNs: Bayesian Physics-Informed Neural Networks for Forward and Inverse PDE Problems with Noisy Data

Liu Yang, Xuhui Meng, George Em Karniadakis

We propose a Bayesian physics-informed neural network (B-PINN) to solve both forward and inverse nonlinear problems described by partial differential equations (PDEs) and noisy data. In this Bayesian framework, the Bayesian neural network (BNN) combined with a PINN for PDEs serves as the prior while the Hamiltonian Monte Carlo (HMC) or the variational inference (VI) could serve as an estimator of the posterior. B-PINNs make use of both physical laws and scattered noisy measurements to provide predictions and quantify the aleatoric uncertainty arising from the noisy data in the Bayesian framework. Compared with PINNs, in addition to uncertainty quantification, B-PINNs obtain more accurate predictions in scenarios with large noise due to their capability of avoiding overfitting. We conduct a systematic comparison between the two different approaches for the B-PINN posterior estimation (i.e., HMC or VI), along with dropout used for quantifying uncertainty in deep neural networks. Our experiments show that HMC is more suitable than VI for the B-PINNs posterior estimation, while dropout employed in PINNs can hardly provide accurate predictions with reasonable uncertainty. Finally, we replace the BNN in the prior with a truncated Karhunen-Loève (KL) expansion combined with HMC or a deep normalizing flow (DNF) model as posterior estimators. The KL is as accurate as BNN and much faster but this framework cannot be easily extended to high-dimensional problems unlike the BNN based framework.

TLDR: Noisy/Real data → Need uncertainty quantification → Bayesian approach is good → combine

Maybe we can see some examples ?

[Submitted on 11 May 2020 ([v1](#)), last revised 16 Sep 2020 (this version, v2)]

SciANN: A Keras/Tensorflow wrapper for scientific computations and physics-informed deep learning using artificial neural networks

Ehsan Haghhighat, Ruben Juanes

In this paper, we introduce SciANN, a Python package for scientific computing and physics-informed deep learning using artificial neural networks. SciANN uses the widely used deep-learning packages Tensorflow and Keras to build deep neural networks and optimization models, thus inheriting many of Keras's functionalities, such as batch optimization and model reuse for transfer learning. SciANN is designed to abstract neural network construction for scientific computations and solution and discovery of partial differential equations (PDE) using the physics-informed neural networks (PINN) architecture, therefore providing the flexibility to set up complex functional forms. We illustrate, in a series of examples, how the framework can be used for curve fitting on discrete data, and for solution and discovery of PDEs in strong and weak forms. We summarize the features currently available in SciANN, and also outline ongoing and future developments.